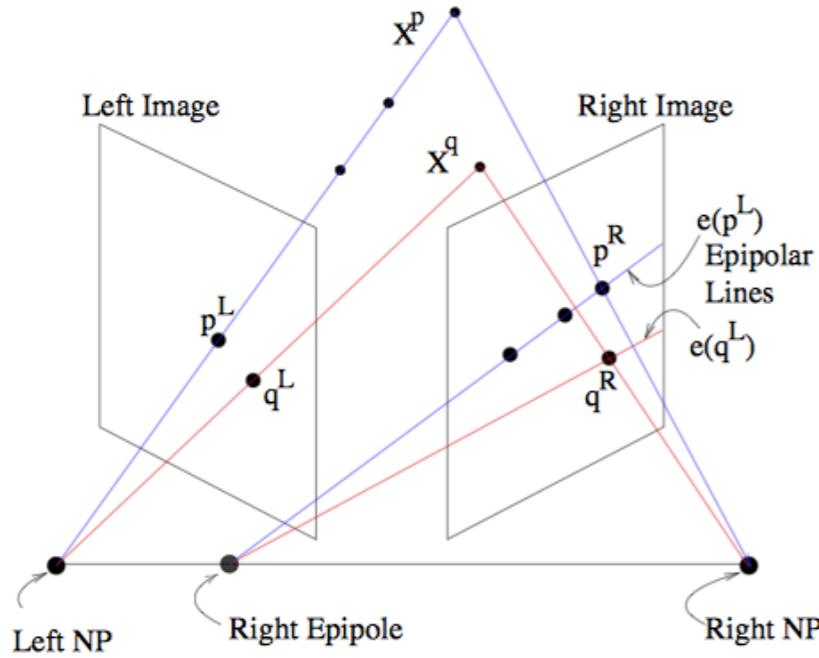


Projection, stereo and panoramic images

X^p and X^q are the physical location of objects.

NP are the nadir points of the cameras.



The projection of the left nadir, NP is seen as a the right epipole with w.r.t. the right image. The epipoles may or may not be within the border of the images.

The projection of X^p to left nadir, NP is seen as an epipole line in the right image. If more than one epipole line is present in the right image, they converge at the right epipole (which might not be within the boundaries of the image).

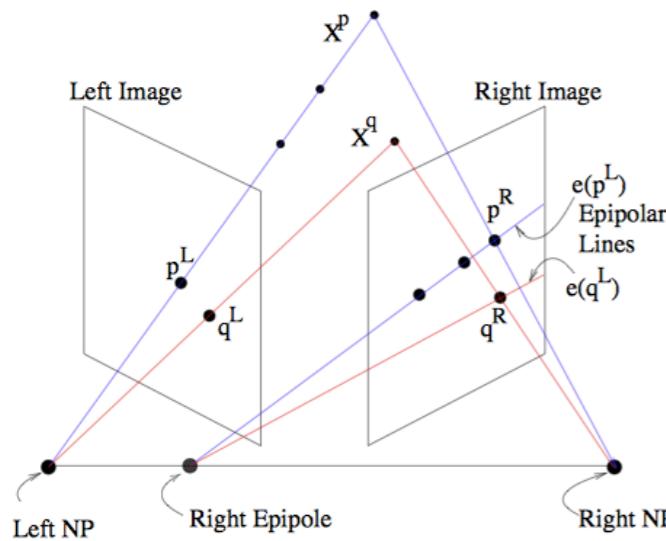
Once the points in the left image, X_L are matched with points in the right image, X_R , if there are at least 7 points, one can determine the “bifocal tensor”,
a.k.a. “fundamental matrix, relating the points in the 2 images using a
3x3 matrix of rank 2.

<http://www.cs.toronto.edu/~jepson/csc420/notes/epiPolarGeom.pdf>
(note, the image is posted on an educational site and copied here without following up on permissions. Any further use of the image should follow up on the origins and permissions.)

$(X_L)^T * F * X_R = 0$ for any pair of points in the images.

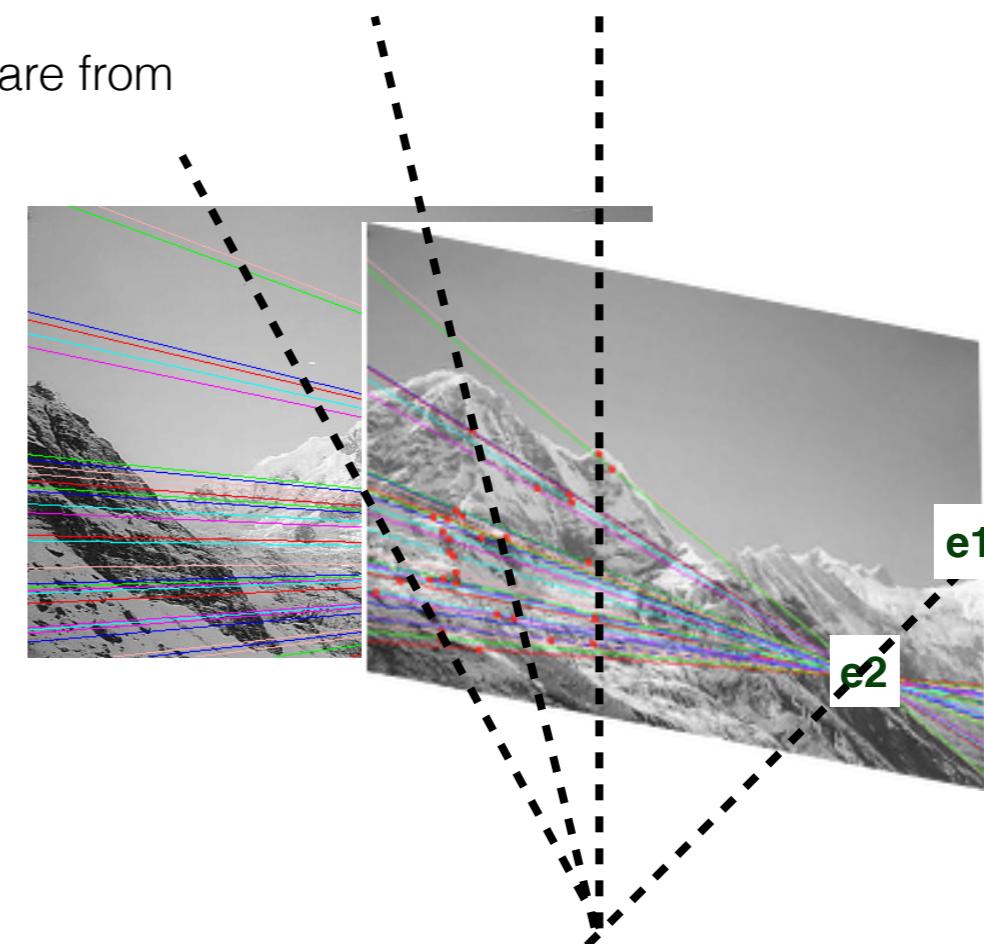
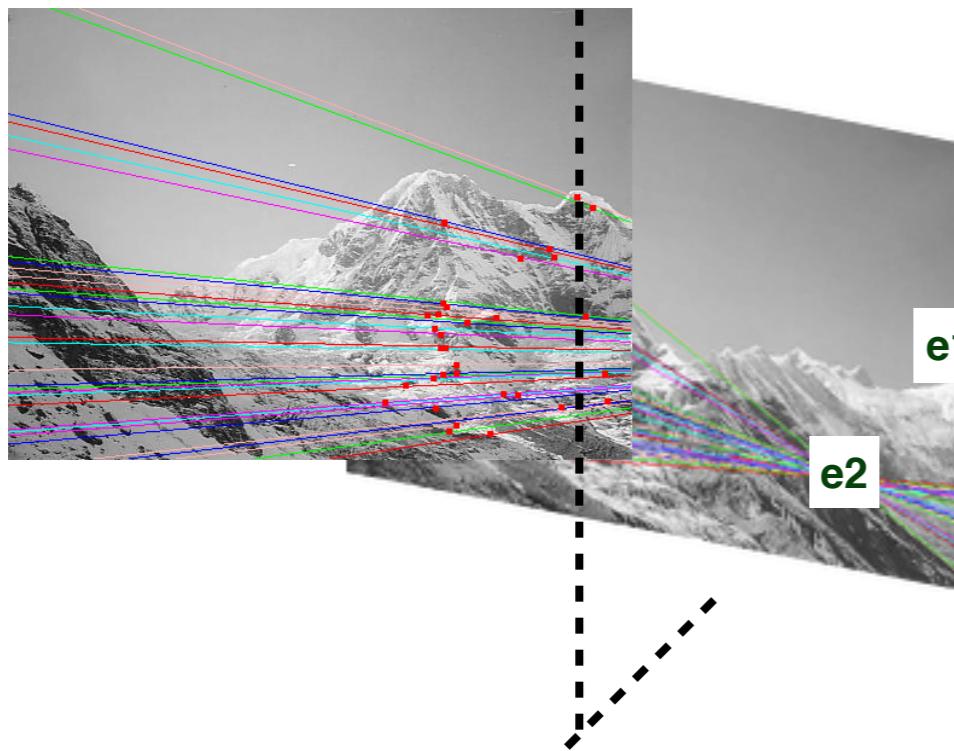
Note: the “Essential matrix” is a matrix used if the camera details are known. The “bifocal tensor”, a.k.a. “fundamental matrix” does not need camera details.

Projection, stereo and panoramic images



<http://www.cs.toronto.edu/~jepson/csc420/notes/epiPolarGeom.pdf>
(note, the image is posted on an educational site and copied here without following up on permissions. Any further use of the image should follow up on the origins and permissions.)

panoramic images here are from
Brown & Lowe 2003

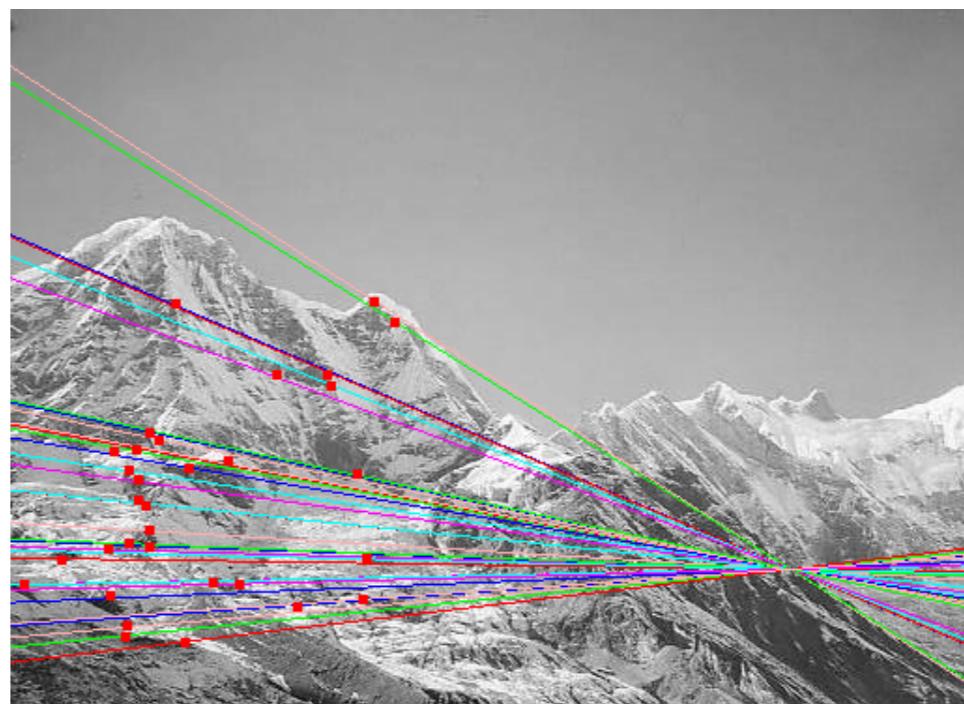
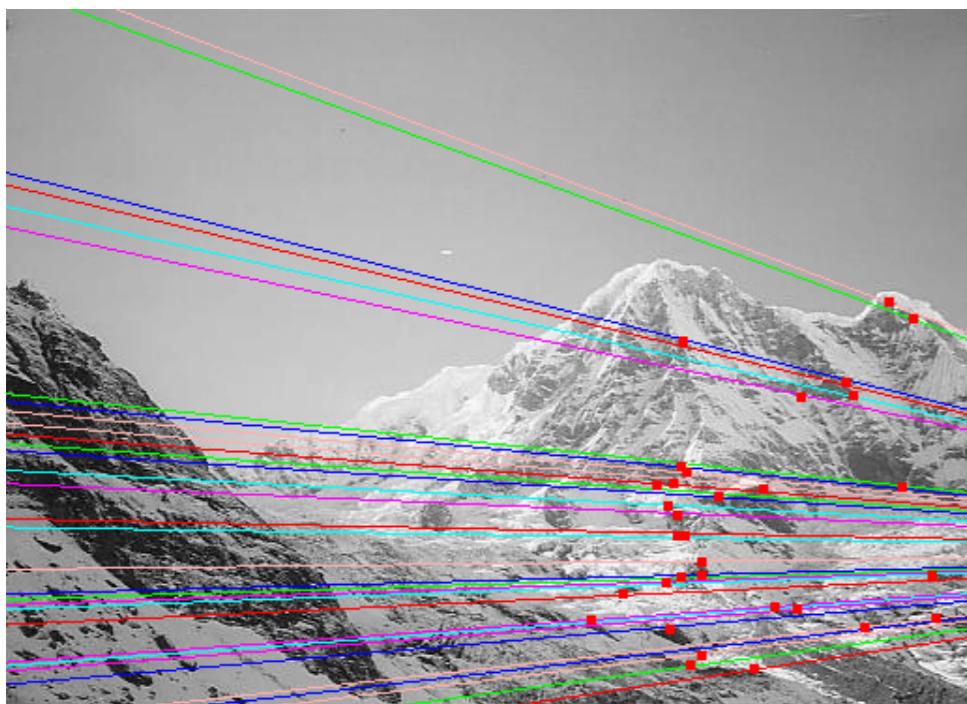


same camera objectives (nadir), but
different orientation for the 2 images
(that is, rotated around the same nadir)

Point Correspondence

Need to create list of matched points between the images in order to solve for the epipolar projection

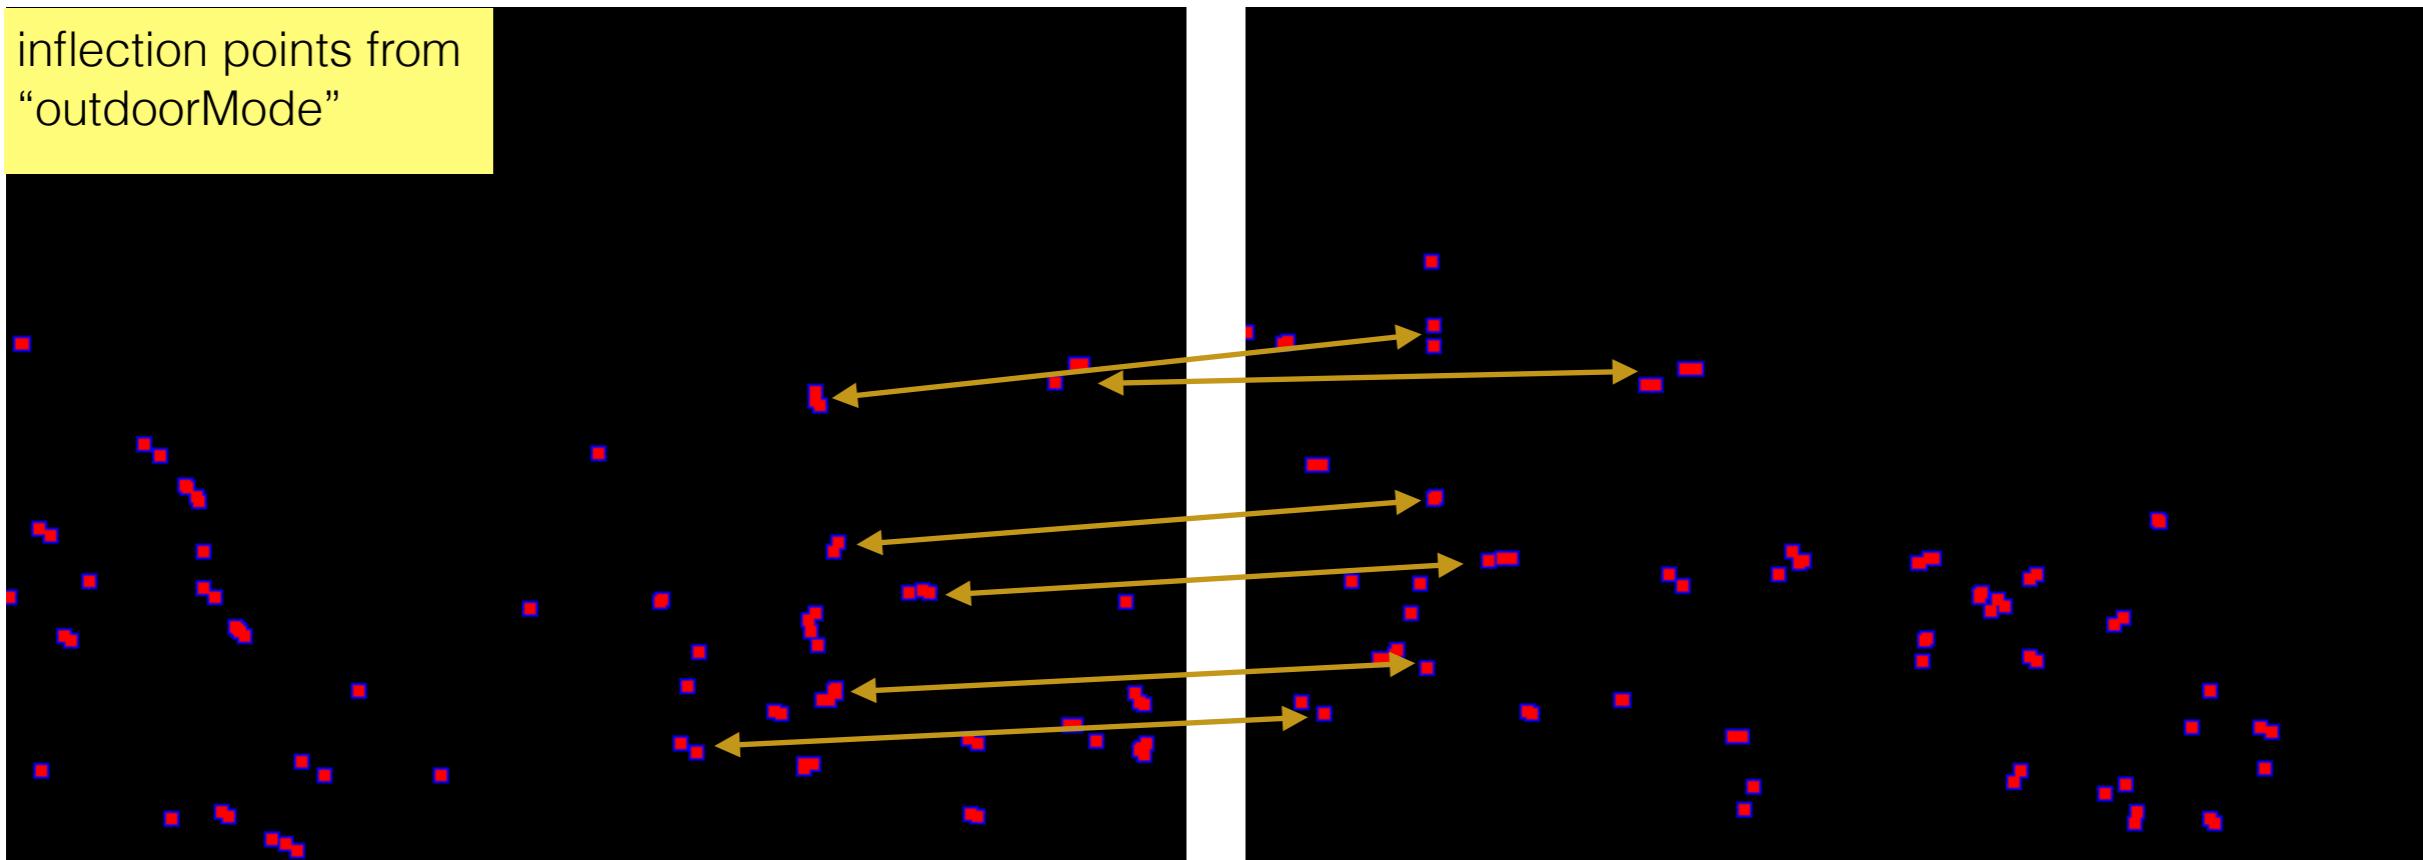
manually making a point list from the corners from the edge extractor used with “outdoor mode”:



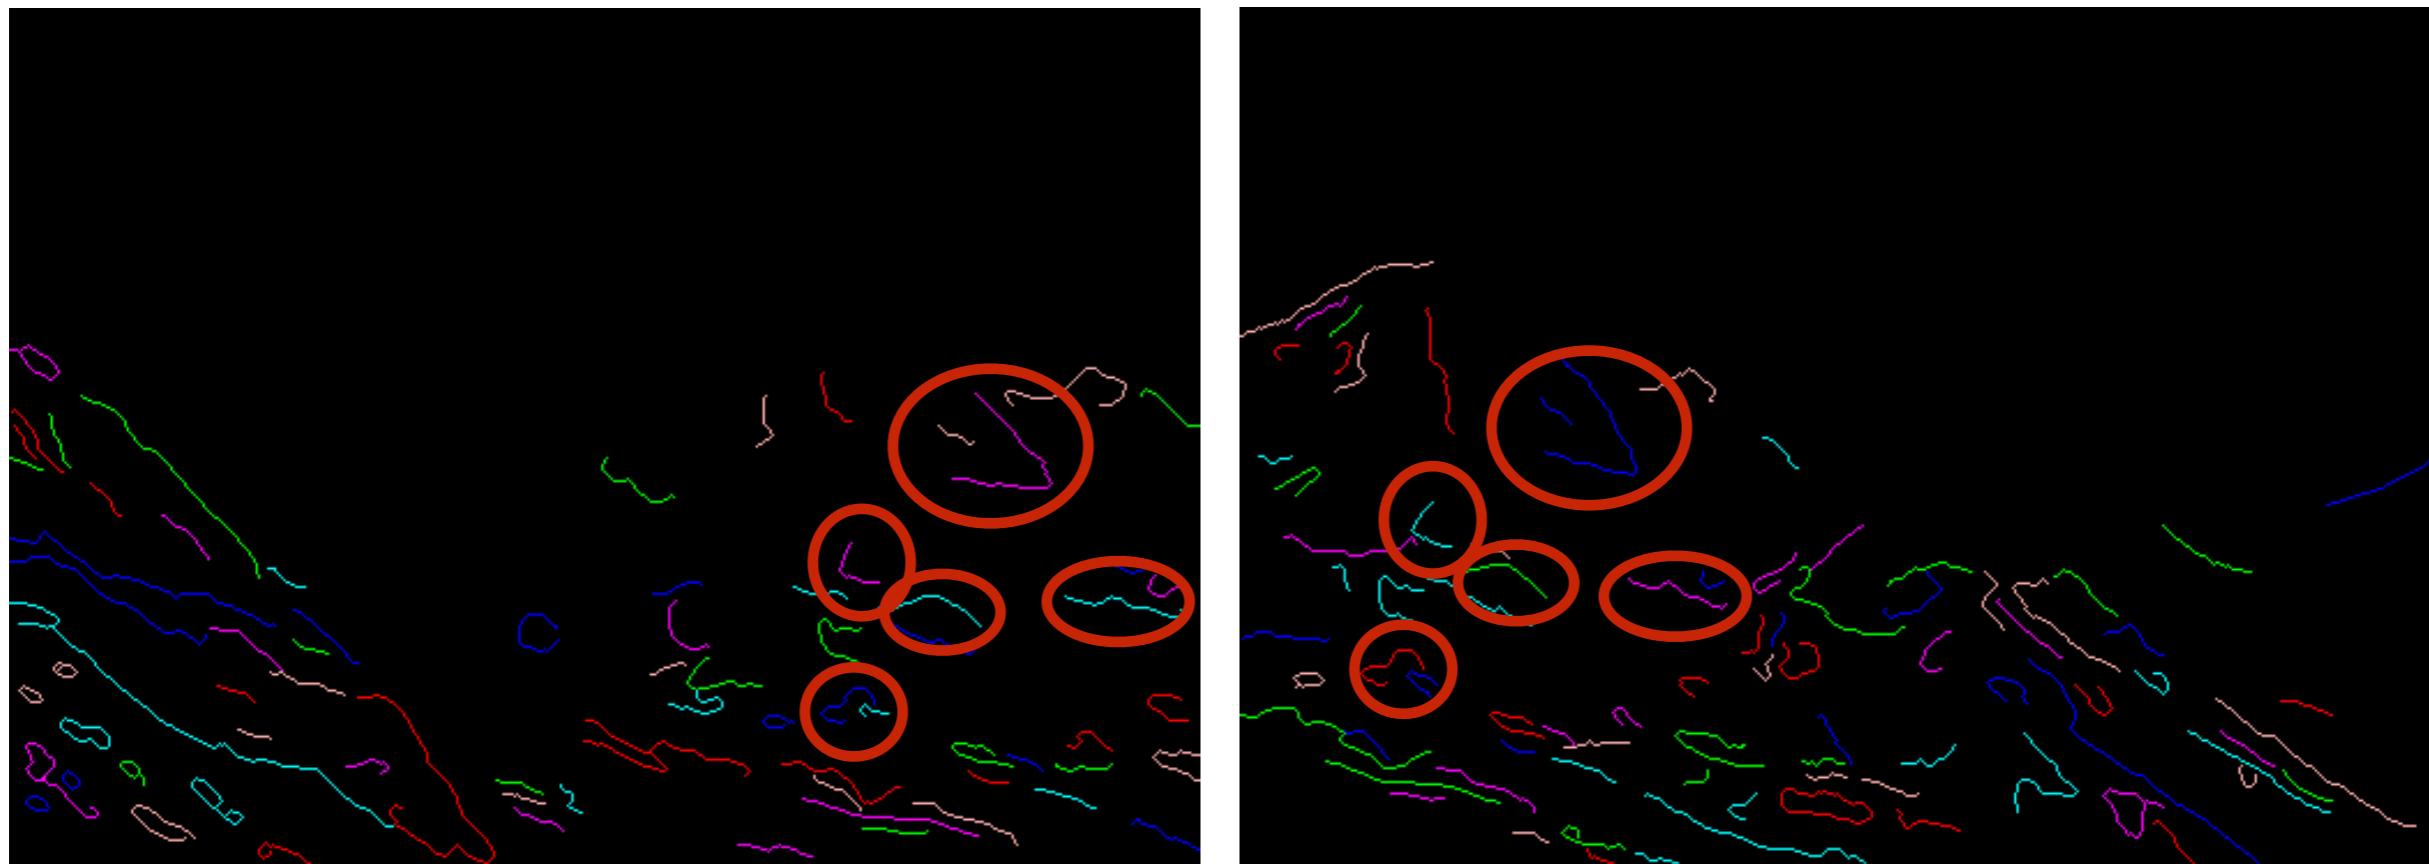
stereo projection fit to 32 points already known to match shows what the epipolar projections should be when the corner find + corner match + stereo projection solve are correctly automated.

**nMatched=32
avgDist=0.281
stDev=0.508**

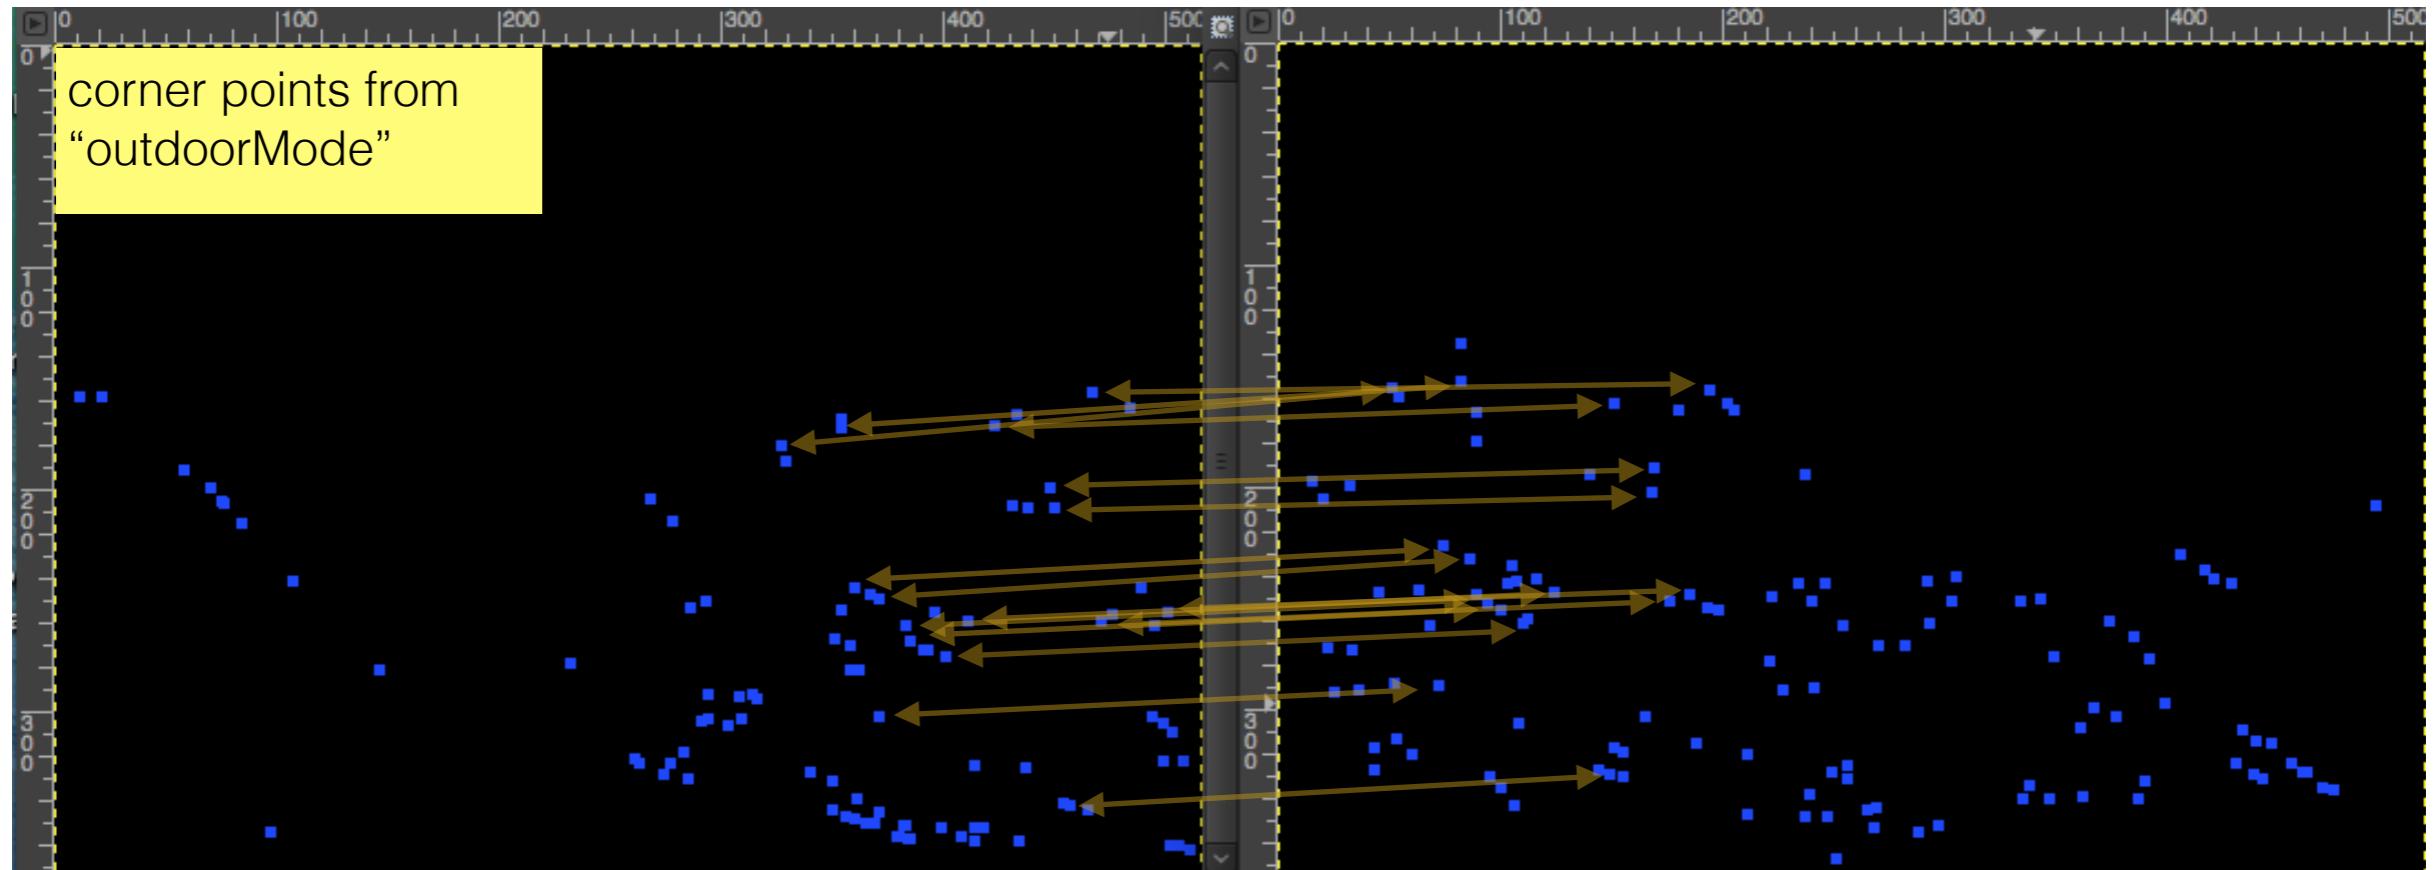
The Brown & Lowe 2003 images: point matching difficult because image intersection << difference



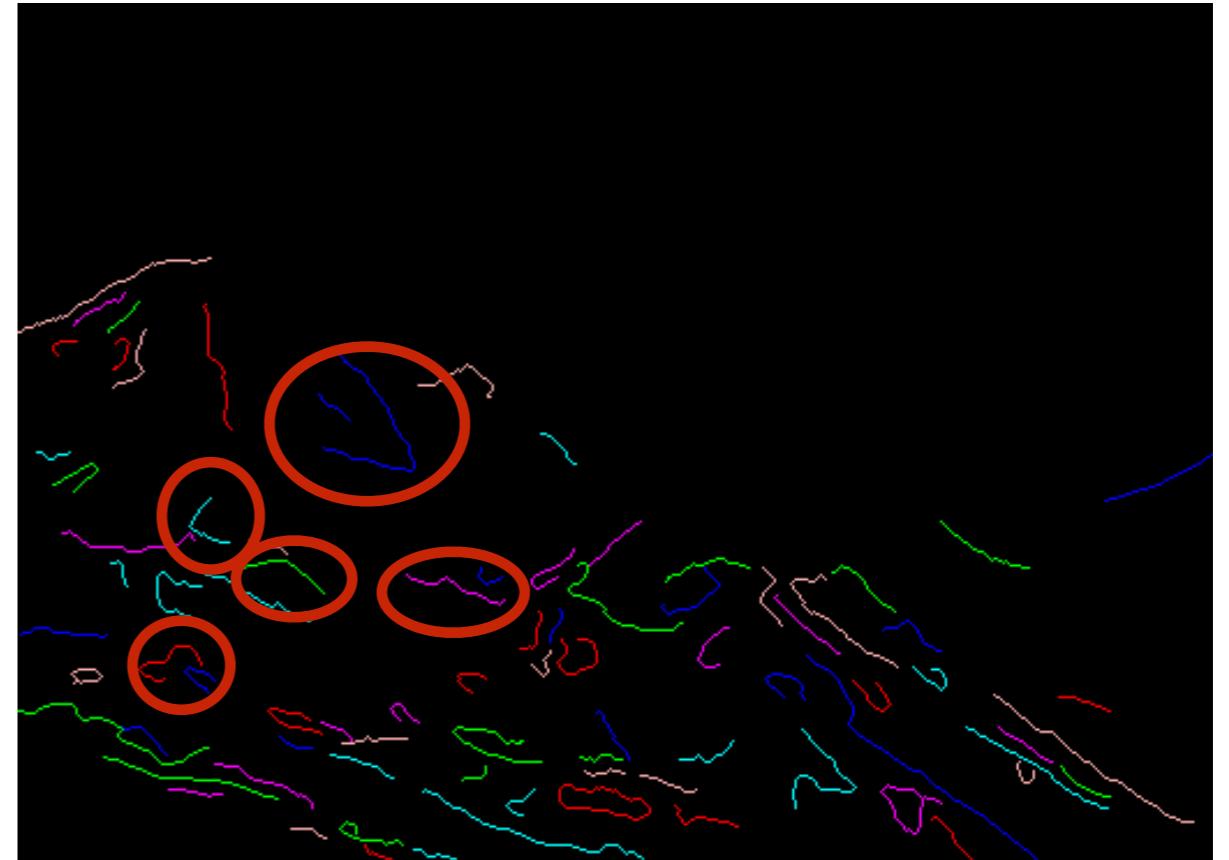
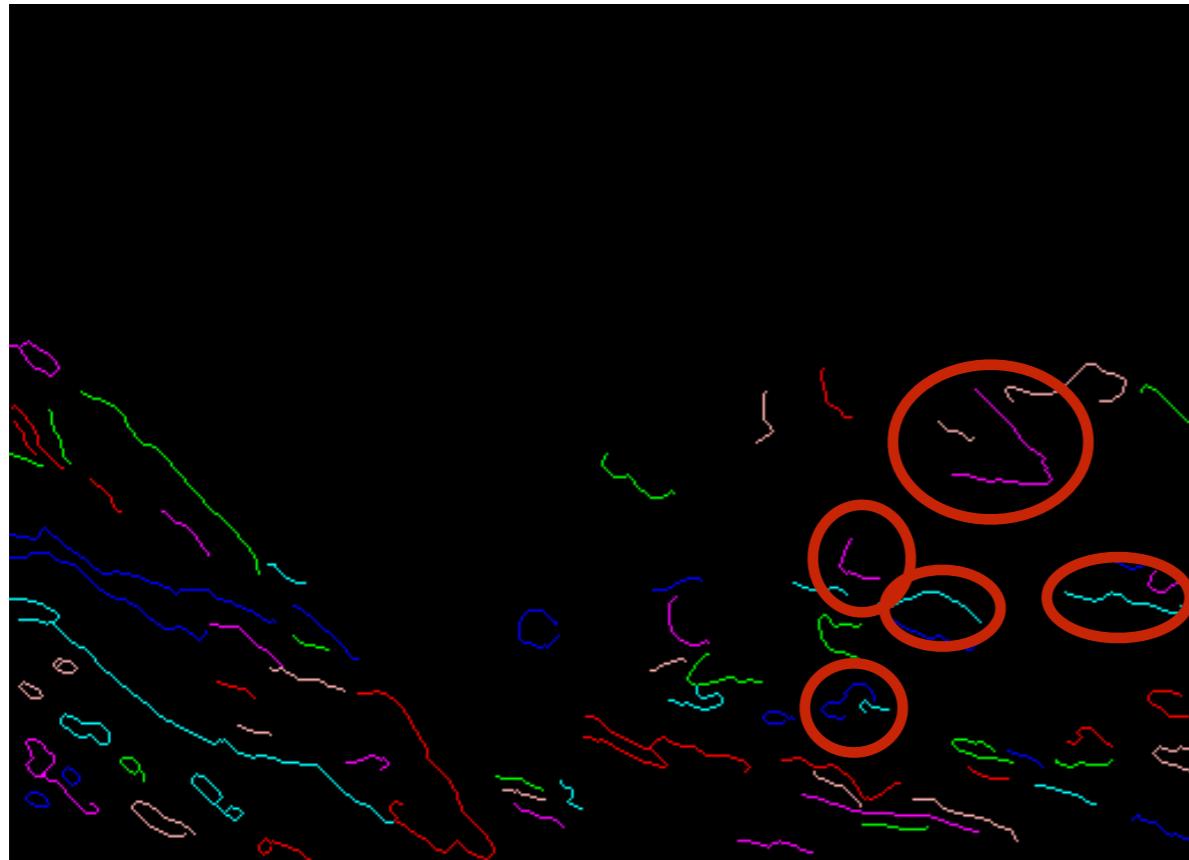
could consider distinct curves and their immediate neighbors, but that would be many more points:



The Brown & Lowe 2003 images: point matching difficult because image intersection < difference

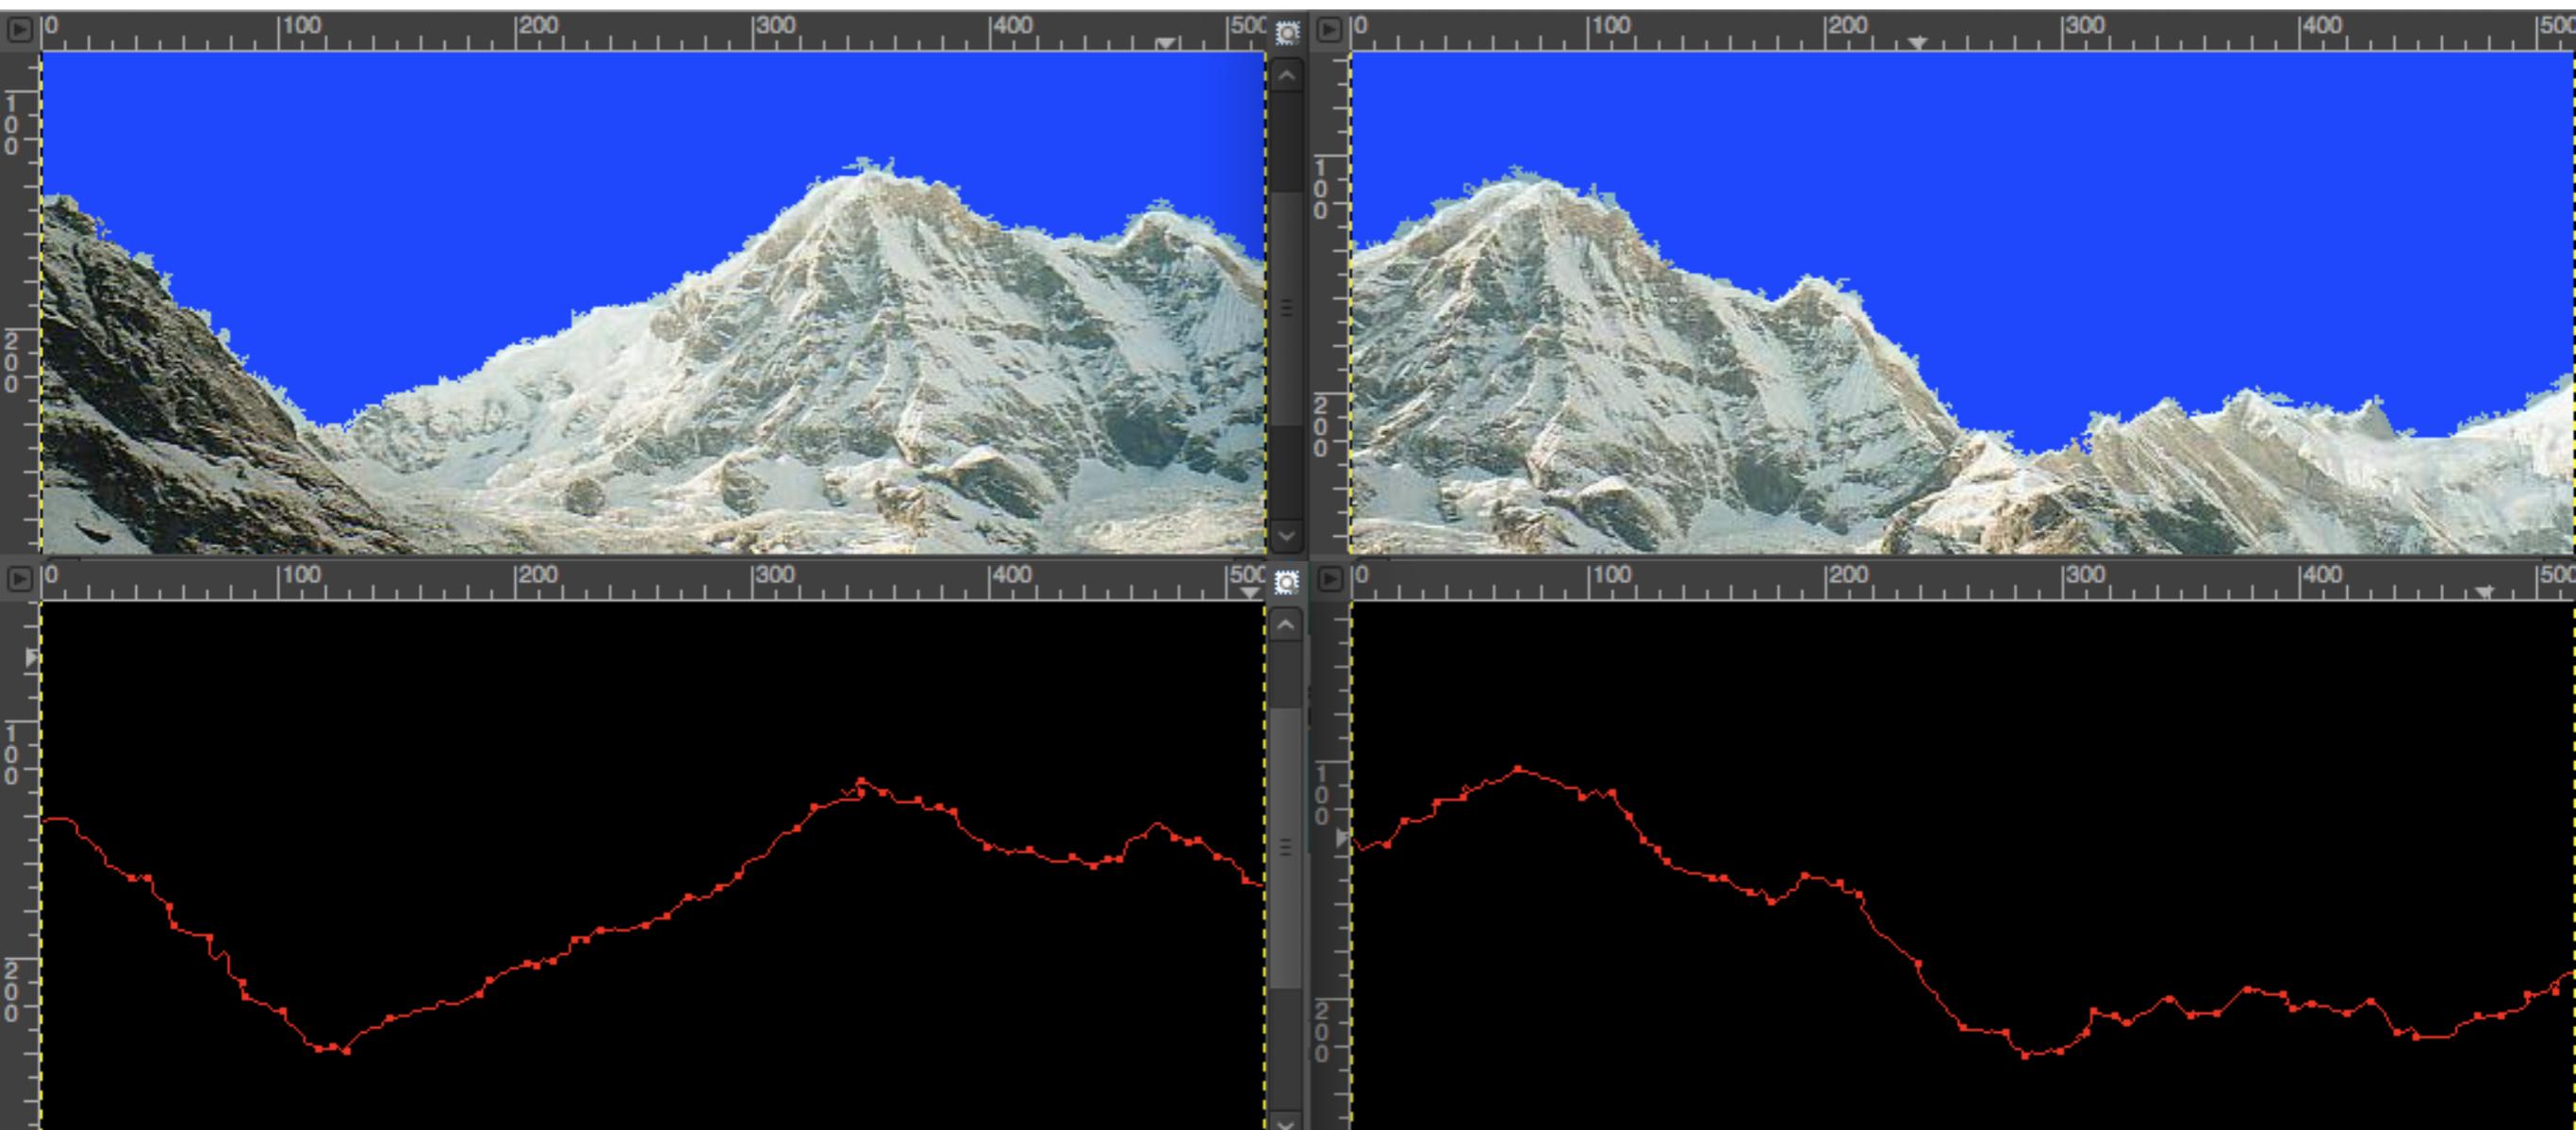


could consider distinct curves and their immediate neighbors, but that would be many more points:



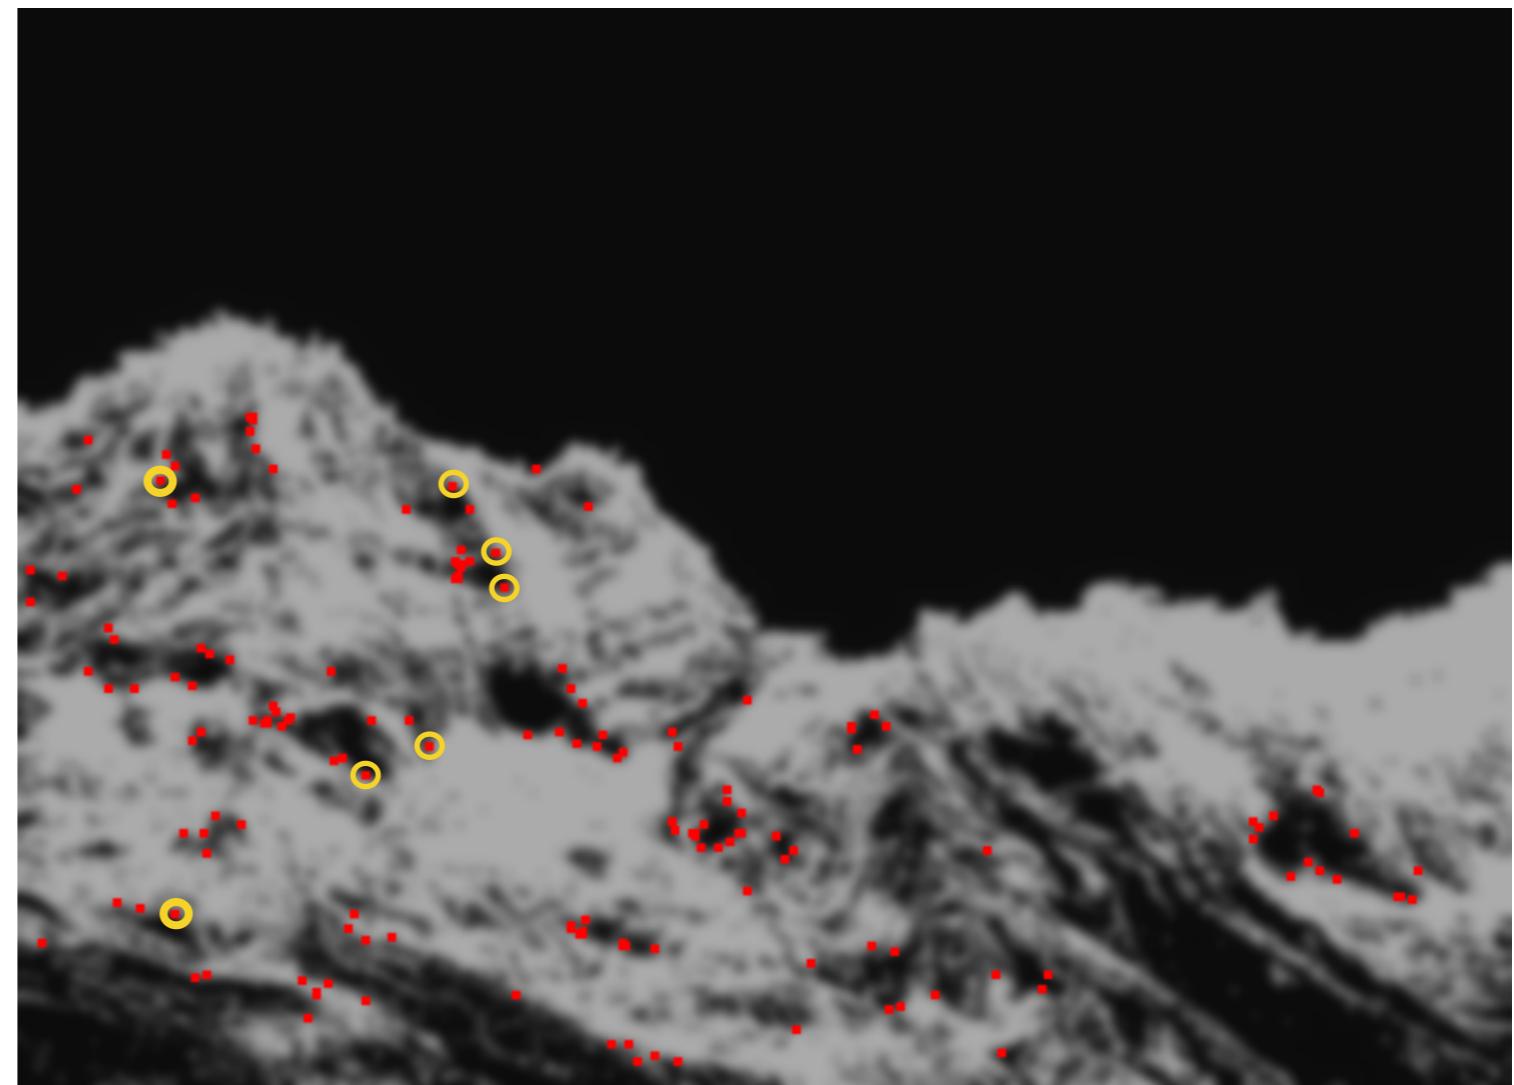
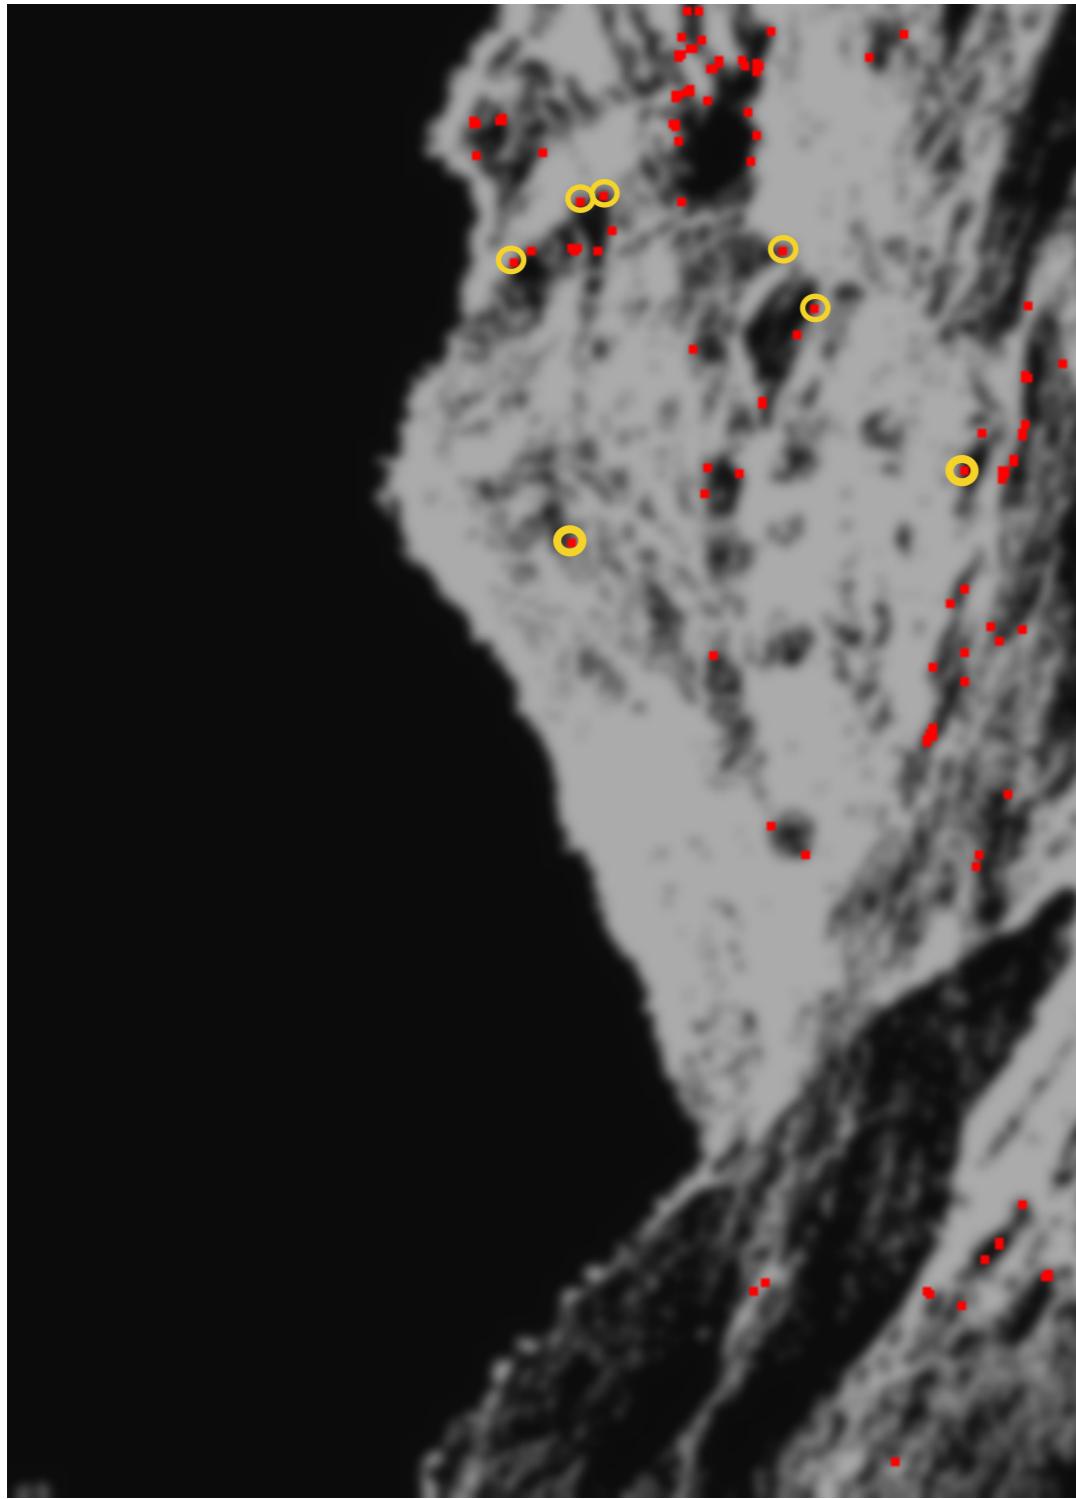
For the outdoor images, can find the sky and create a sky mask and also create corners from just the skyline.
(see skyline_extraction.pdf)

This sky mask helps to pre-process the image before feature finding and correspondence.

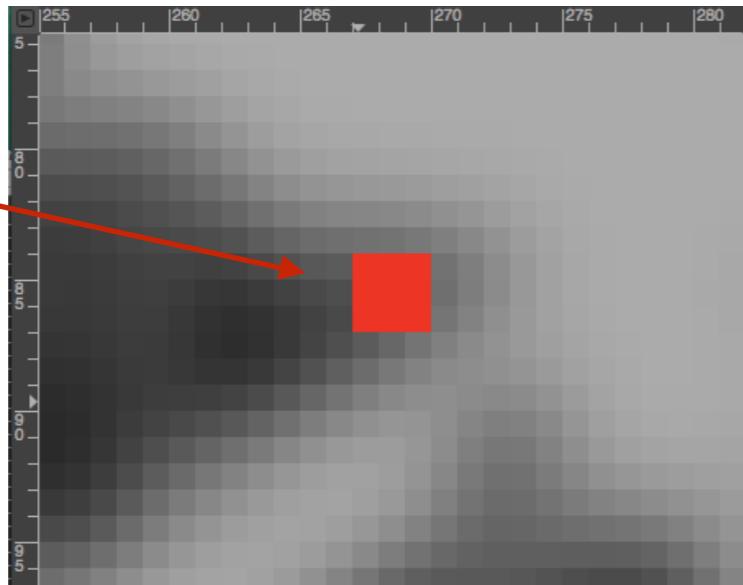
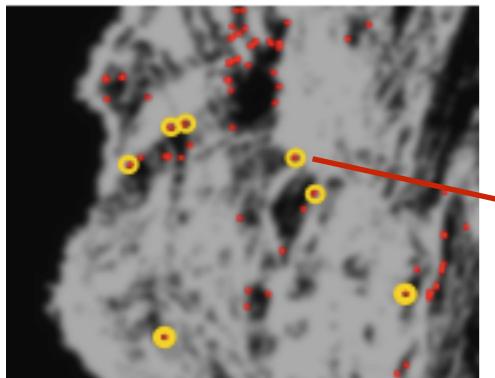


Need to match the corners using features because their locations alone are often not enough to match correctly, especially when the number of truly matchable is smaller than the number of corners or when features such as ridges exist which are false matches.

For features, need to determine an orientation of the corner region which is consistent w.r.t. the corner in any image. The highlighted corners are some examined in detail. The left is rotated to make sure the methods correctly handle rotation.



orientation: since the corners are made with CSS the edge information already exists and can be used here.



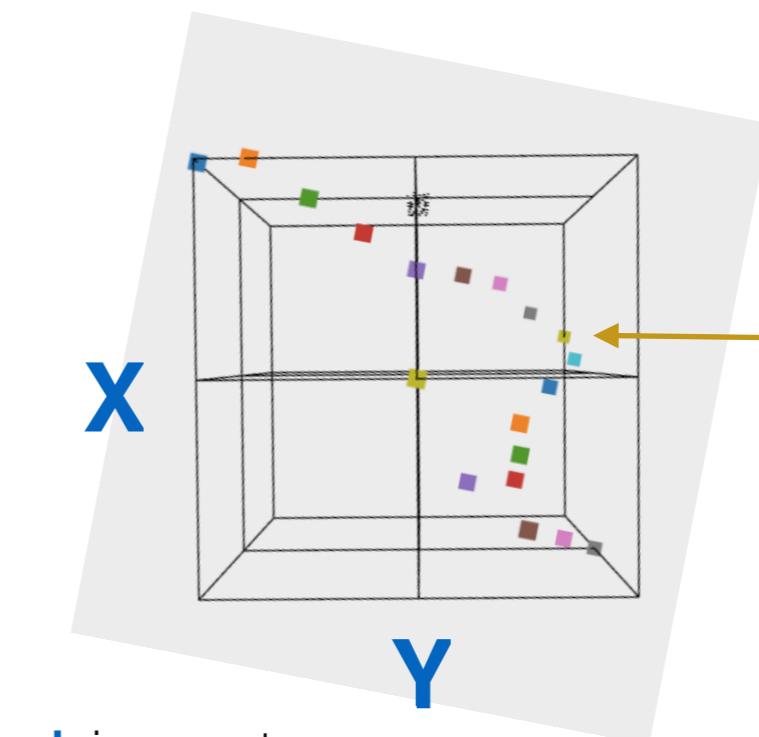
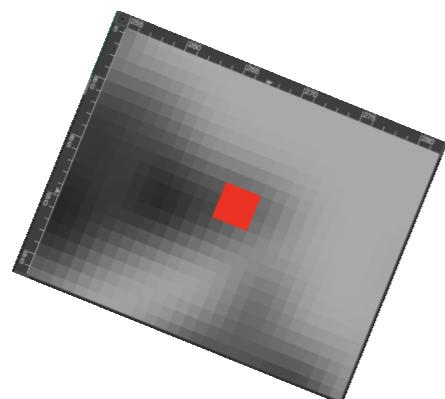
A vector perpendicular to k_{\max} on this edge can be used to calculate the orientation of this region (where the region is defined by being a CSS corner).

deriving the perpendicular angle from the points directly to the left and right of the maximum of curvature for points which have curvature

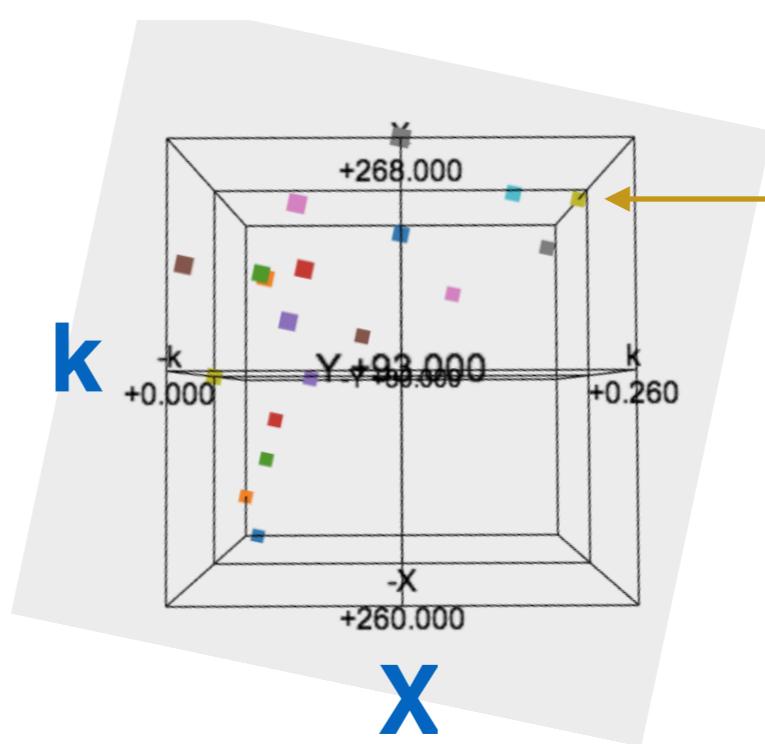
k	x	y
0.24, 267, 84	$dx, dy=(1,1)$,	$perp=45$
0.26, 268, 85	$<-max k$	
0.21, 268, 86	$dx, dy=(0,1)$,	$perp==0$

orientation = $(45 + 0)/2. = 22.5$

rotated



\mathbf{k} is curvature



maximum in curvature

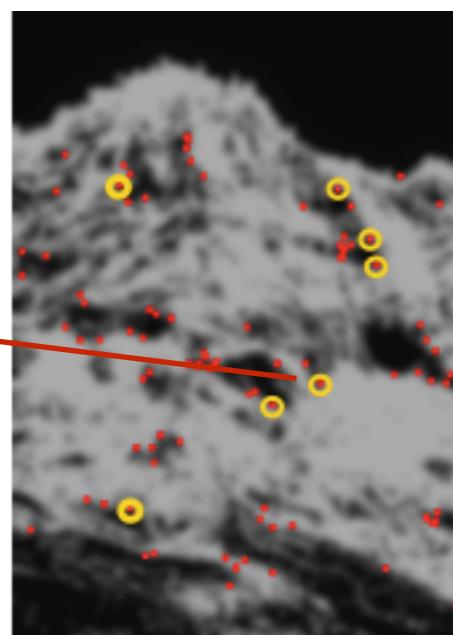
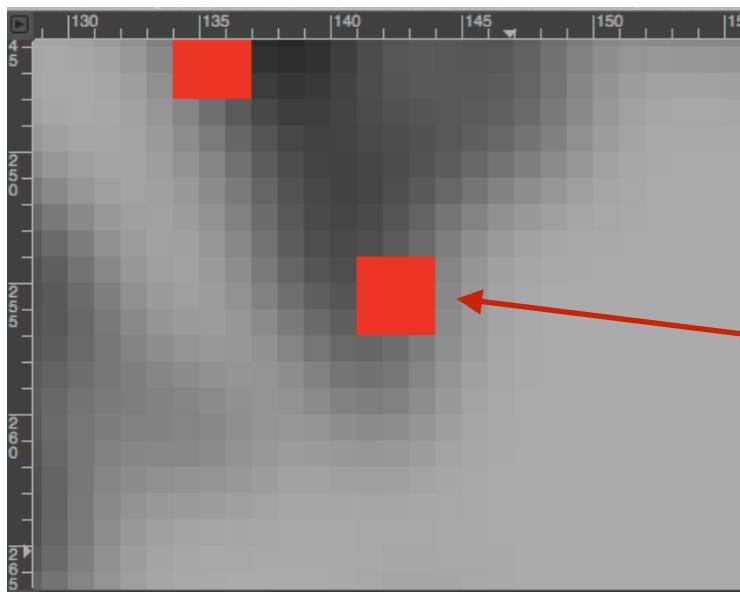
extract curvature from the highest resolution scale space curves:

```

idx=164 (268.0,85.0)
k      x      y
0.08   for (258, 79)
0.04   for (259, 79)
0.01   for (260, 80)
0.00   for (261, 80)
0.02   for (262, 81)
0.03   for (263, 82)
0.06   for (264, 83)
0.10   for (265, 83)
0.17   for (266, 83)
0.24   for (267, 84)
0.26   for (268, 85)
0.21   for (268, 86)
0.13   for (267, 87)
0.04   for (266, 88)
0.04   for (266, 89)
0.07   for (266, 90)
0.06   for (265, 90)
0.00   for (266, 91)
0.07   for (267, 92)
0.13   for (268, 93)
0.15   for (268, 94)

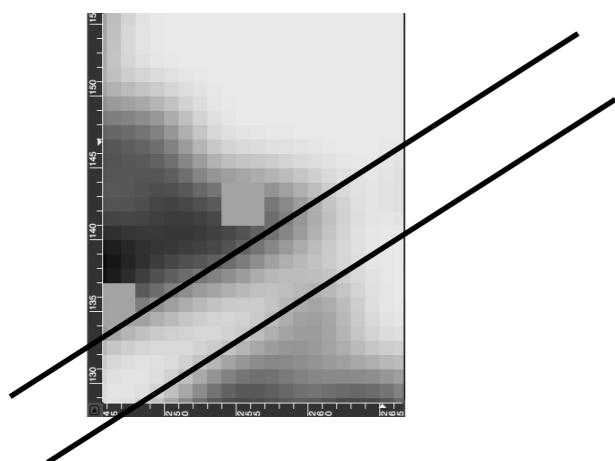
```

orientation.



extract the final curvature from the highest resolution
scale space curves:
idx=112 (143.0, 255.0)

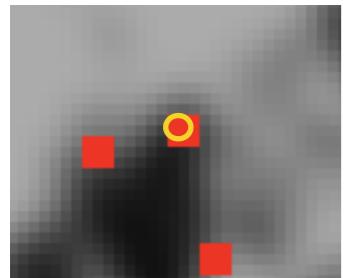
```
k[110]=0.01  for (143, 253)  
k[111]=0.13  for (143, 254)  
k[112]=0.34  for (143, 255)  dx,dy=(-1,0)  
k[113]=0.43  for (142, 255) <-- k_max  
k[114]=0.35  for (141, 255)  dx,dy=(-1,0)  
k[115]=0.28  for (140, 255)  
k[116]=0.24  for (139, 255)
```



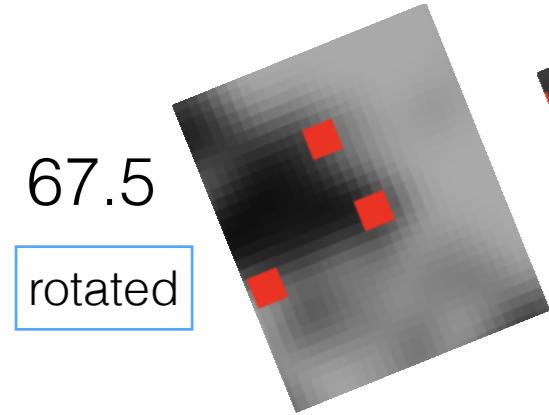
$$\text{orientation} = ((-90)+(-90))/2. = -90$$

rotated

orientation. corners from Brown & Lowe 2003 left
and right panorama images



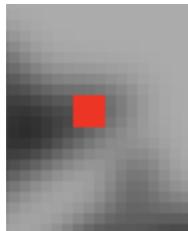
k[0]=0.25 x,y=(204, 66)
k[1]=0.32 x,y=(205, 66)
k[2]=0.32 x,y=(206, 66)
k[3]=0.30 x,y=(207, 67)
k[4]=0.26 x,y=(208, 68)



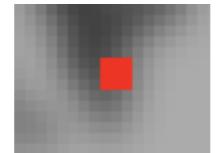
67.5

rotated

k[0]=0.17 x,y=(266, 83)
k[1]=0.24 x,y=(267, 84)
k[2]=0.26 x,y=(268, 85)
k[3]=0.21 x,y=(268, 86)
k[4]=0.13 x,y=(267, 87)

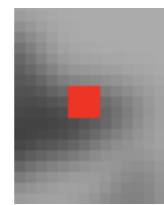
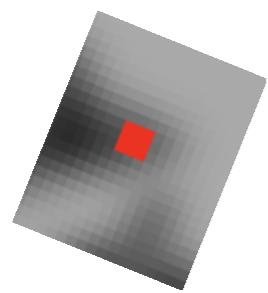


k[0]=0.13 x,y=(143, 254)
k[1]=0.34 x,y=(143, 255)
k[2]=0.43 x,y=(142, 255)
k[3]=0.35 x,y=(141, 255)
k[4]=0.28 x,y=(140, 255)



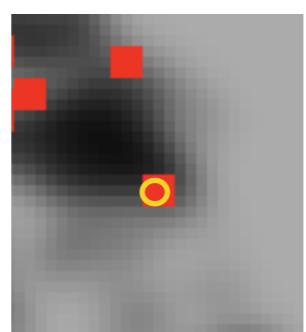
22.5

rotated

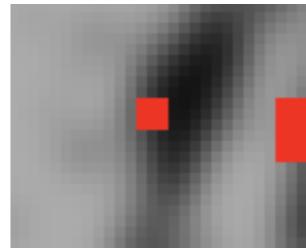


270.0

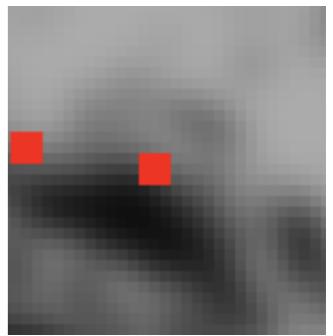
rotated



k[0]=0.22 x,y=(169, 197)
k[1]=0.32 x,y=(169, 198)
k[2]=0.41 x,y=(169, 199)
k[3]=0.38 x,y=(168, 200)
k[4]=0.27 x,y=(167, 200)



k[0]=0.11 x,y=(332, 161)
k[1]=0.37 x,y=(331, 161)
k[2]=0.58 x,y=(330, 161)
k[3]=0.47 x,y=(330, 162)
k[4]=0.30 x,y=(330, 163)



k[0]=0.07 x,y=(56, 315)
k[1]=0.09 x,y=(55, 314)
k[2]=0.31 x,y=(54, 313)
k[3]=0.23 x,y=(53, 313)
k[4]=0.03 x,y=(52, 313)

337.5

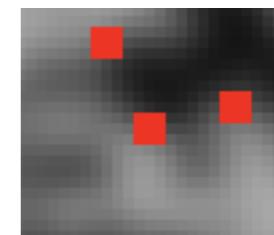
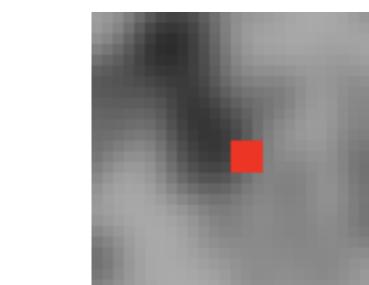
rotated

135.0

rotated

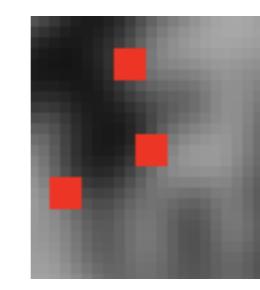
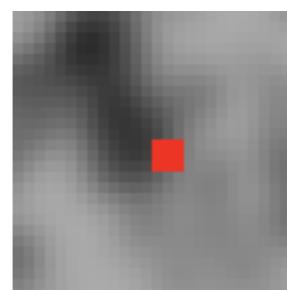
67.5

rotated



k[0]=0.16 x,y=(195, 183)
k[1]=0.26 x,y=(195, 184)
k[2]=0.32 x,y=(195, 185)
k[3]=0.29 x,y=(195, 186)
k[4]=0.25 x,y=(194, 187)

k[0]=0.20 x,y=(52, 170)
k[1]=0.23 x,y=(53, 171)
k[2]=0.25 x,y=(54, 171)
k[3]=0.21 x,y=(55, 171)
k[4]=0.13 x,y=(56, 171)



0.0

rotated

270.0

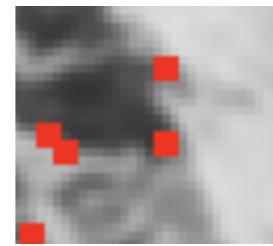
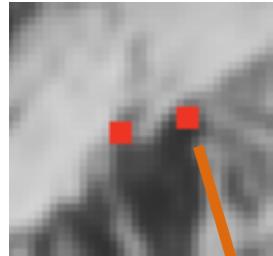
rotated

orientation and junctions. corners
from Brown & Lowe 2003 left and right panorama images

k[0]=0.09 x,y=(204, 65)
k[1]=0.19 x,y=(205, 65)
k[2]=0.21 x,y=(206, 65)
k[3]=0.09 x,y=(206, 64)
k[4]=0.03 x,y=(206, 63)

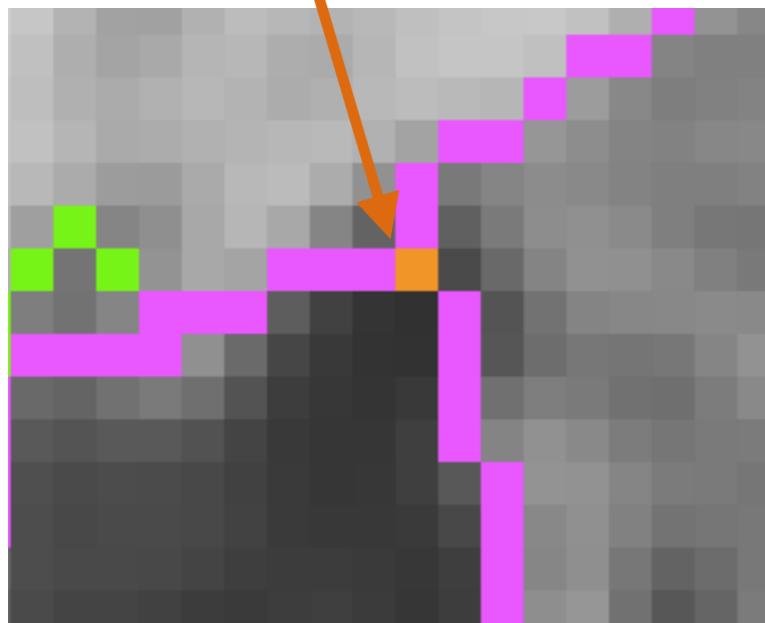
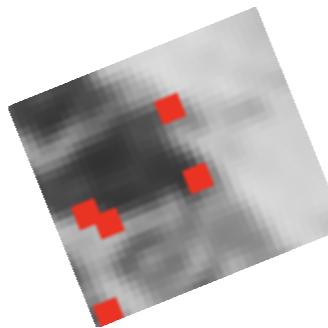
k[0]=0.21 x,y=(170, 198)
k[1]=0.32 x,y=(170, 199)
k[2]=0.41 x,y=(170, 200)
k[3]=0.38 x,y=(169, 201)
k[4]=0.27 x,y=(168, 201)

315.0



337.5

rotated

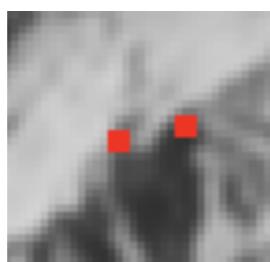


The left image corner is in a *junction* and that requires more complex analysis to determine here that the better edge for it is the lower right edge instead of the upper edge (the later gives an inconsistent orientation w.r.t. the right image).

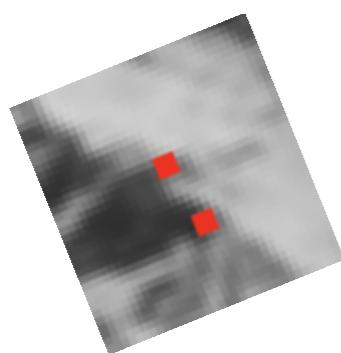
Without the more complex analysis for now, will just create a CornerRegion for each combination and one of the 3 will match the right.

orientation and junctions (continued)

k[0]=0.02 x,y=(205, 65)
k[1]=0.21 x,y=(206, 65)
k[2]=0.02 x,y=(207, 66)



k[0]=0.21 x,y=(170, 198)
k[1]=0.32 x,y=(170, 199)
k[2]=0.41 x,y=(170, 200)
k[3]=0.38 x,y=(169, 201)
k[4]=0.27 x,y=(168, 201)



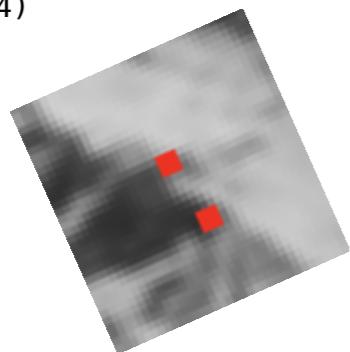
67.5

rotated

337.5

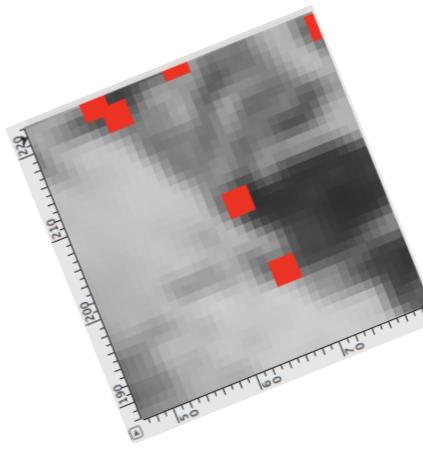
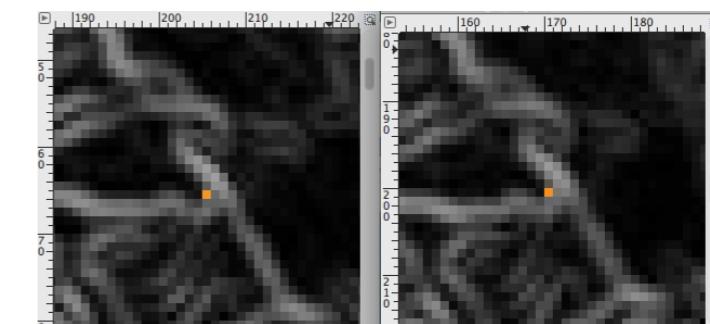
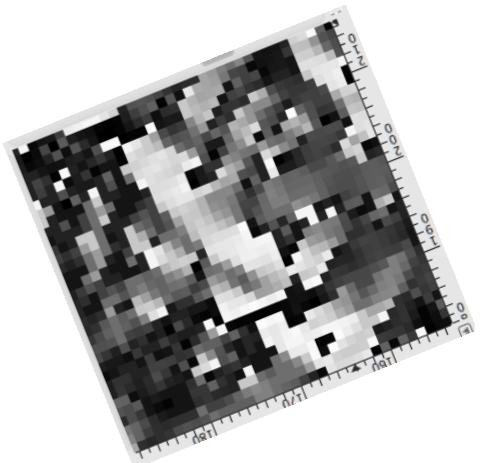
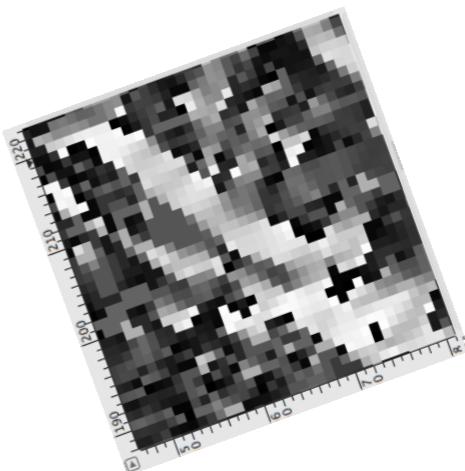
rotated

k[0]=0.02 x,y=(205, 65)
k[1]=0.21 x,y=(206, 65)
k[2]=0.02 x,y=(206, 64)

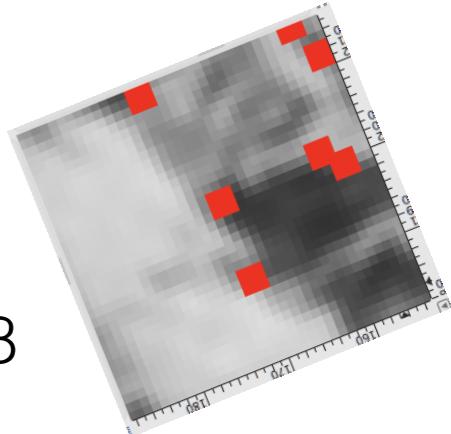


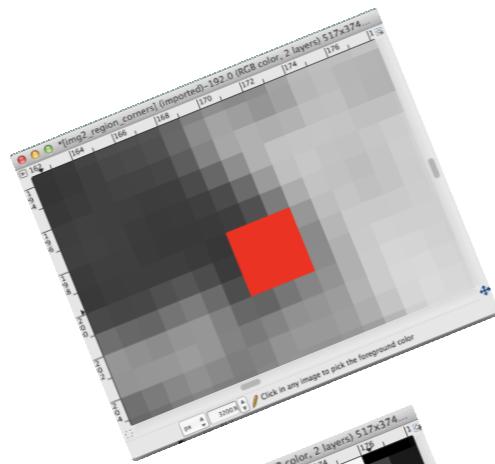
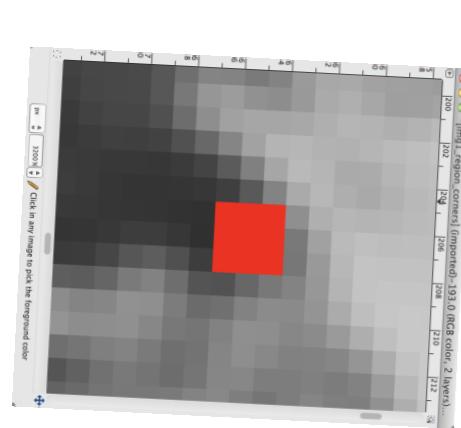
66

rotated

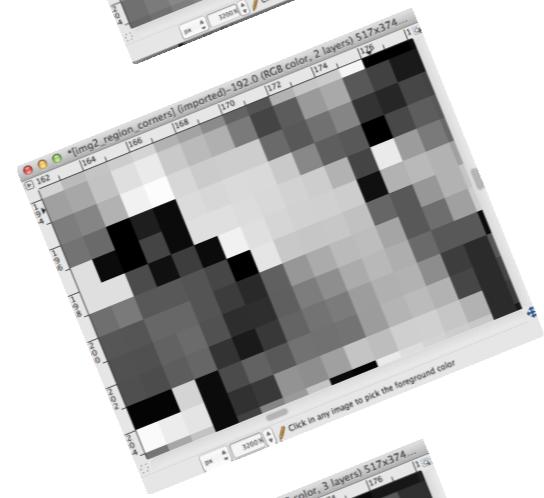
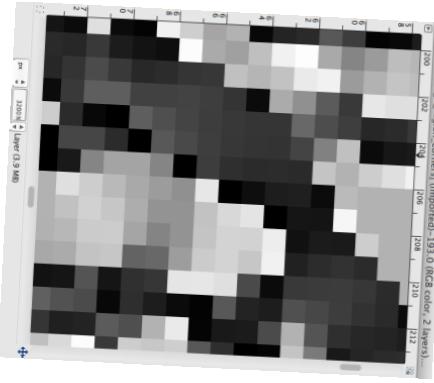


250 158

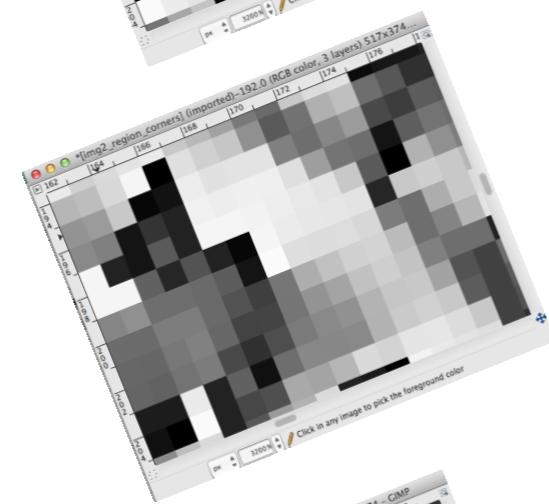
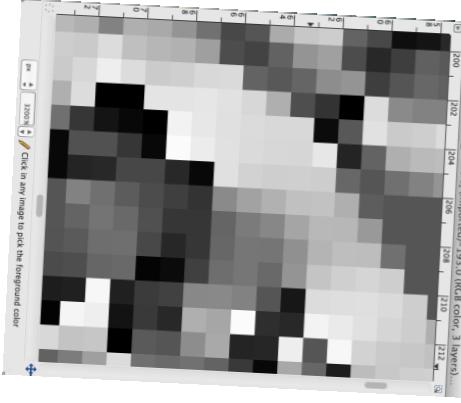




greyscale rotated to “dominant orientation”

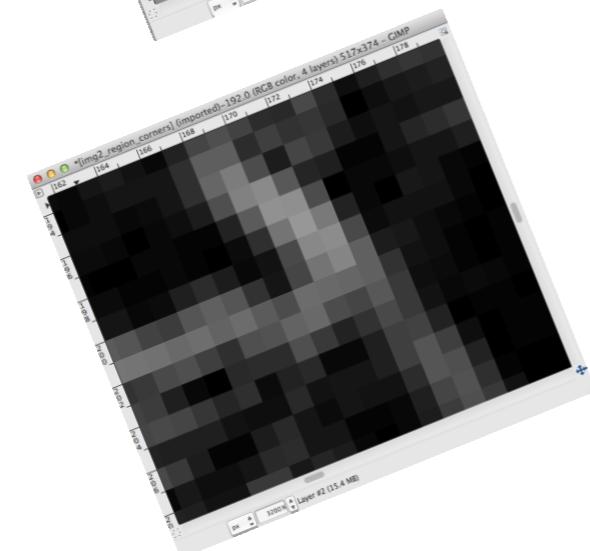
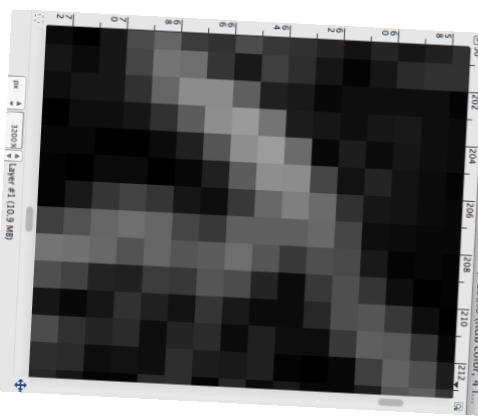


gradient theta rotated to “dominant orientation”



gradient theta rotated to “dominant orientation” with “dominant orientation” subtracted from theta

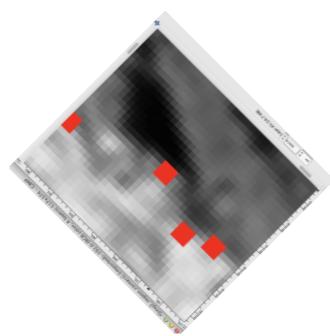
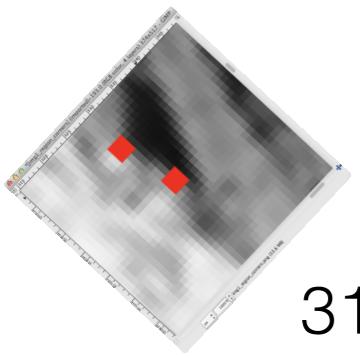
can see that image 2 should have preferred match w/ orientation 67.5 instead of 93



can see that SSD on gradient theta could be a good descriptor with the orientation corrections and quadrant consideration (e.g. 0 diff 350 is 10, not 350)

can see that SSD on gradient could be a good descriptor but may not be as tolerant of skew as subcells of histograms of orientation?

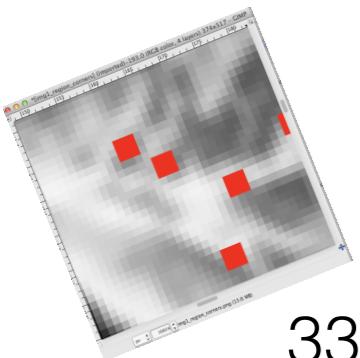
will use gradient, but binned into 2x2 cells, and 16 of them surrounding the corner.



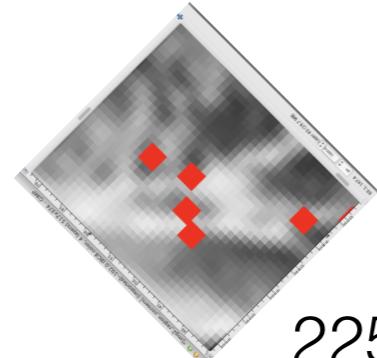
[1]

313

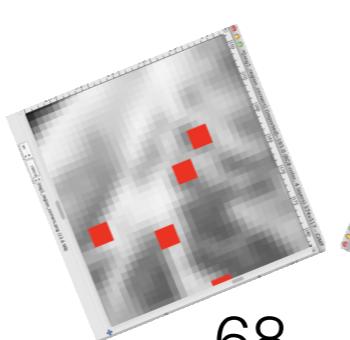
225



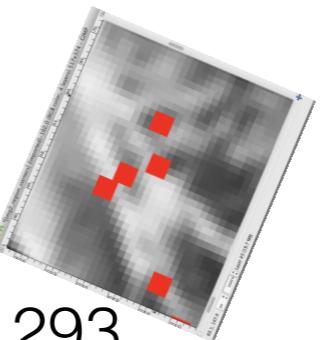
338



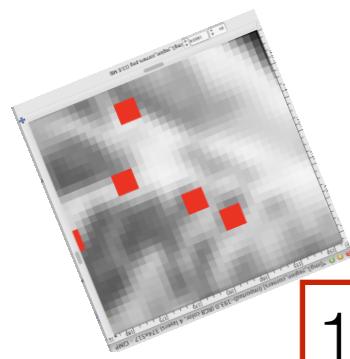
[2]



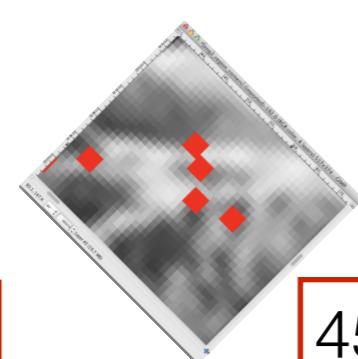
68



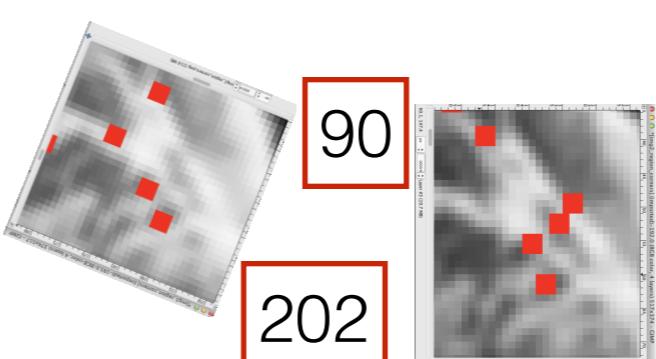
293



158



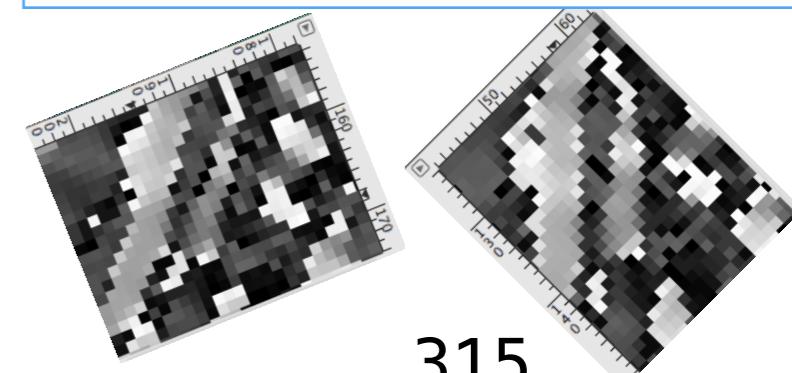
45



90

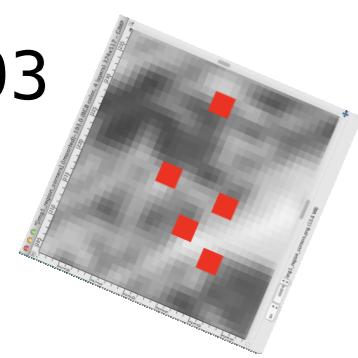
202

gradient theta rotated to “dominant orientation” with “dominant orientation” subtracted from theta (ref for method?)

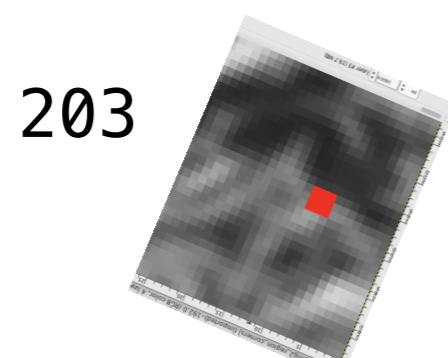


68

315



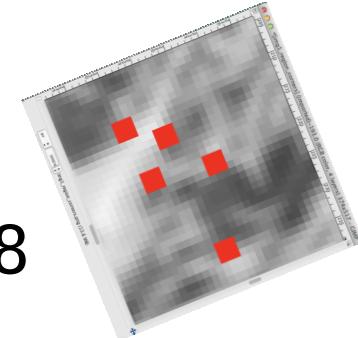
293



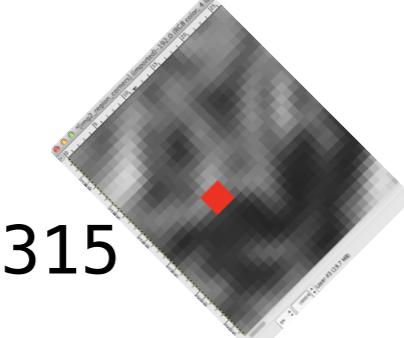
203

[3]

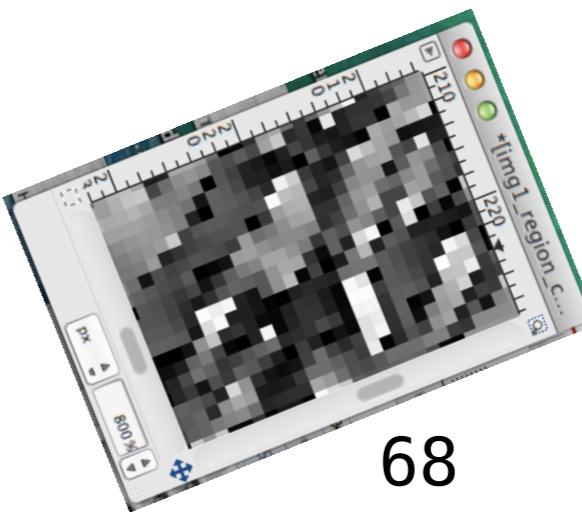
good for
normalization
test or preference
for gradient theta
(220, 220)(9, 194)



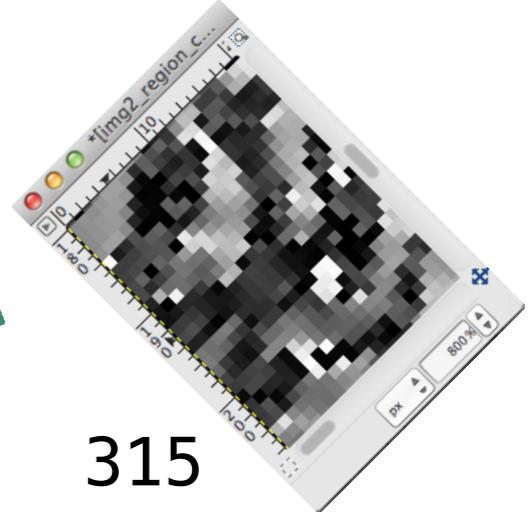
68



315

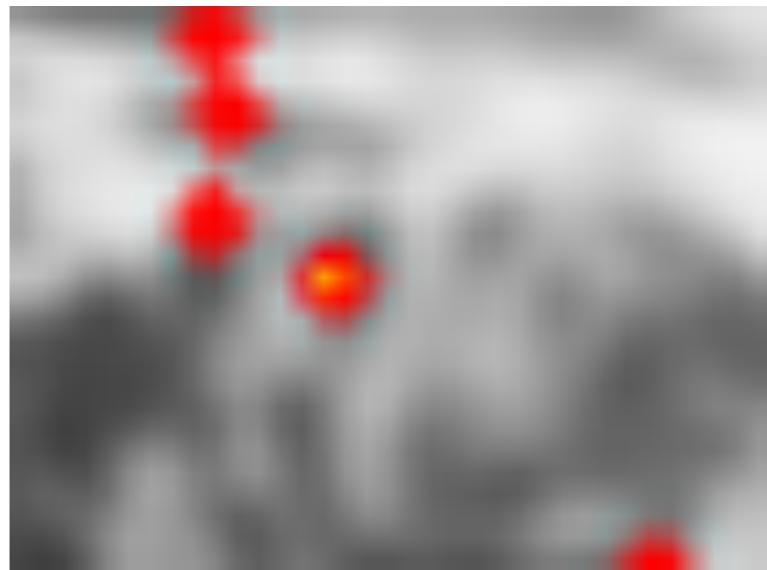


68



315

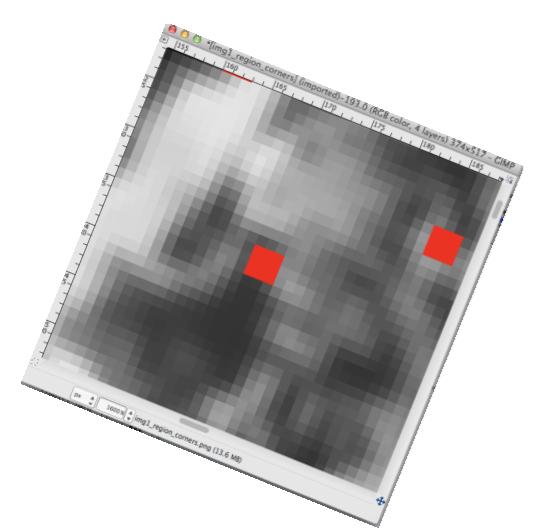
gradient theta rotated to “dominant orientation” with “dominant orientation” subtracted from theta (ref for method?)



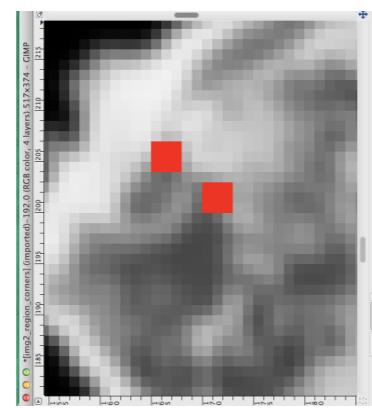
left and right clips from the Brown & Lowe showing projection and illumination differences.

Binning the intensity descriptor pixels into 2x2 cells and subtracting the mean of the descriptor intensity values from itself helps to remove the illumination differences and projection differences for those within that small range.

Note that the rectangular radius around the central pixel is somewhat large, 12 pixels (note too that the snapshot above is larger than a descriptor's region).

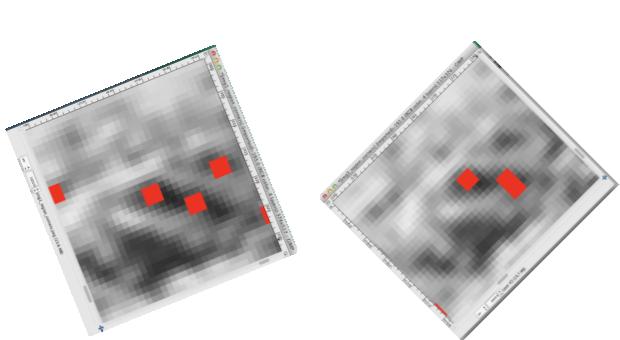


0 or 22.5

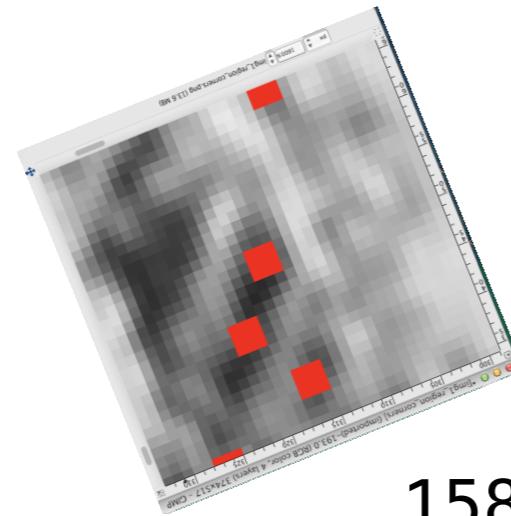


270

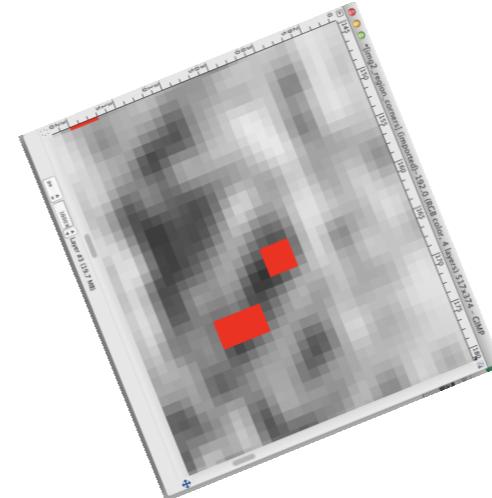
[4]



68

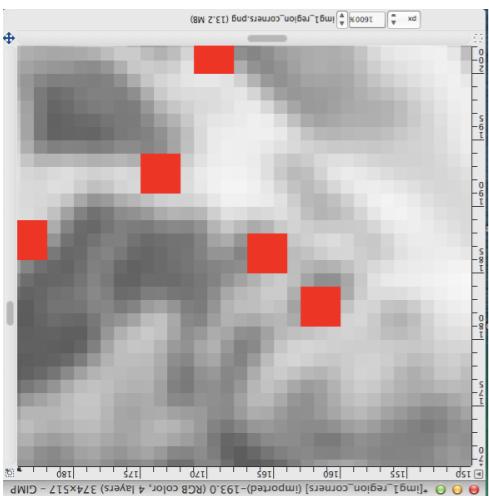


315

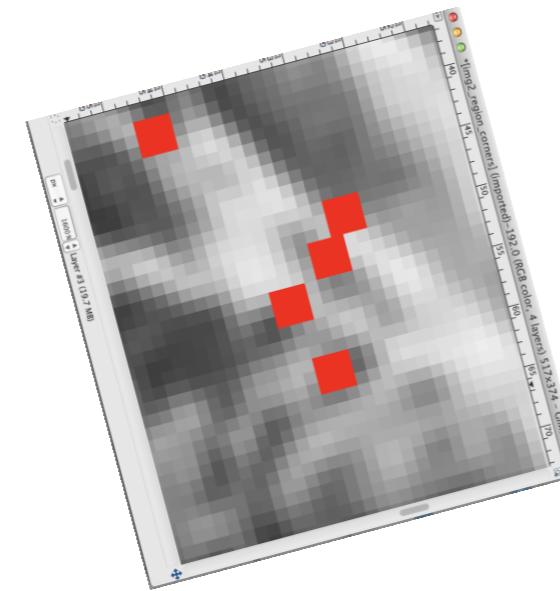


158

[5]



180

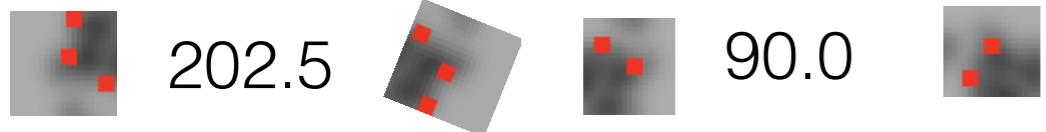


75?

[6]

orientation. corners from Venturi dataset

$k[0]=0.15$ x,y=(461, 449)	$k[0]=0.13$ x,y=(157, 466)
$k[1]=0.23$ x,y=(460, 448)	$k[1]=0.19$ x,y=(156, 466)
$k[2]=0.29$ x,y=(459, 447)	$k[2]=0.22$ x,y=(155, 466)
$k[3]=0.26$ x,y=(459, 446)	$k[3]=0.20$ x,y=(154, 466)
$k[4]=0.16$ x,y=(459, 445)	$k[4]=0.16$ x,y=(153, 467)



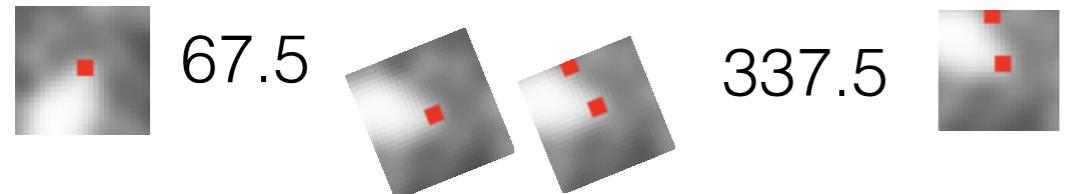
rotated

$k[0]=0.18$ x,y=(293, 496)	$k[0]=0.18$ x,y=(105, 298)
$k[1]=0.24$ x,y=(292, 495)	$k[1]=0.33$ x,y=(105, 297)
$k[2]=0.29$ x,y=(291, 494)	$k[2]=0.36$ x,y=(106, 296)
$k[3]=0.25$ x,y=(291, 493)	$k[3]=0.26$ x,y=(107, 296)
$k[4]=0.13$ x,y=(291, 492)	$k[4]=0.14$ x,y=(108, 296)

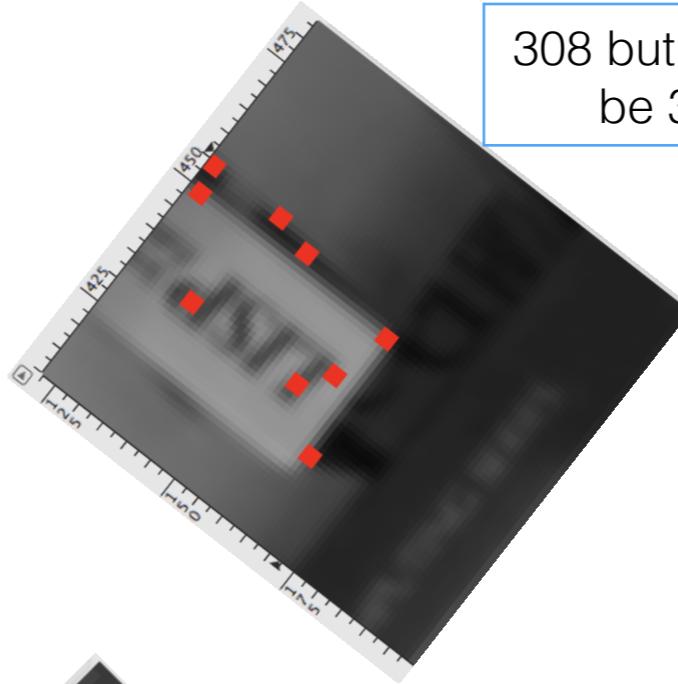
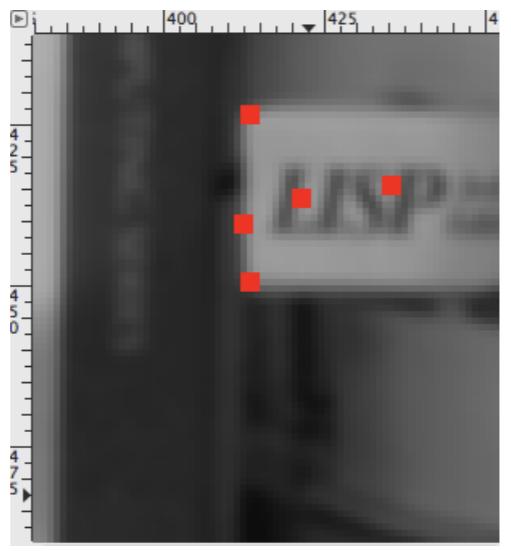
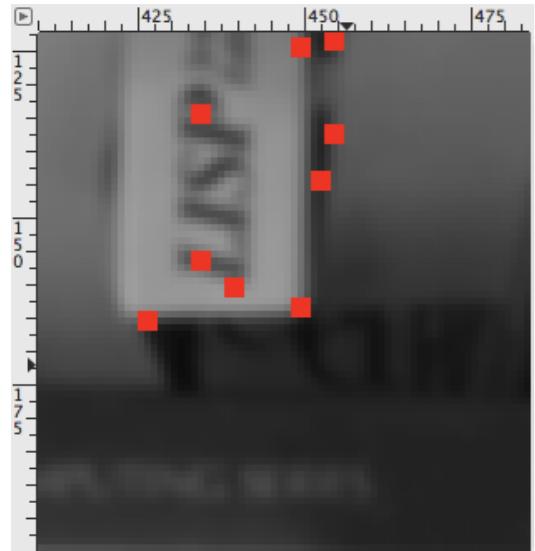


rotated

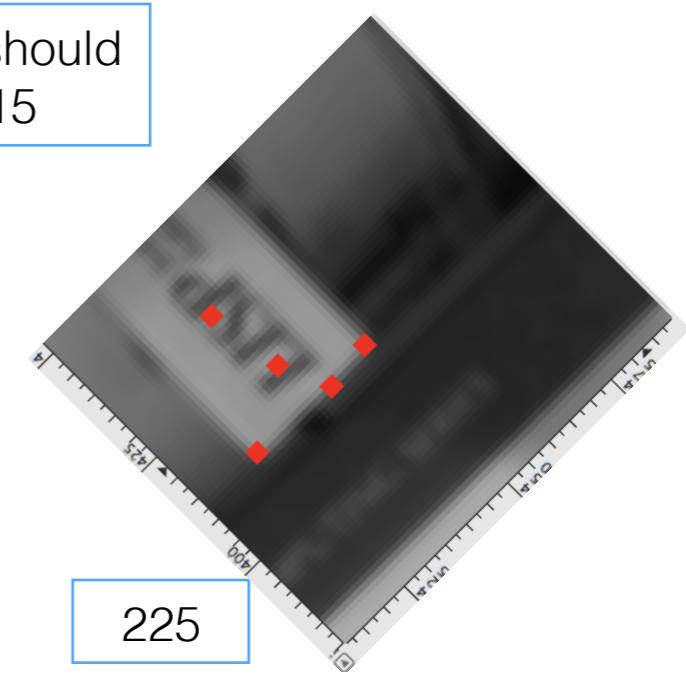
$k[0]=0.27$ x,y=(449, 222)	$k[0]=0.26$ x,y=(382, 448)
$k[1]=0.33$ x,y=(448, 221)	$k[1]=0.29$ x,y=(383, 447)
$k[2]=0.37$ x,y=(447, 220)	$k[2]=0.31$ x,y=(384, 446)
$k[3]=0.35$ x,y=(446, 220)	$k[3]=0.29$ x,y=(384, 445)
$k[4]=0.26$ x,y=(445, 220)	$k[4]=0.23$ x,y=(384, 444)



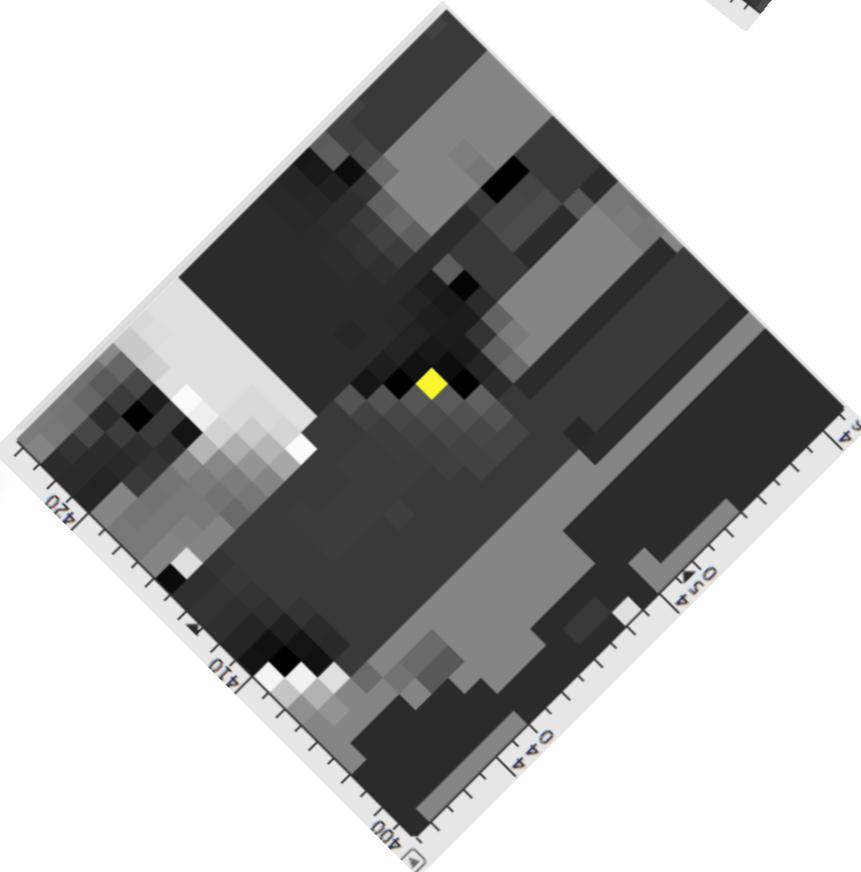
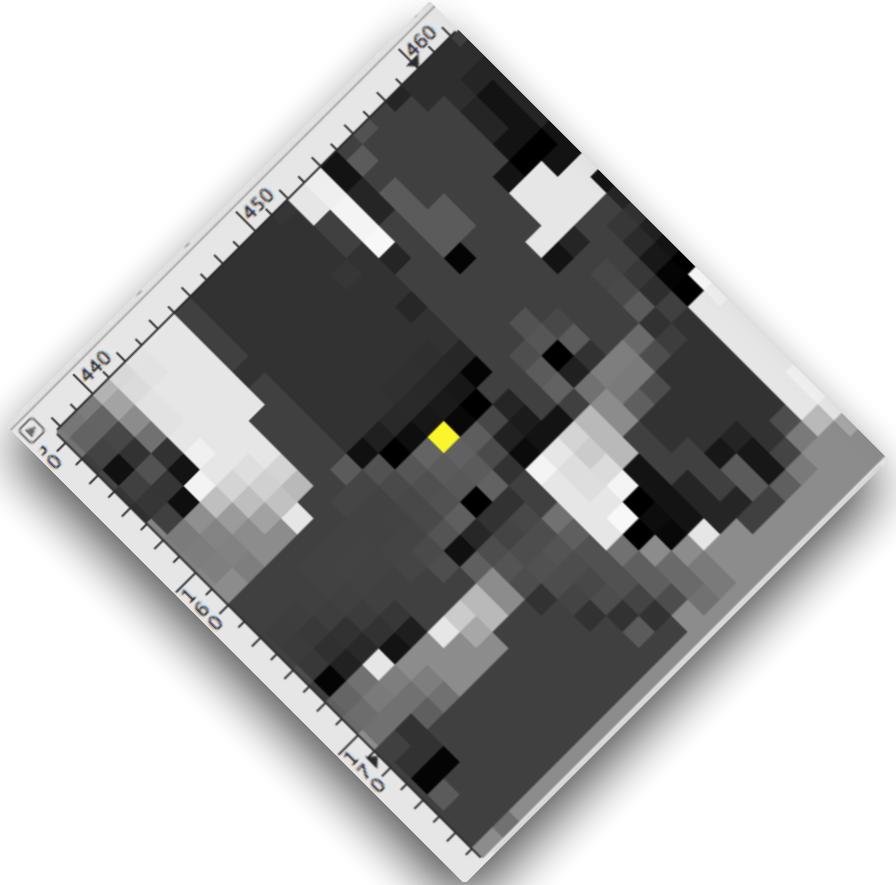
rotated



308 but should
be 315

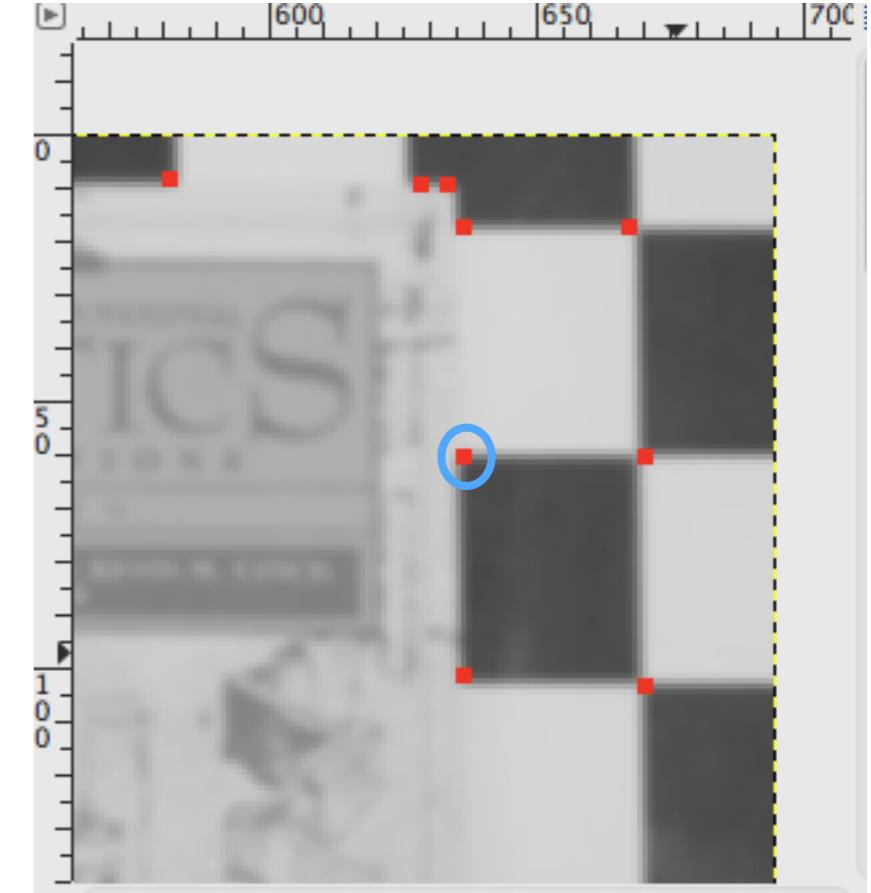
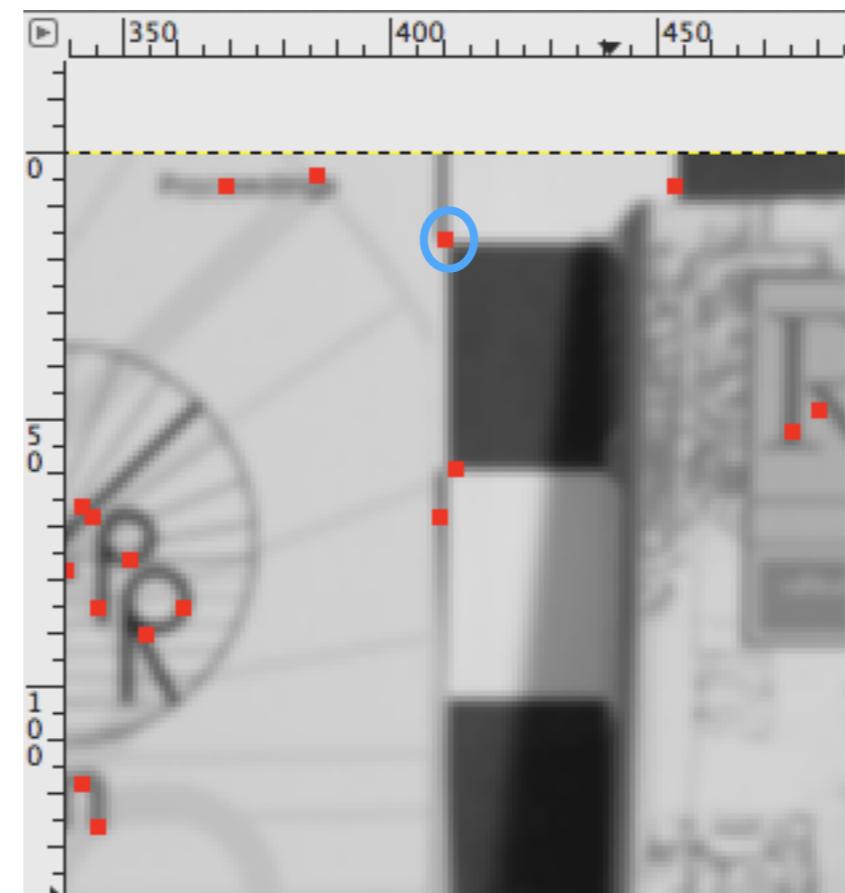
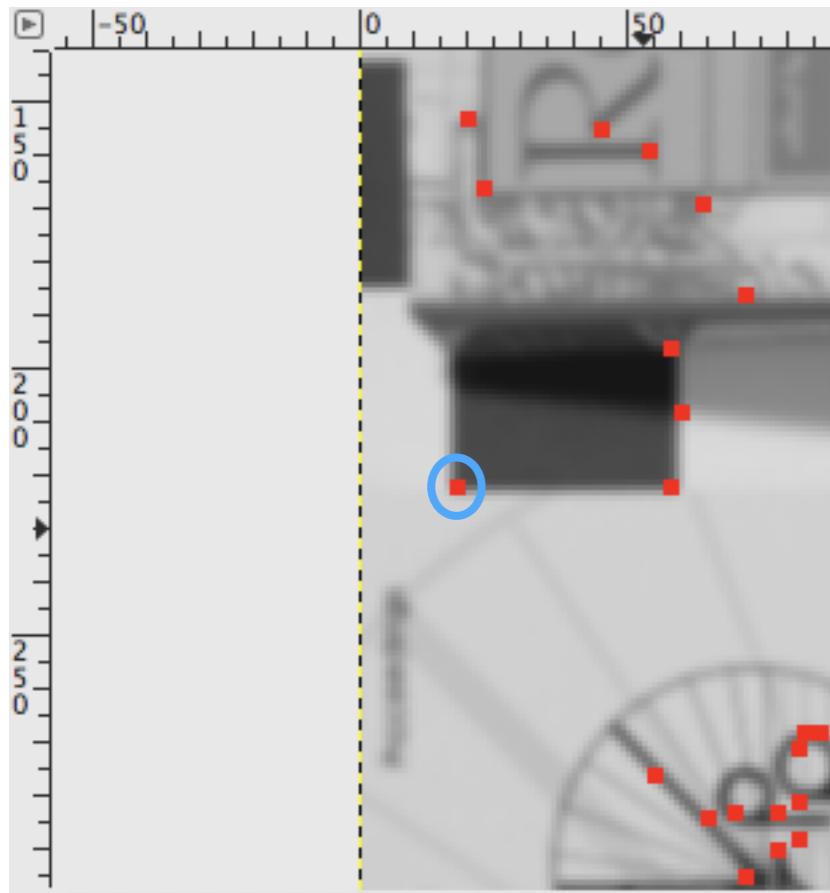


225



the theta descriptor SSD is larger than the error from auto-correlation.
discard in favor of clearer matches and note for a future algorithm which
recognizes contours and matches within...

projection shows a different neighboring region for the book. so theta is very good for identifying the book consistently, but object needs to be on an unvarying background/foreground or one should only compare theta within the contour (bounds) of the book.



true match (not found)

false match (found)

this is an example where a repeated pattern is better matched elsewhere in the image due to 2 things: (1) it's a pattern (very big texture) and (2) the background/foreground changes near the corner change the appearance enough that the true match isn't found.

where “found” is referring to the matching algorithm to make rough correspondence lists before use of the epipolar projection solver.

might need to consider the top k matches, especially if the image is known to contain repeated patterns/textures.

Euclidean transformation from pairwise calc of known points:

rotationInRadians=4.5554905

rotationInDegrees=261.01037890524134 scale=1.004318

translationX=266.80643

translationY=5.852234 originX=0.0 originY=0.0

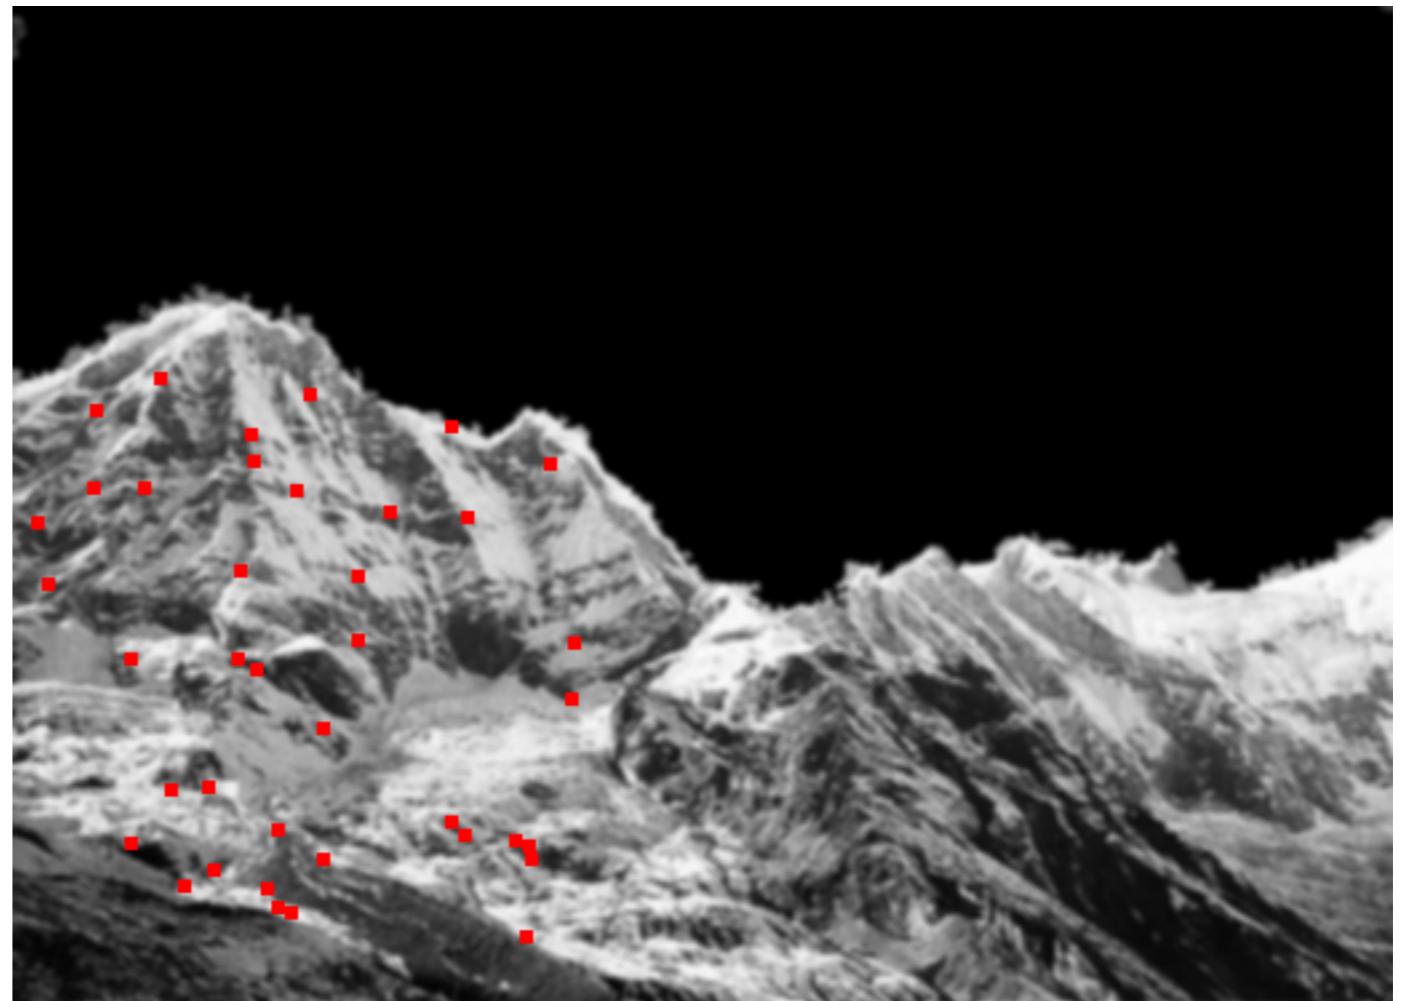
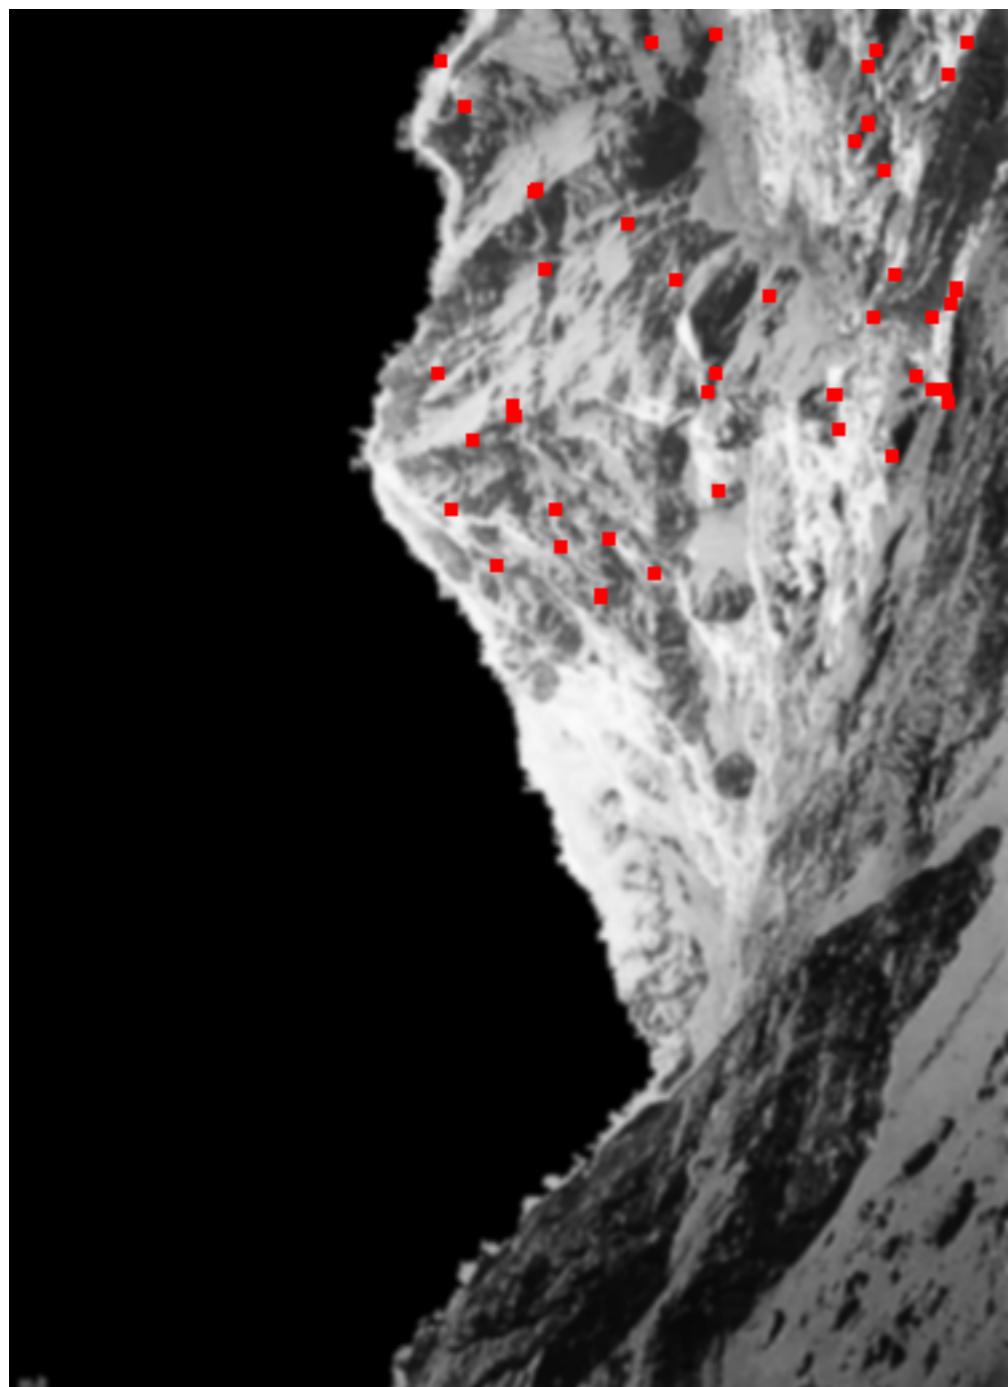
Feature matches of corner regions followed by a look at the implied rotations by frequency then filtering with that to see the most frequent translation in X then filtering by that to find the most frequent translation in Y:

solution

rotation=270.0+-20

translationX=215.0+-55.847068786621094

translationY=-25.0+=29.851972579956055



result is a list of correspondence, usable as input to the epipolar solver which uses RANSAC to discard outliers while solving the fundamental matrix.

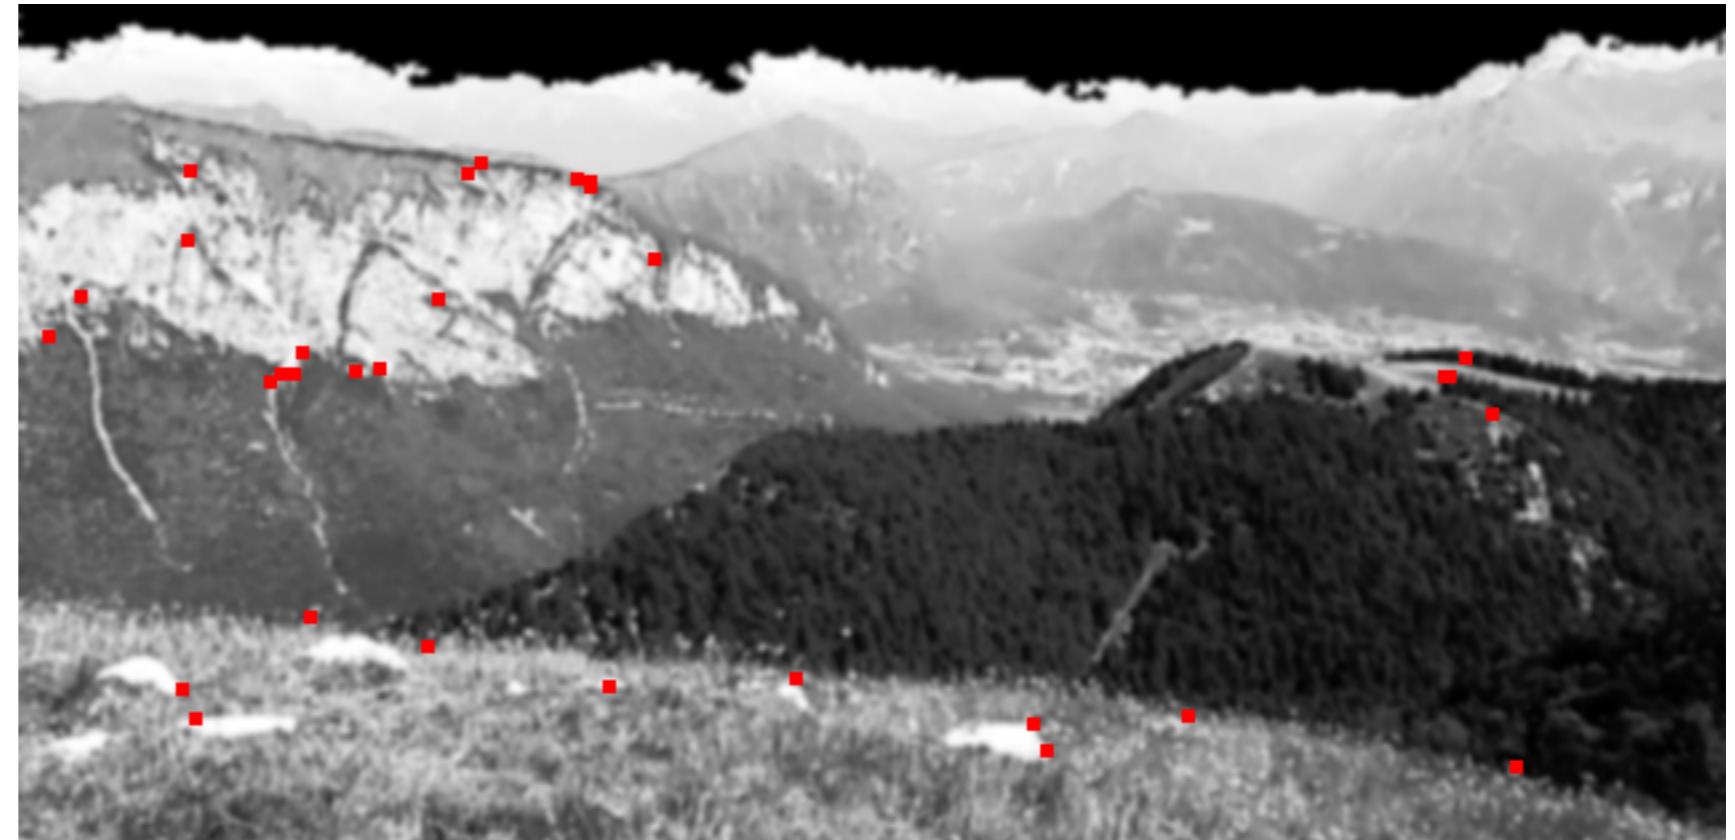
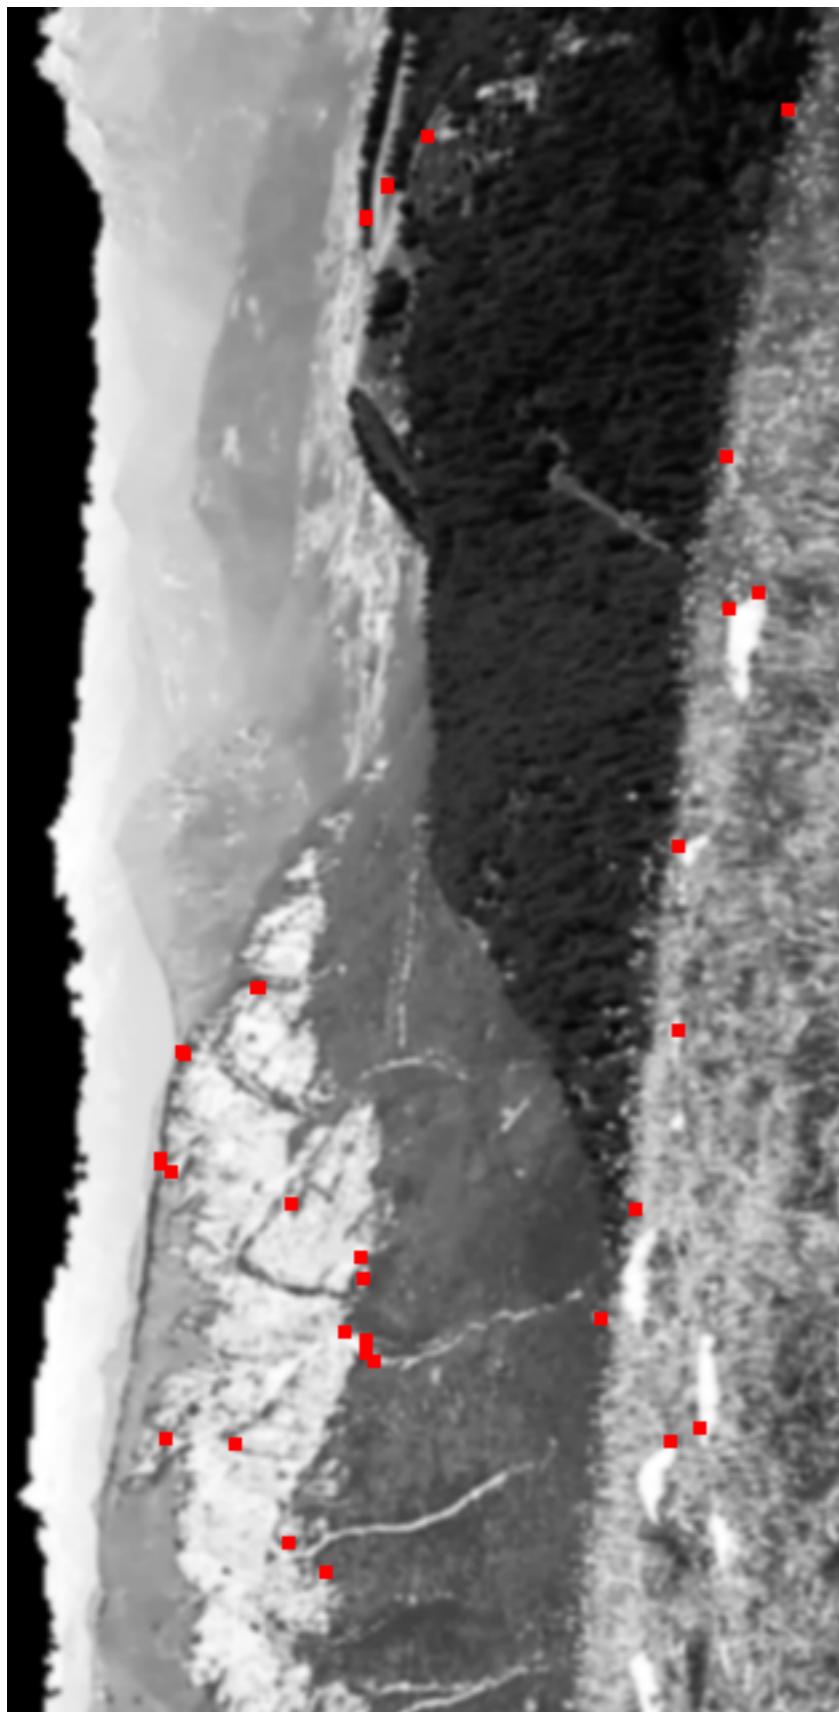
Euclidean transformation from pairwise calc of known points:

```
params=rotationInRadians=4.712166  
rotationInDegrees=269.9872145691191 scale=1.0279127  
translationX=614.1122 translationY=-4.460491  
originX=0.0 originY=0.0
```

Feature matches of corner regions followed by a look at the implied rotations by frequency then filtering with that to see the most frequent translation in X then filtering by that to find the most frequent translation in Y:

solution

```
rotation=280.0+-20  
translationX=625.0+-31.55487060546875  
translationY=5.0+=0.0
```



result is a list of correspondence, usable as input to the epipolar solver which uses RANSAC to discard outliers while solving the fundamental matrix.

Euclidean transformation from pairwise
calc of known points:

params=rotationInRadians=4.512667
rotationInDegrees=258.556783710076
scale=1.0265783
translationX=639.77747 translationY=12.565971
originX=0.0 originY=0.0

Feature matches of corner regions followed by a
look at the implied rotations by frequency then
filtering with that to see the most frequent
translation in X then filtering by that to find
the most frequent translation in Y:

solution
rotation=270.0+-20
translationX=625.0+-55.98248291015625
translationY=25.0+=44.471153259277344



On already rectified images, the rough euclidean solution that is used to remove more distant matched pairs probably needs a 2nd order term, at least for translations.

Have algorithm to match features, but still need to know scale between the images.

For many sets this is already known to be '1', but for images of different cameras, etc, it isn't known. A characteristic scale for both images can be found using scale space, but that still requires identifying similar objects in both images or a similar population of objects.

With CSS, can make scale space maps with inflection points for a large range of Gaussian sigma to find the peaks of the inflection points in the scale space maps. The same point's sigma divided by the same in the other image's scale space maps gives scale (and location of those points further gives rotation and translation with 2 sets of points).

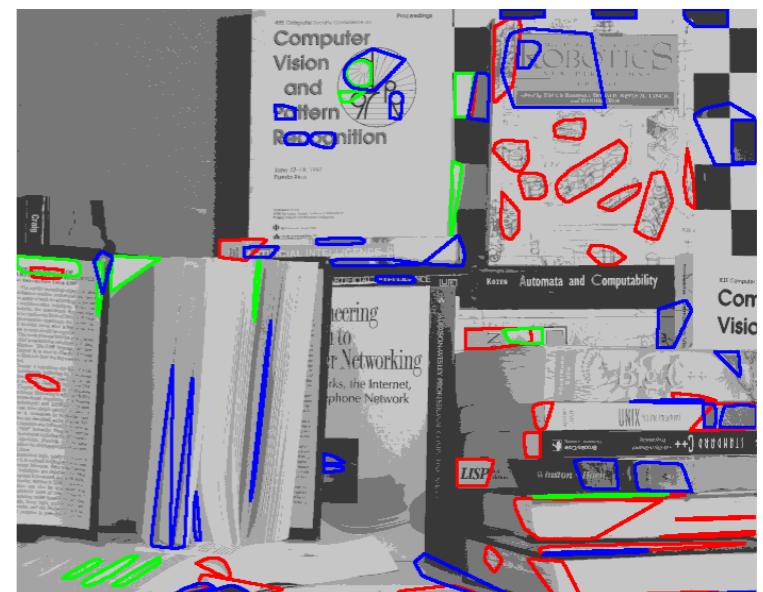
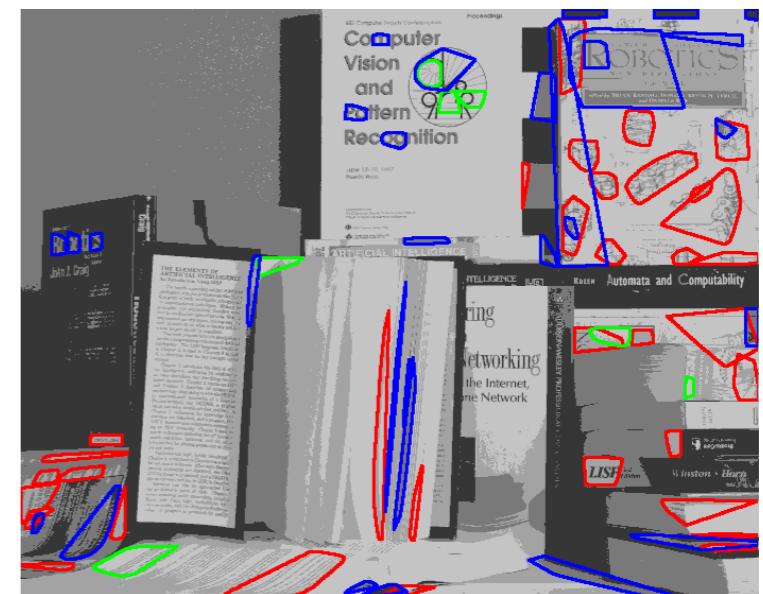
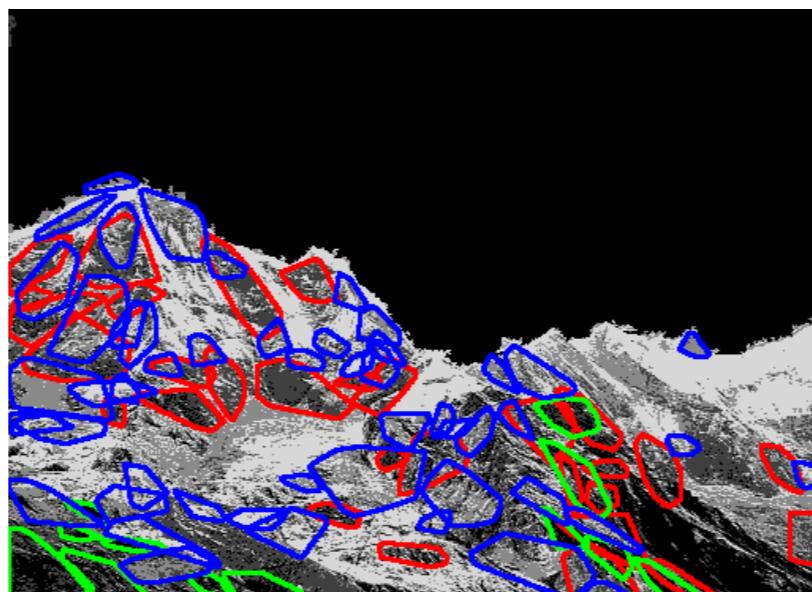
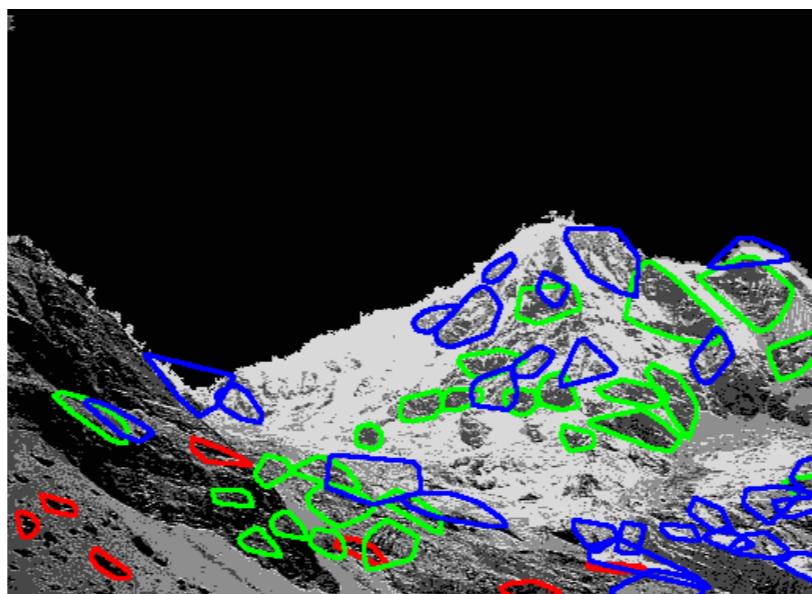
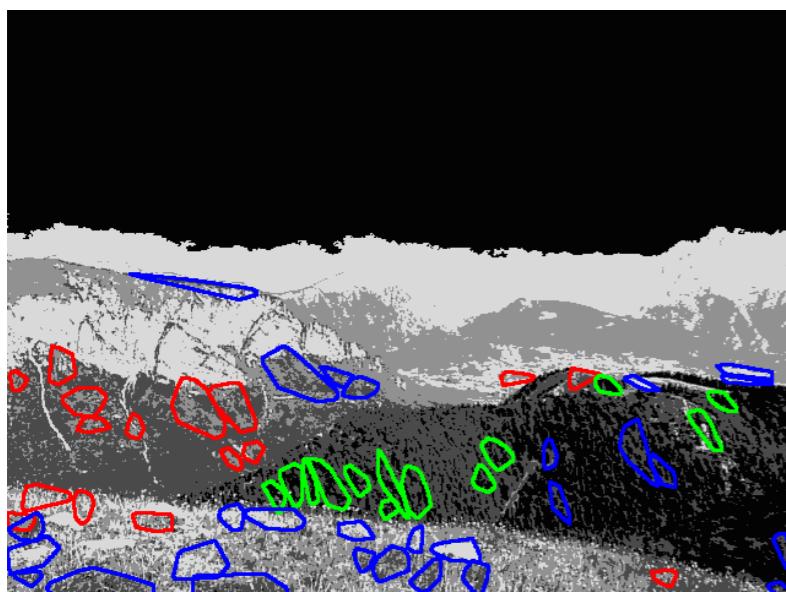
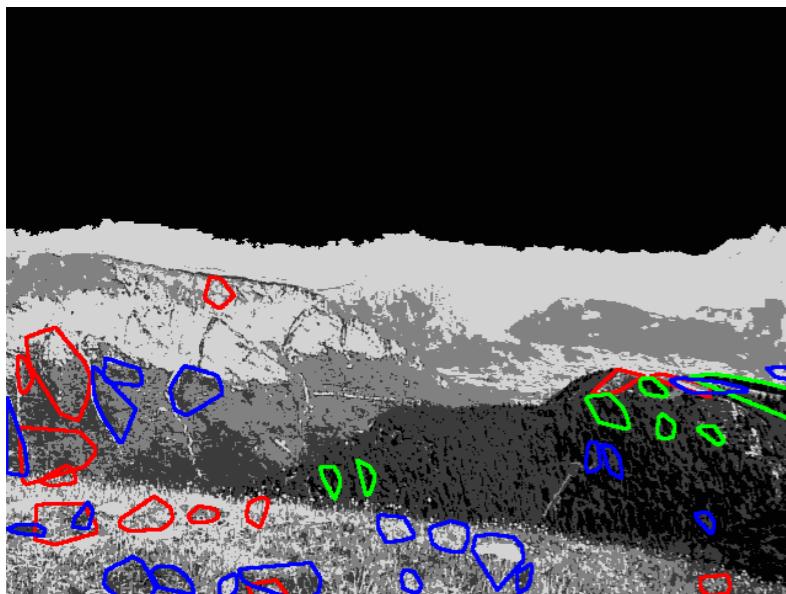
In order to use scale space to solve for scale, need to make sure that the comparison is among same object population or specifically identify the objects with methods such as in the ContourMatcher, ContourFinder and InflectionMapper (see doc/contours.pdf too).

Because edges don't necessarily by default in this project create contours (closed curves enclosing objects of interest), will try the segmentation technique and find contiguous pixels of a given intensity level within a range of sizes. Will either use the convex hulls of those or the concave hull approximation (PerimeterFinder) to create closed curves for those. Then will find the inflection points on those and create scale space maps for those and use the contour matcher to solve for the scale. This looks to be feasible with the start images binned to a size < 200 x 200 and segmentation somewhere between k=2 and k = 5.

Trying blob finding in next slides.

- perform histogram equalization if mean is too far from median or the two images are too different for those params.
- color segmentation of **k=4** to reduce image to 4 bands of intensities.
- contiguous pixel group finder for each of the 4 bands using point limits of:
smallestGroupLimit = **100**; largestGroupLimit = **1000**;

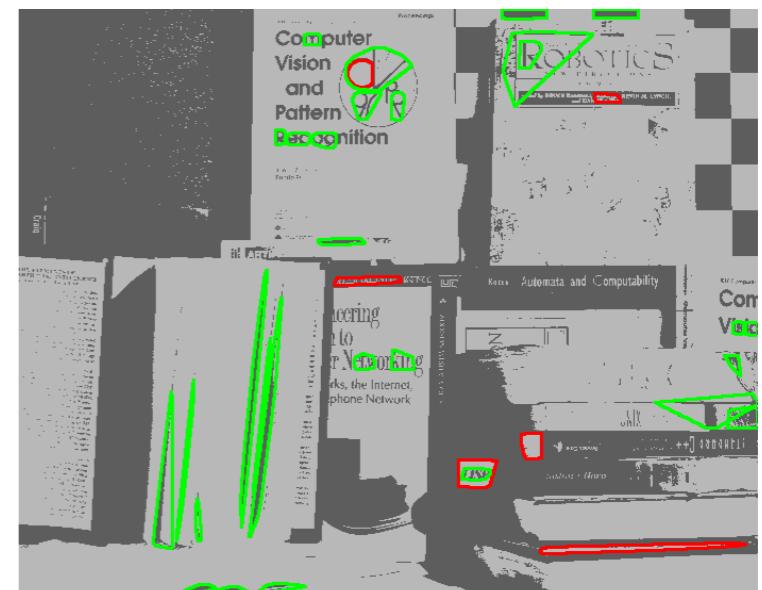
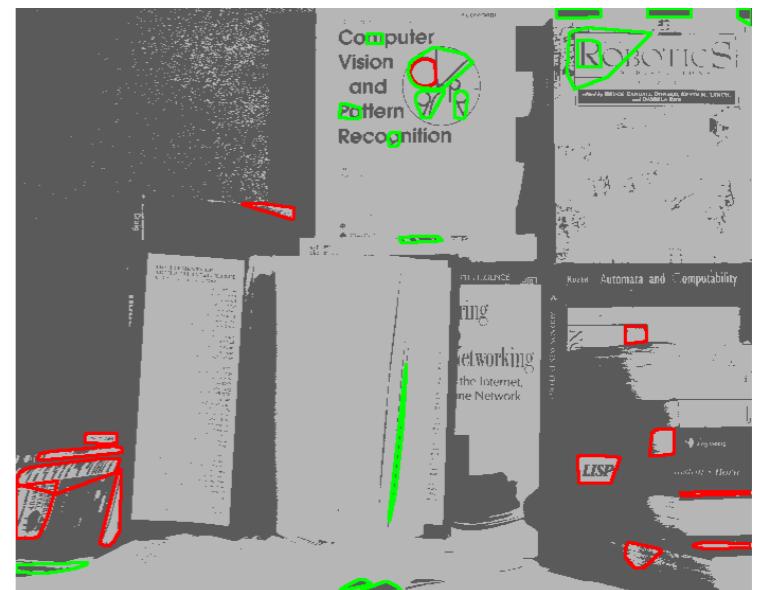
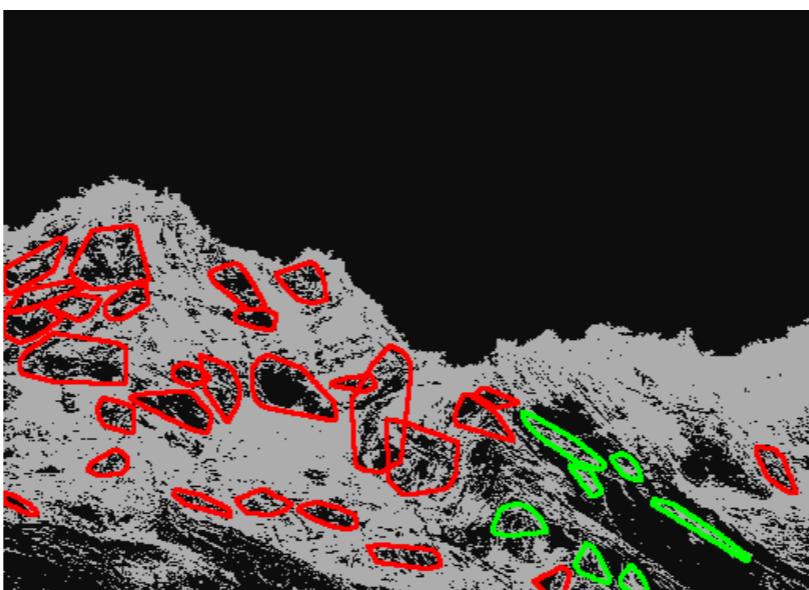
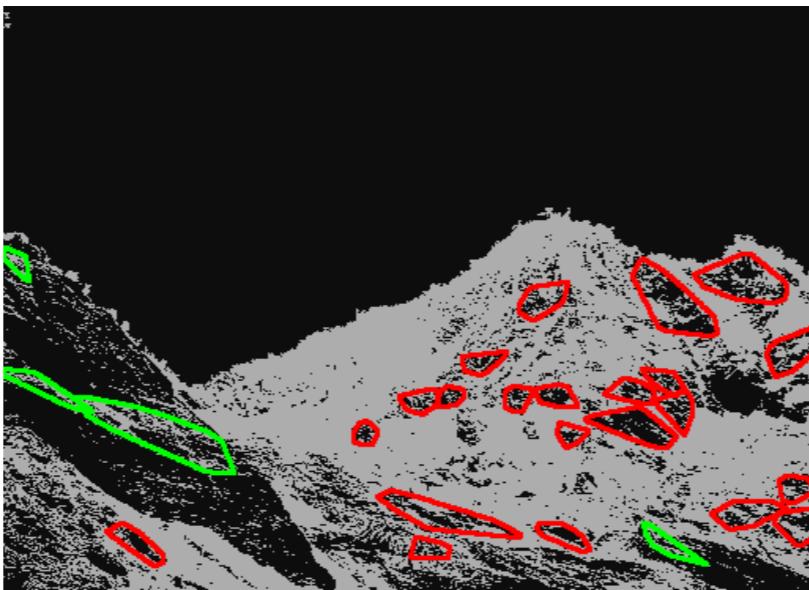
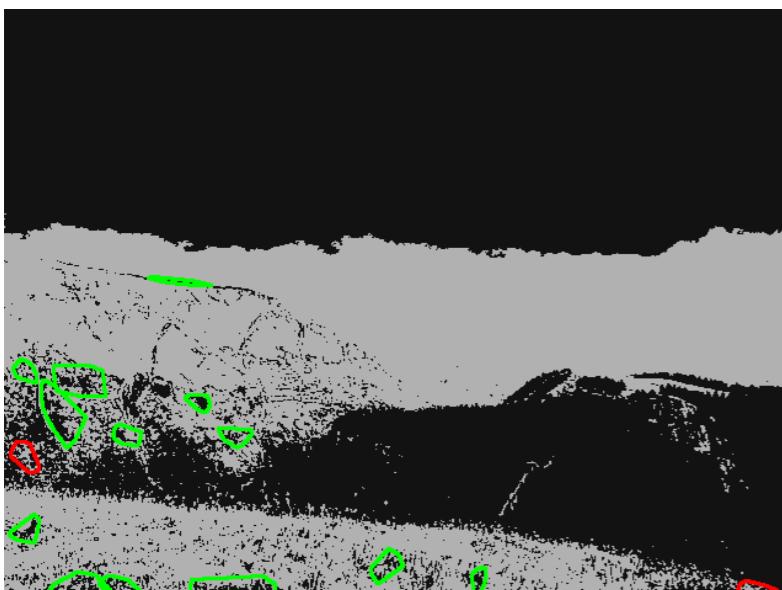
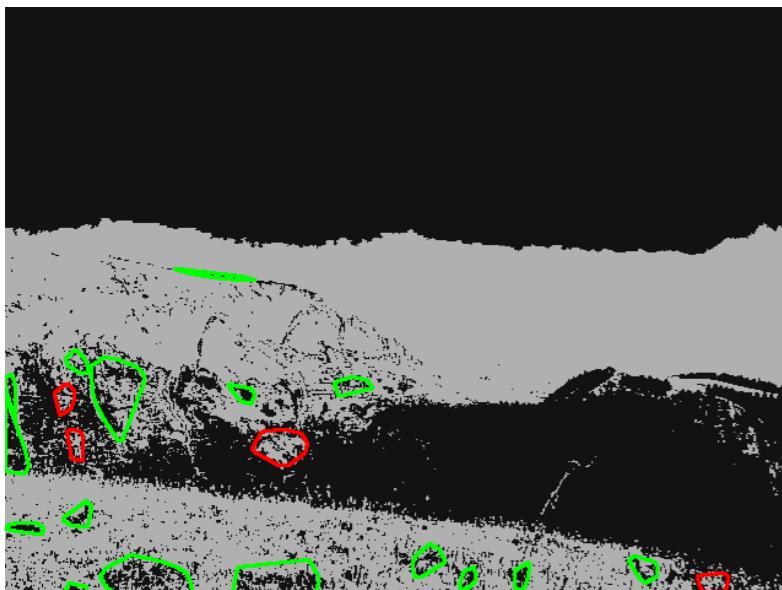
Would prefer fewer features in the Brown & Lowe.



Even better:

- perform histogram equalization if mean is too far from median or the two images are too different for those params.
- color segmentation of **k=2** to reduce image to 2 bands of intensities.
- contiguous pixel group finder for each of the 2 bands using point limits of:
smallestGroupLimit = **100**; largestGroupLimit = **1000**;

This is a smaller number of features that almost has all of the major blobs one would want to use for matching.

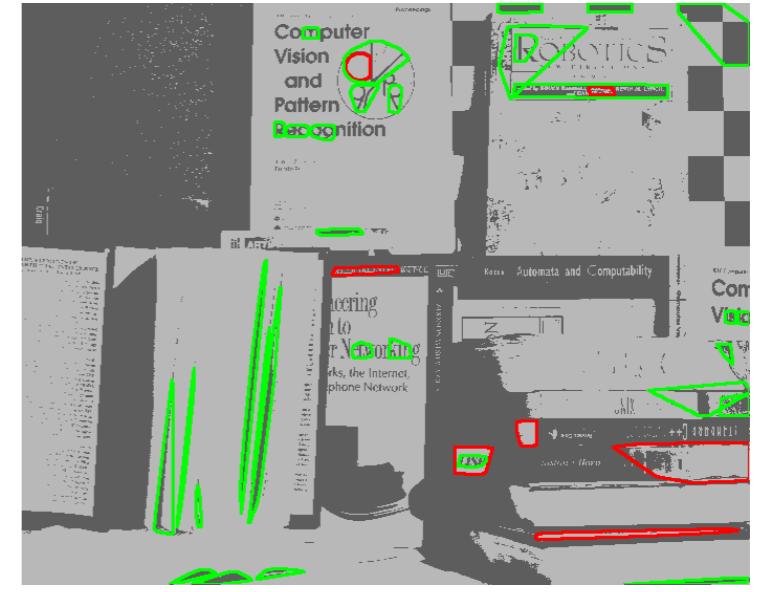
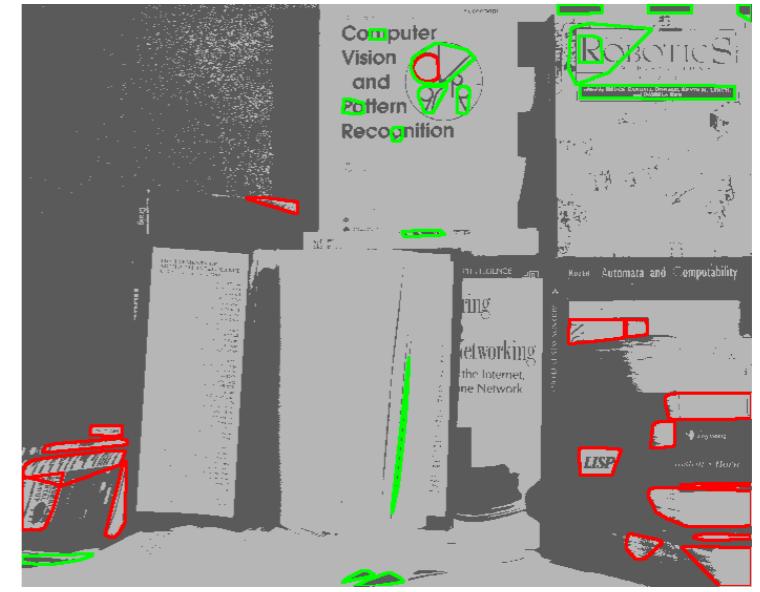
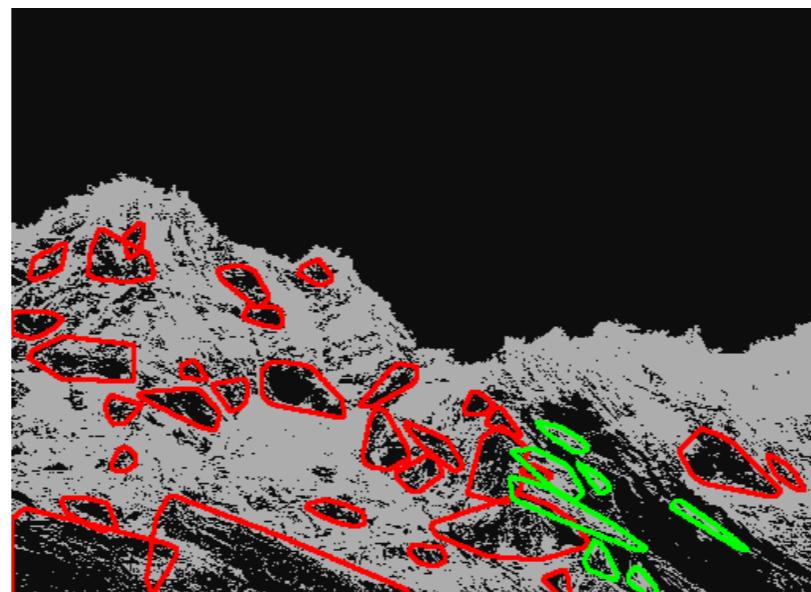
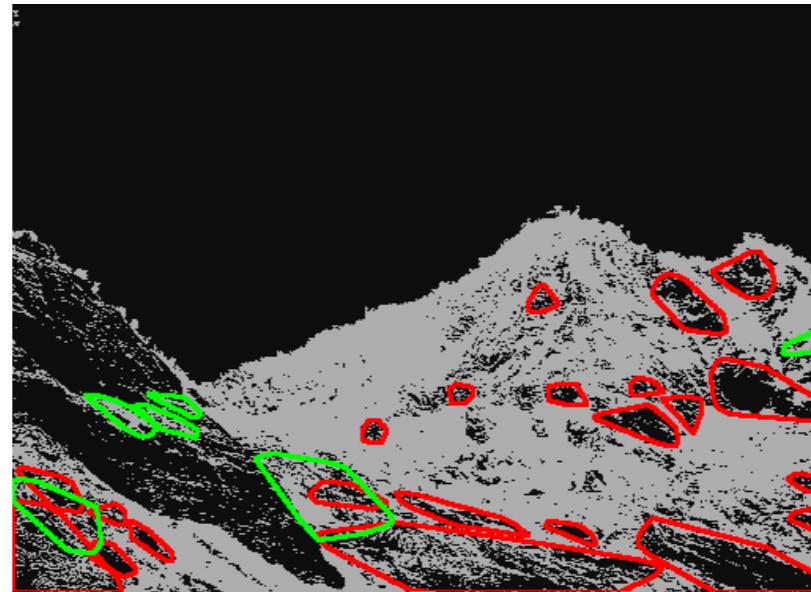
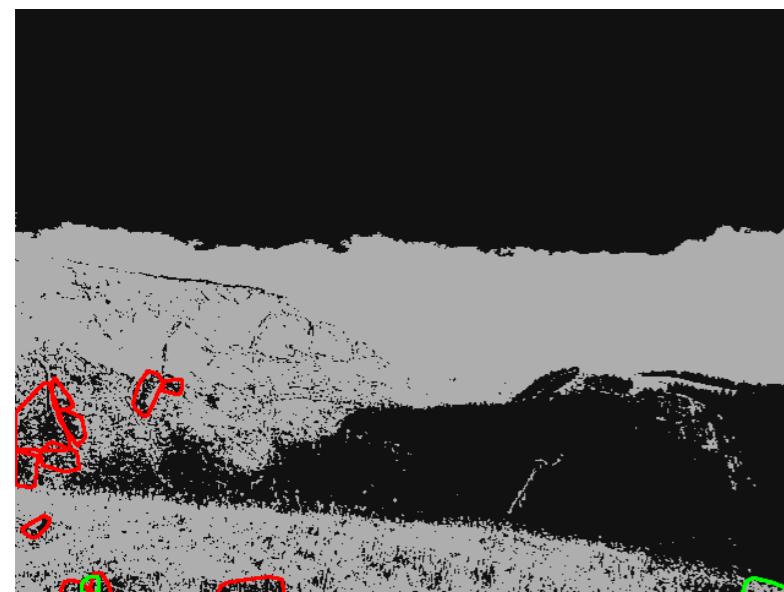
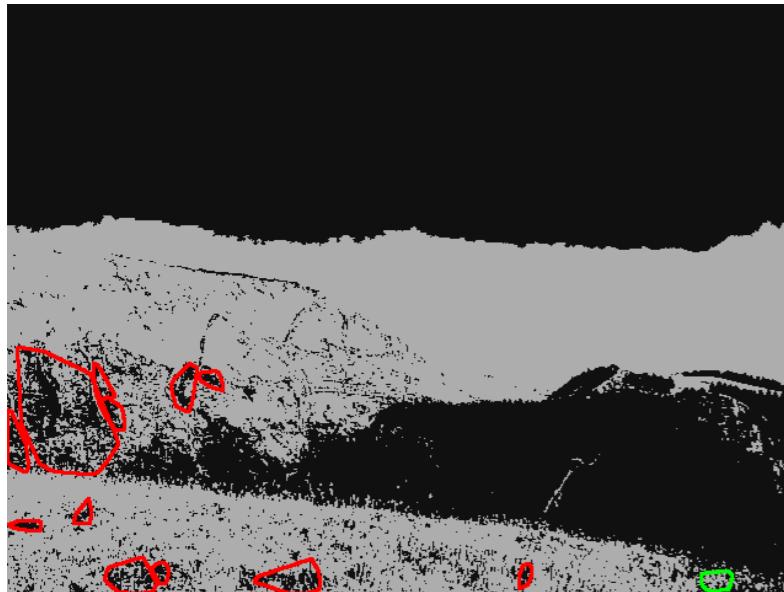




Even better:

- perform histogram equalization if mean is too far from median or the two images are too different for those stats.
- color segmentation of **k=2** to reduce image to 2 bands of intensities.
- contiguous pixel group finder for each of the 2 bands using point limits of: smallestGroupLimit = **100**; largestGroupLimit = **5000**;

correspondence looks like enough for an initial Euclidean solution

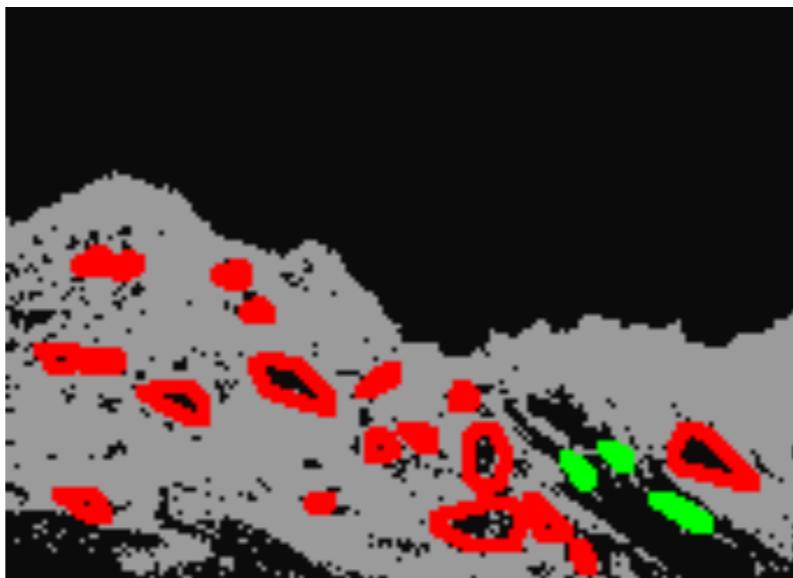
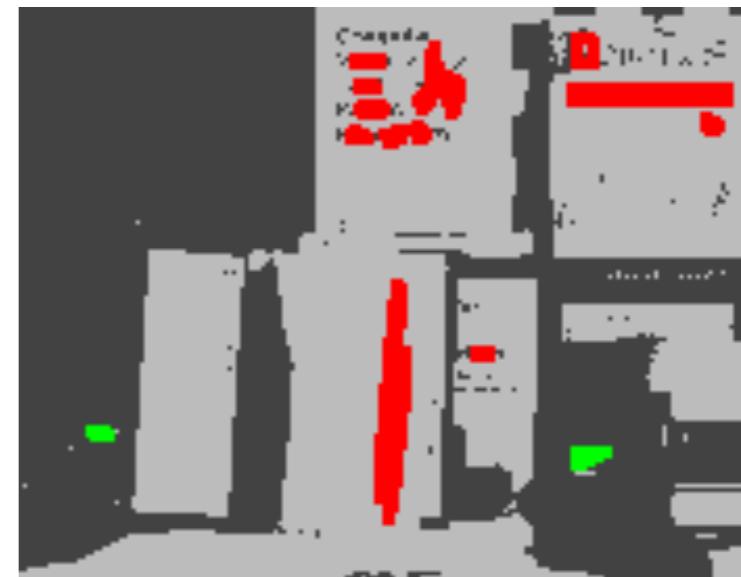
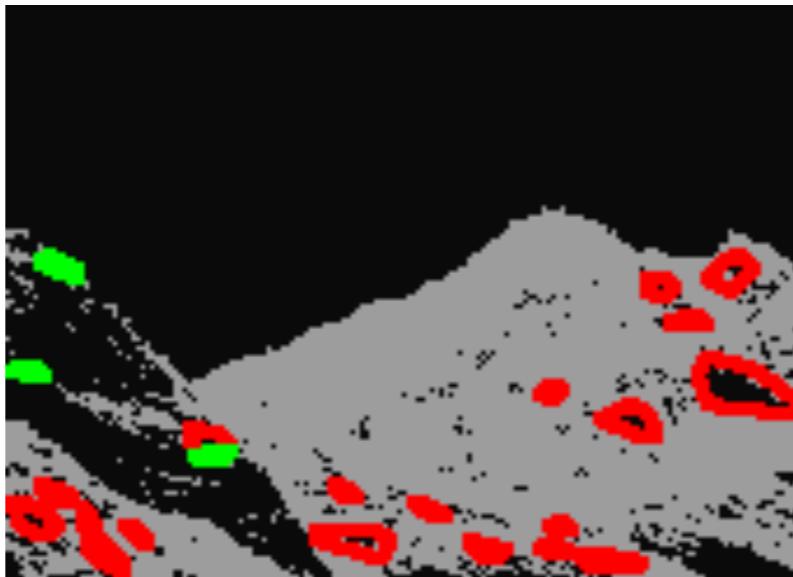
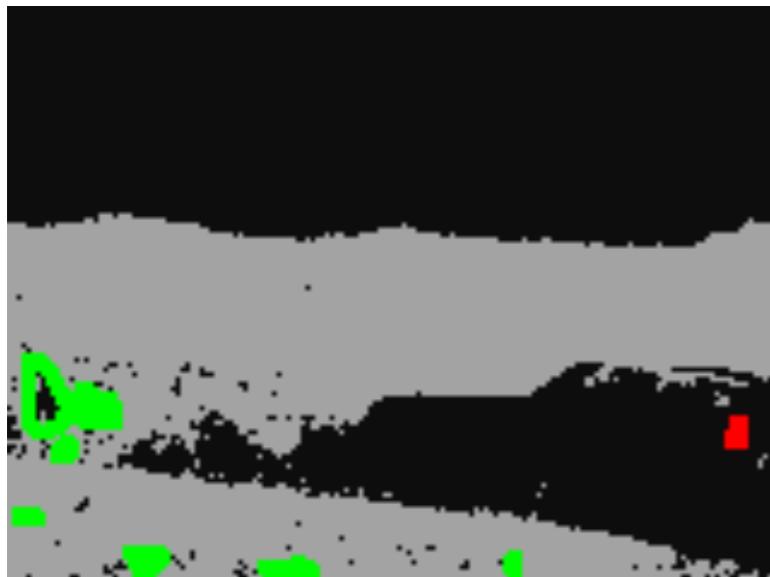


NOTE that blobs as first transformation solution works better for images like the middle where in contrast if done with corners, the ridge line in the differences would match more strongly than the true corner matches.

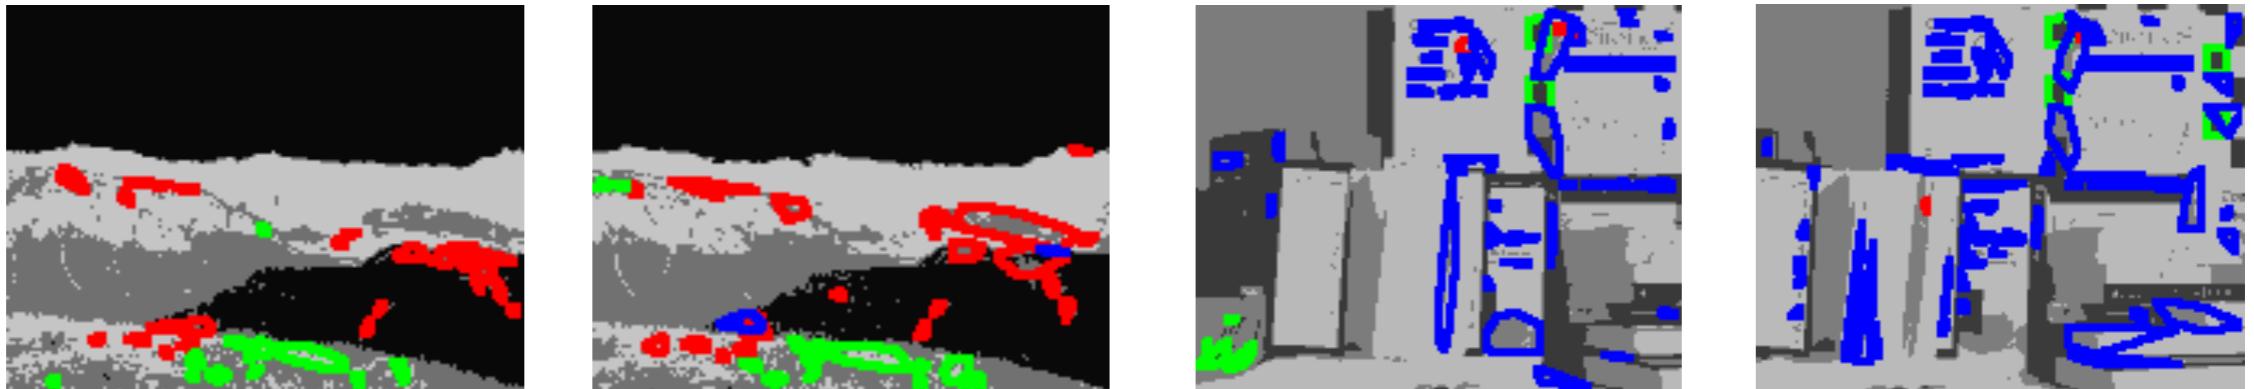
Alternatively for the blob (feature) finding, can reduce the image size:

- perform histogram equalization
- ***bin the image to < 200 X 200***
- color segmentation of ***k=2*** to reduce image to 2 bands of intensities.
- contiguous pixel group finder for each of the 2 bands using point limits of:
smallestGroupLimit = ***100/(binFactor^2)***; largestGroupLimit = ***5000/(binFactor^2)***;

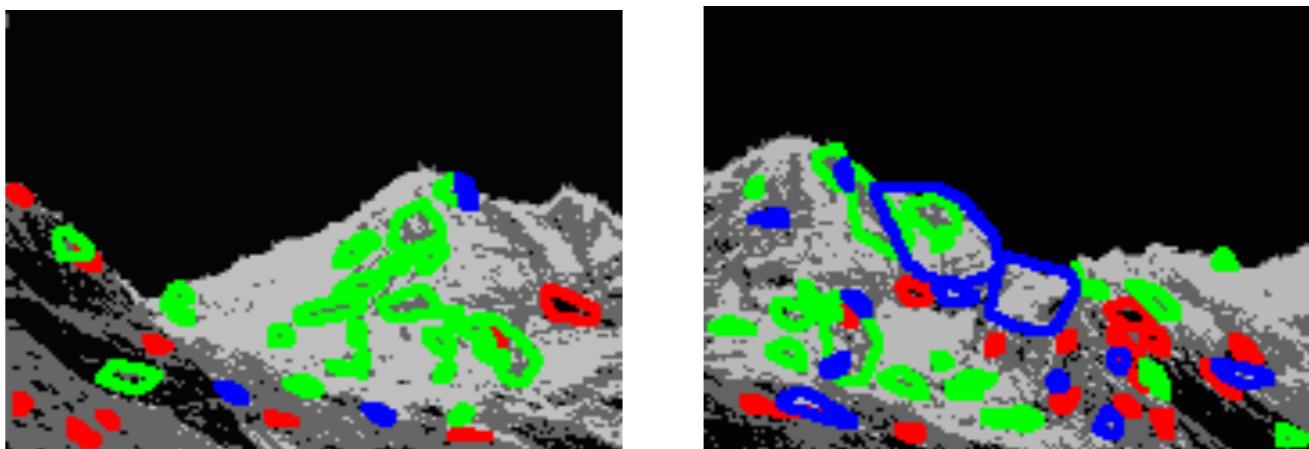
This solution doesn't have a better matchable to non-matchable blob fraction than the full image, but making scale space images should be much faster so will try binned first.



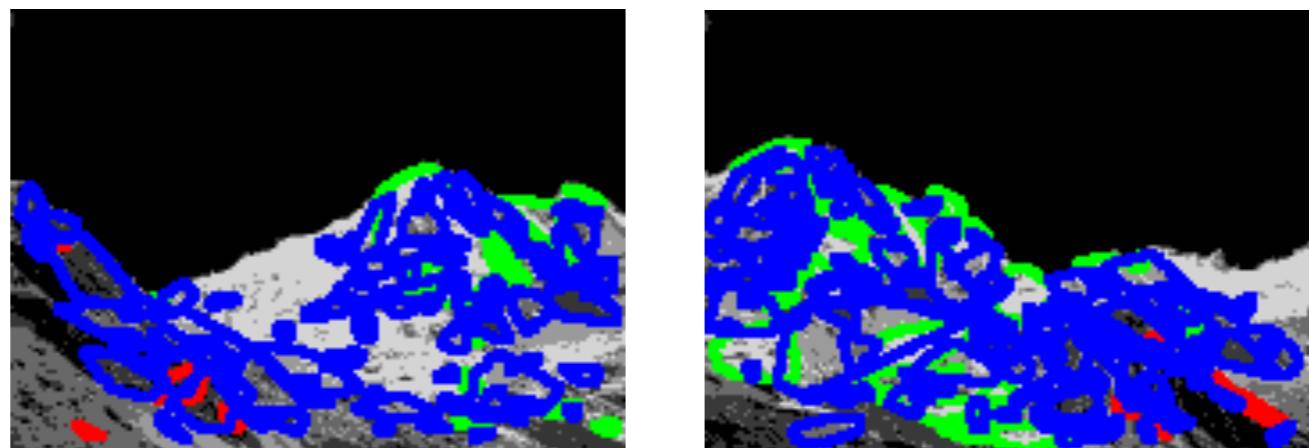
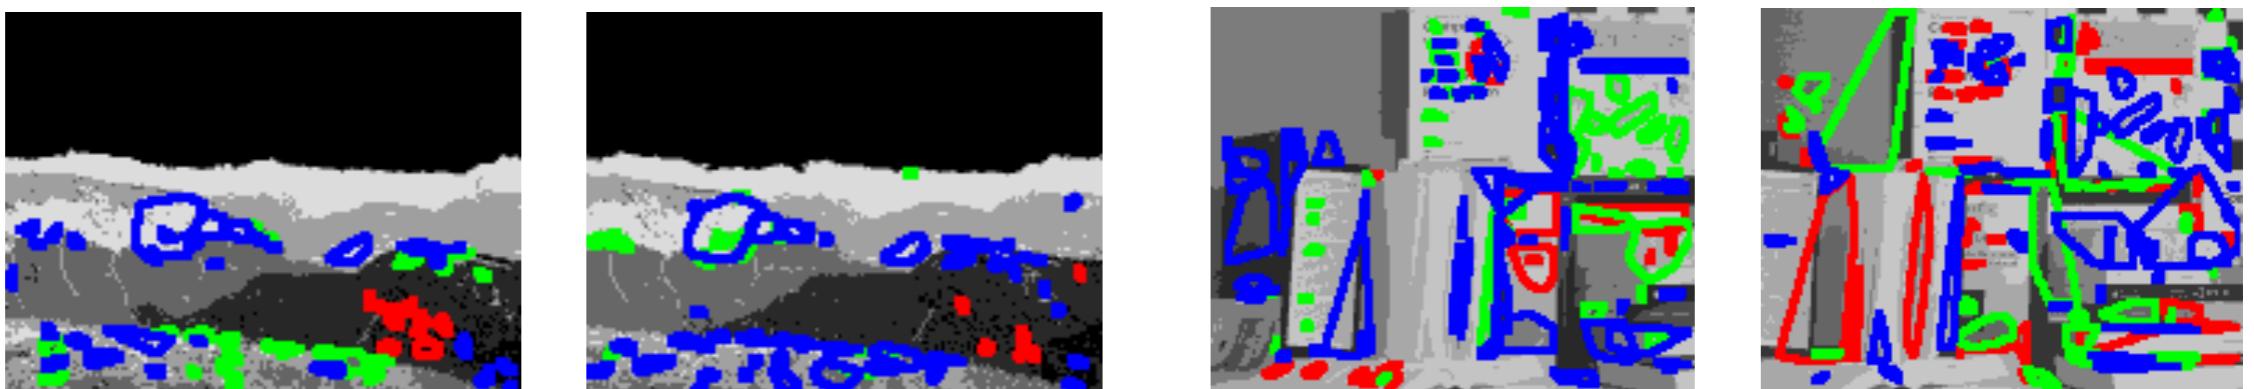
$k=3$



$k=3$ is very good
for venturi but
not for brown &
lowe



$k=5$



Looks like $k=2$ segmented images will produce blobs findable in both images but when the number of blobs is too low, can add the blobs for $k=3$ segmentation too.

in progress...



Implemented a blob ordered border closed curve extractor to create contours. Then the CSS contour matching code attempts to match the contours from the first image to the second.

Ideally, will be able to find objects in the first image and second image in similar ways and in smaller numbers, but large enough to match.

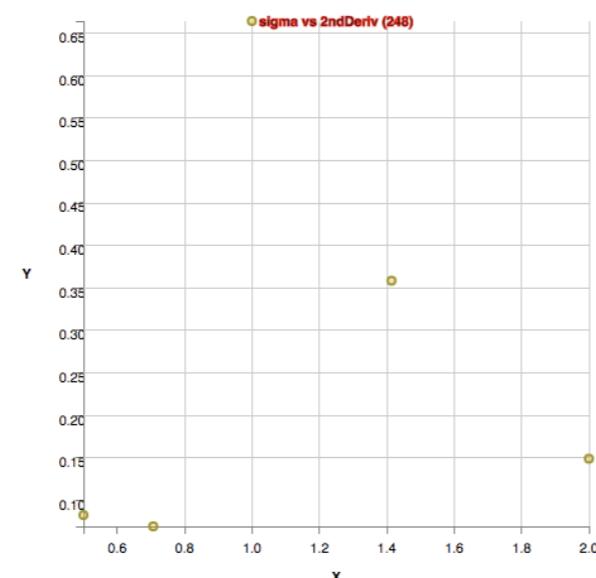
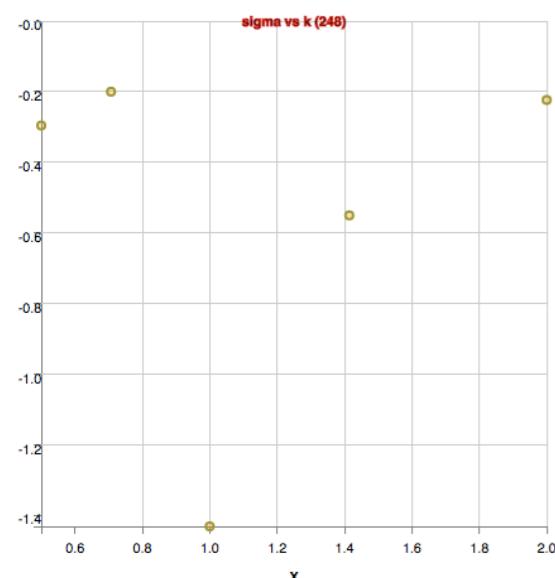
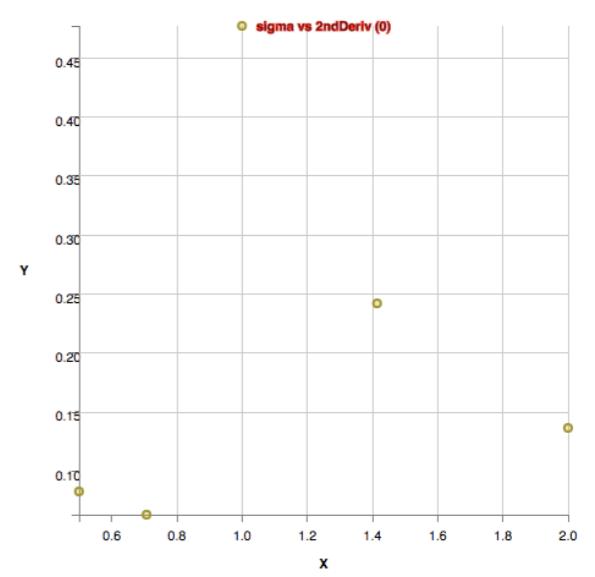
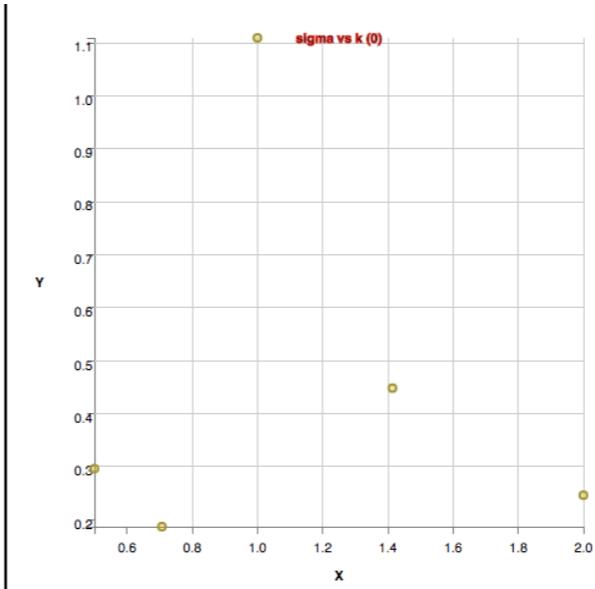
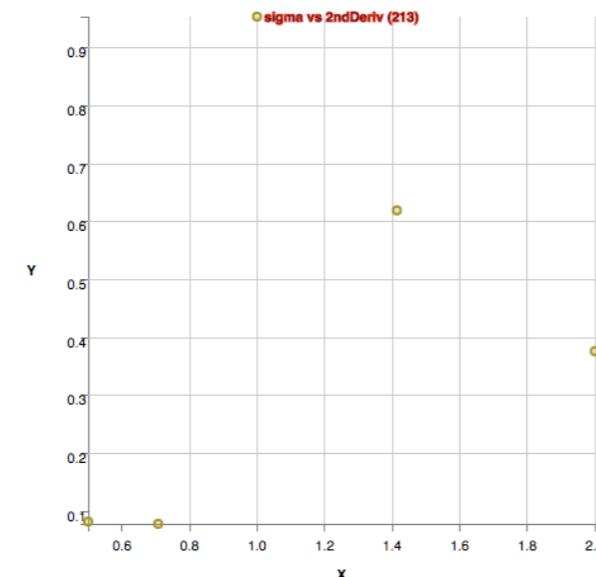
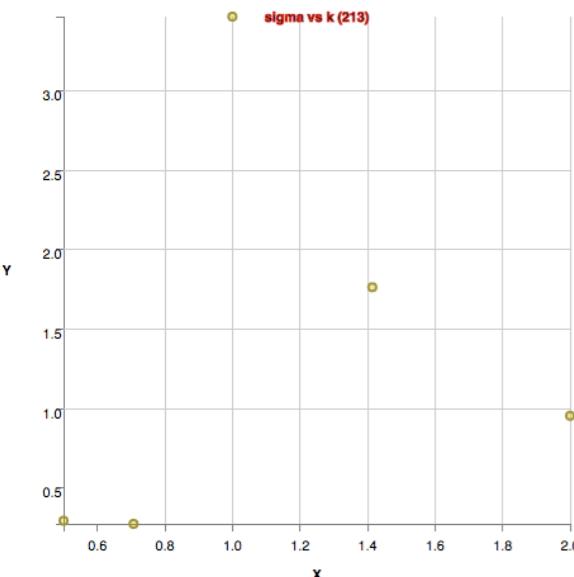
The result is **scale** and if more than one match is obtained, **rotation and translation** are also solved.

A quick look at “characteristic scale” from 2nd derivatives and curvature of corners in scale space maps.

The 2 images of the same objects are from the same camera and same camera settings, so the responses are similar. Shown are not the full range to sigma=256.

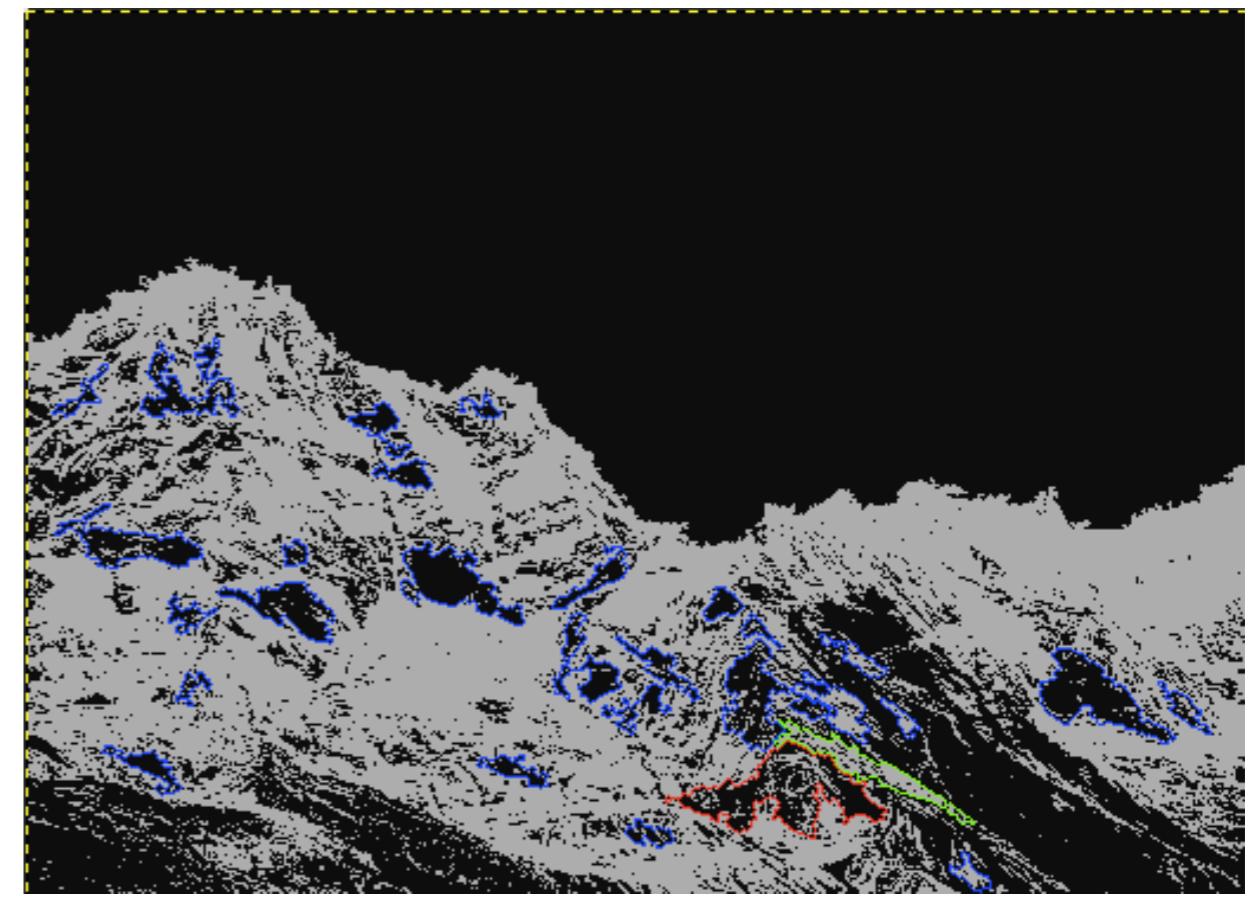
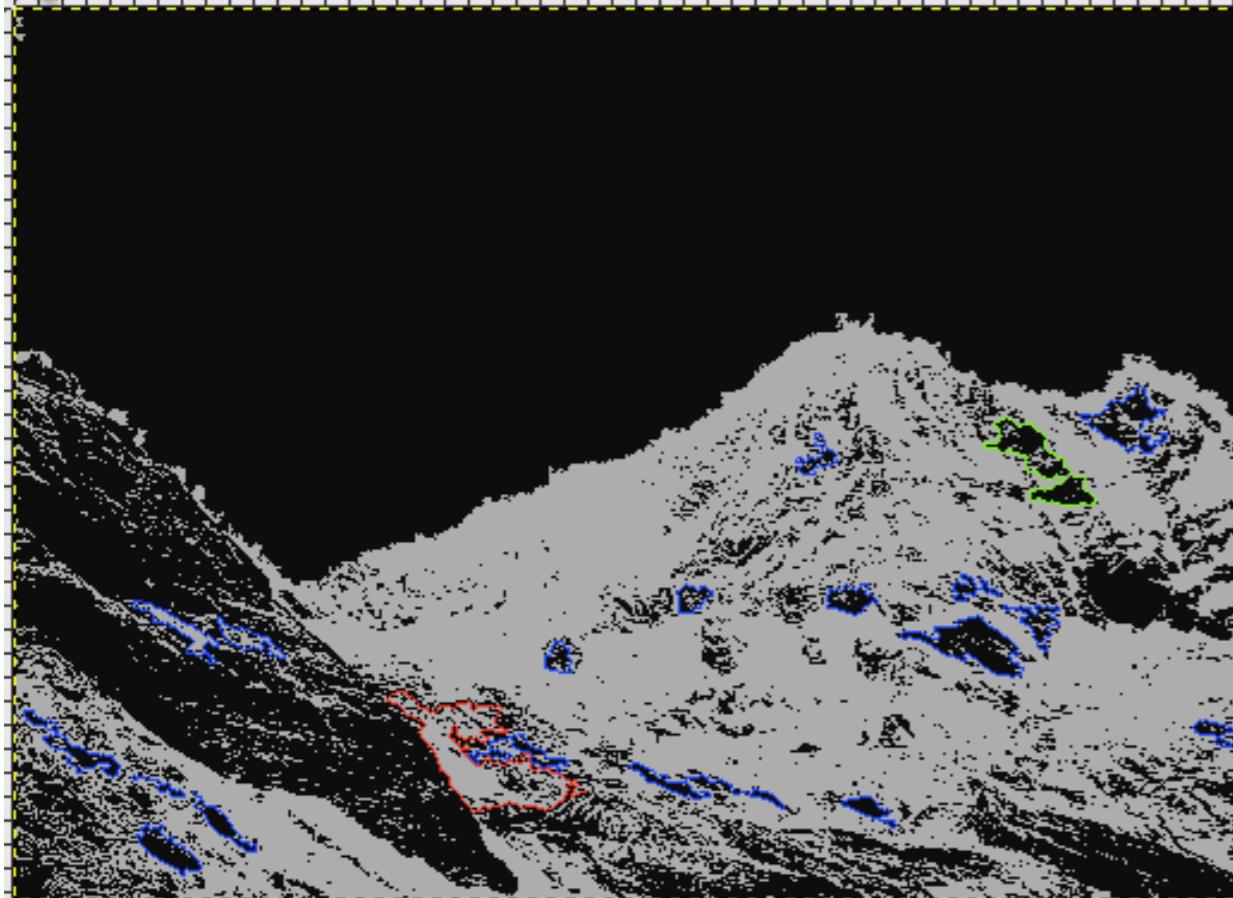
X' is gaussian first derivative, X'' is gaussian second derivative. t is unit along path in scale space. $\sigma \sim$ is gaussian sigma.

$$k(t, \sigma \sim) = \frac{X'(t, \sigma \sim) * Y''(t, \sigma \sim) - Y'(t, \sigma \sim) * X''(t, \sigma \sim)}{(X'^2(t, \sigma \sim) + Y'^2(t, \sigma \sim))^{1.5}}$$

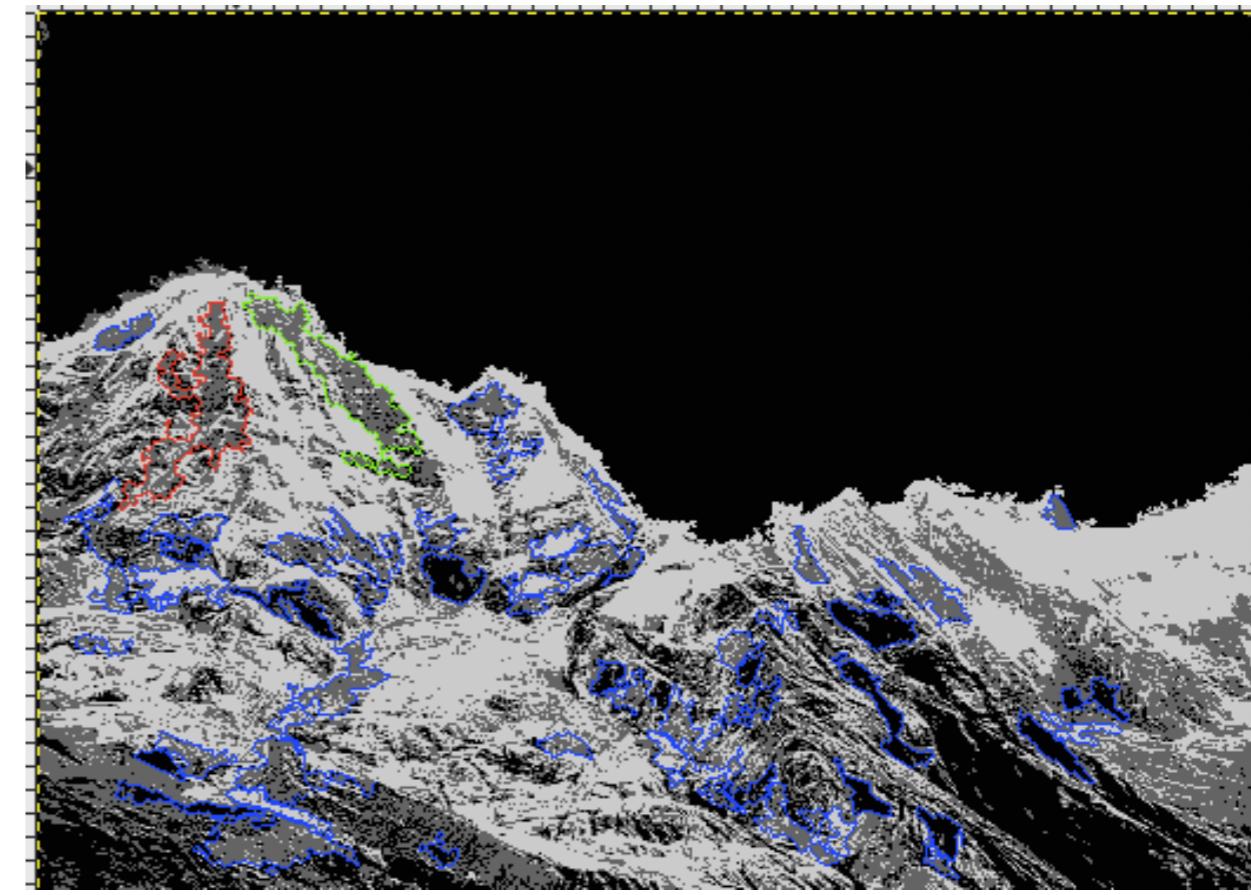
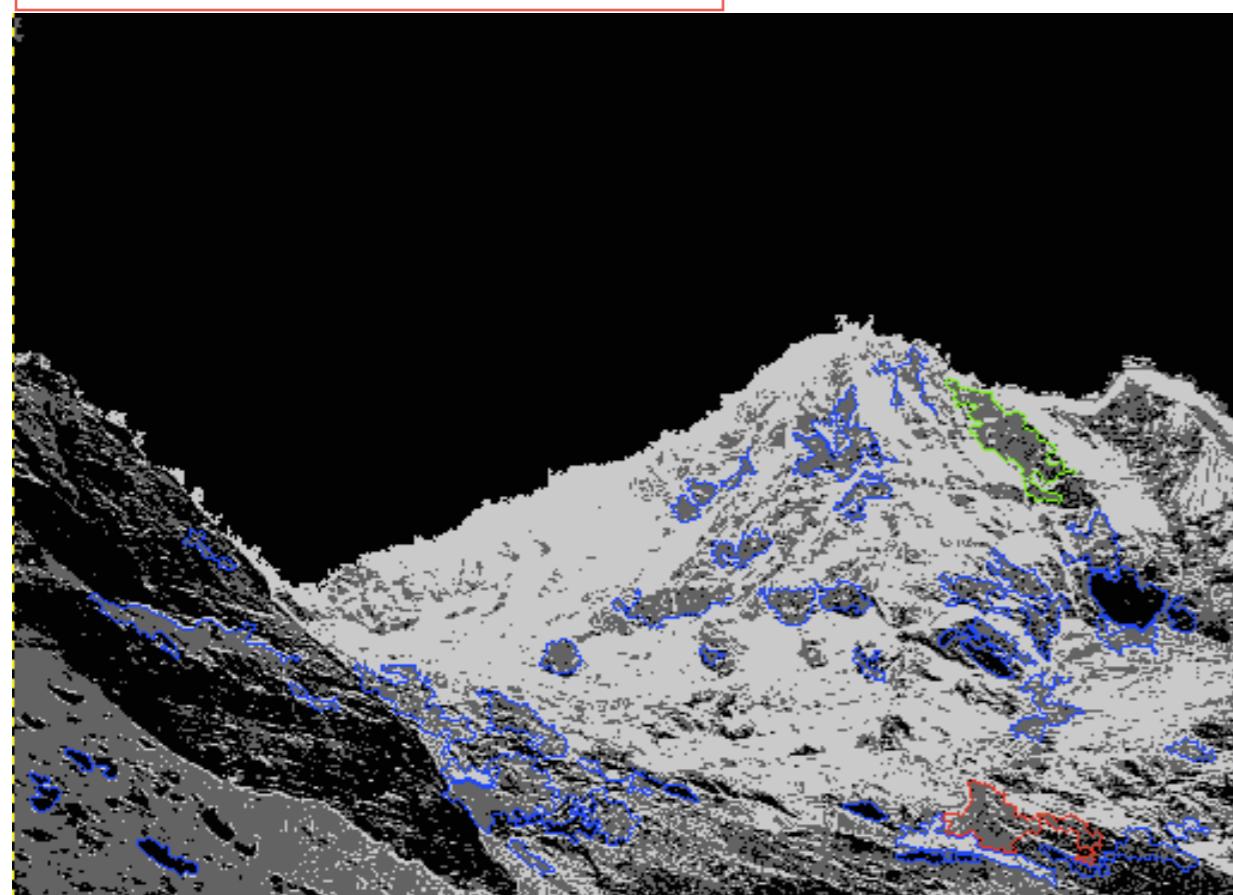




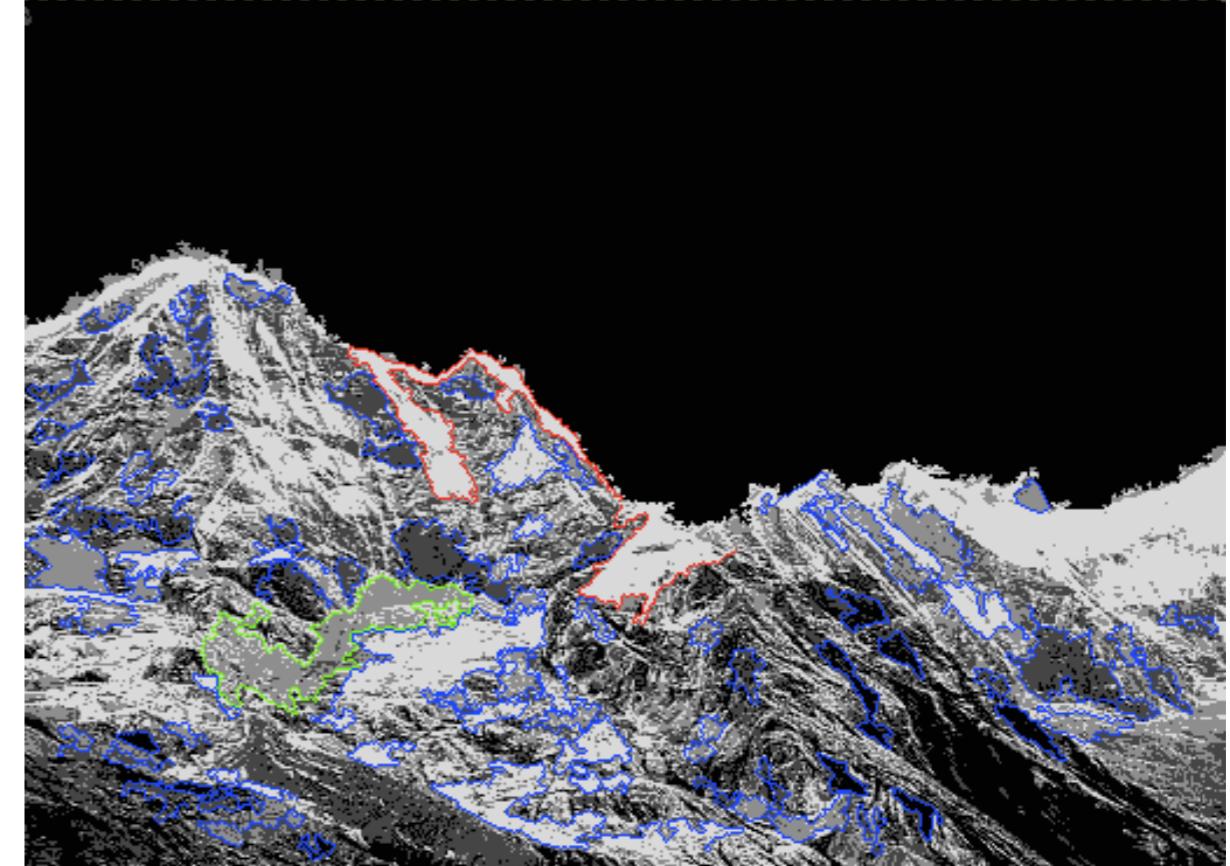
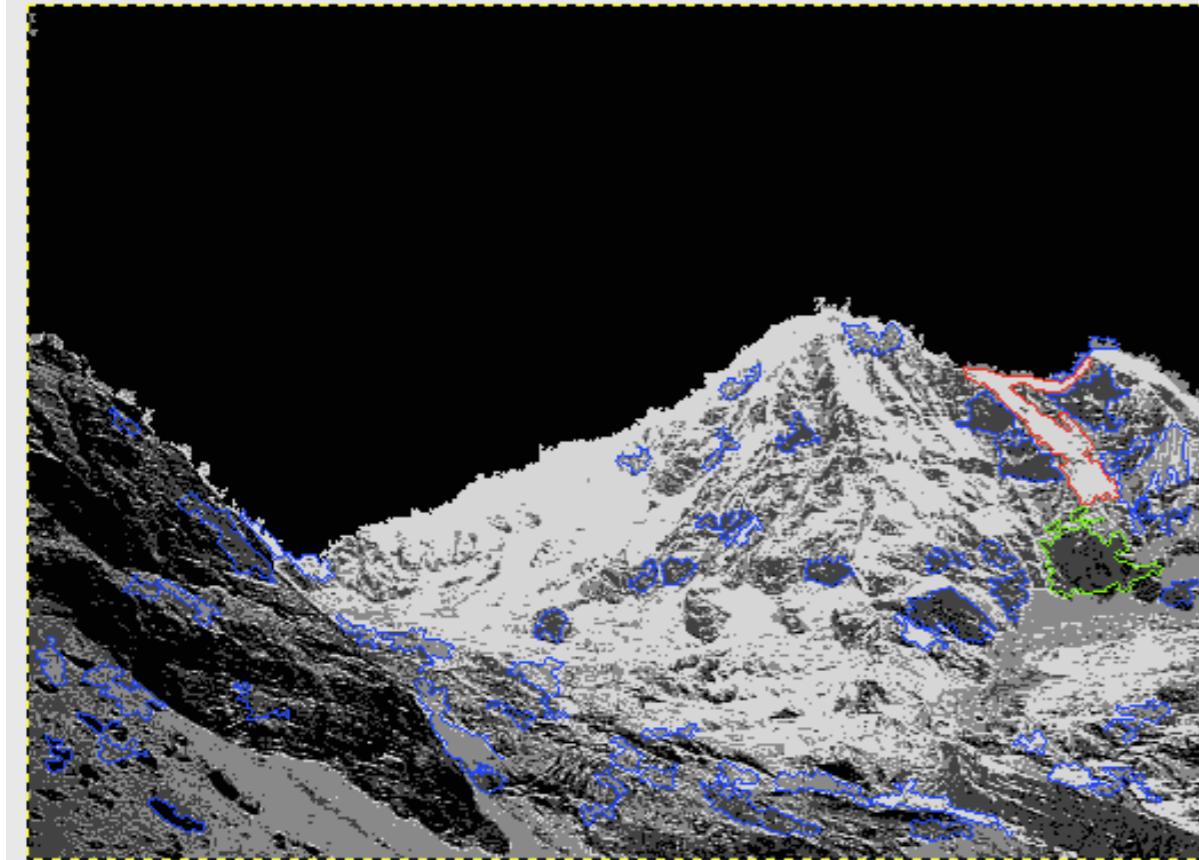
k=2 segmentation blobs. a look at blobs for use with contour matching to solve for euclidean scale. will combine with feature descriptors to only compare similar appearance contours.



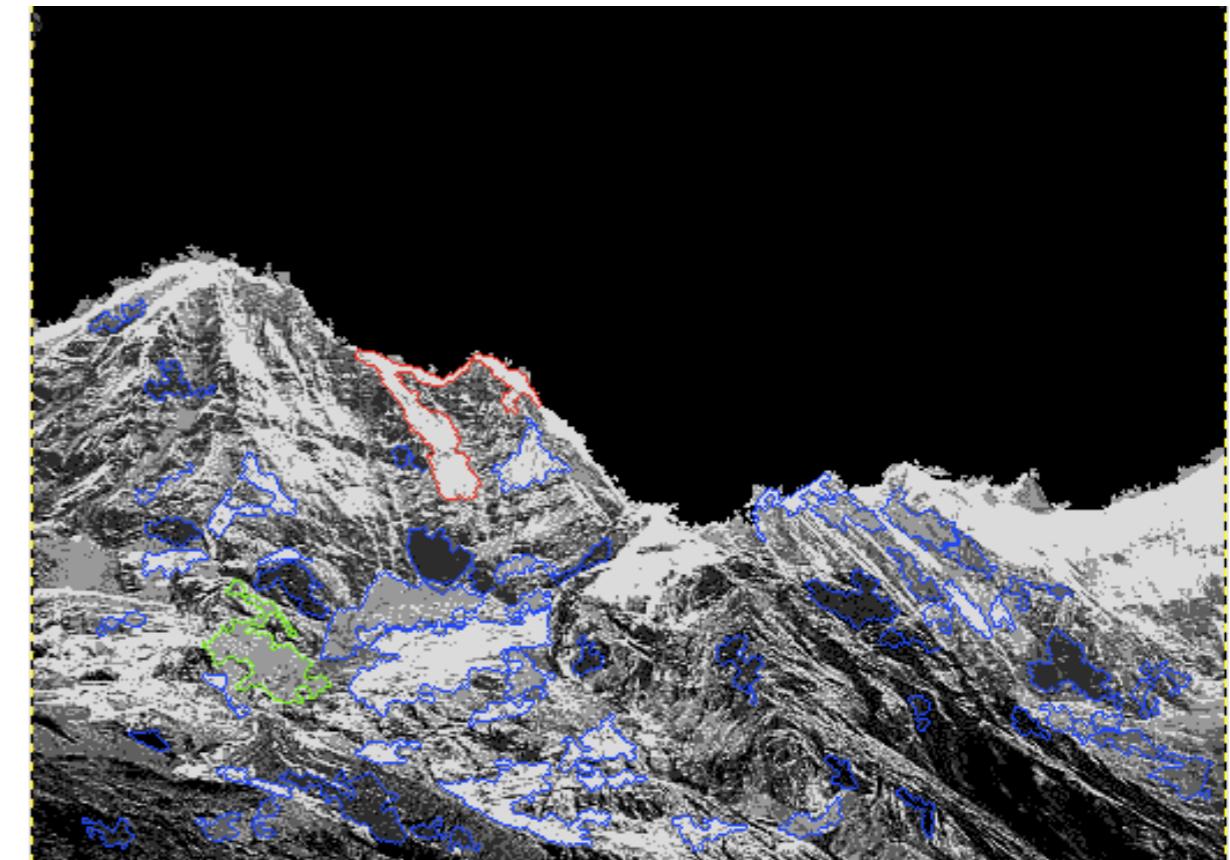
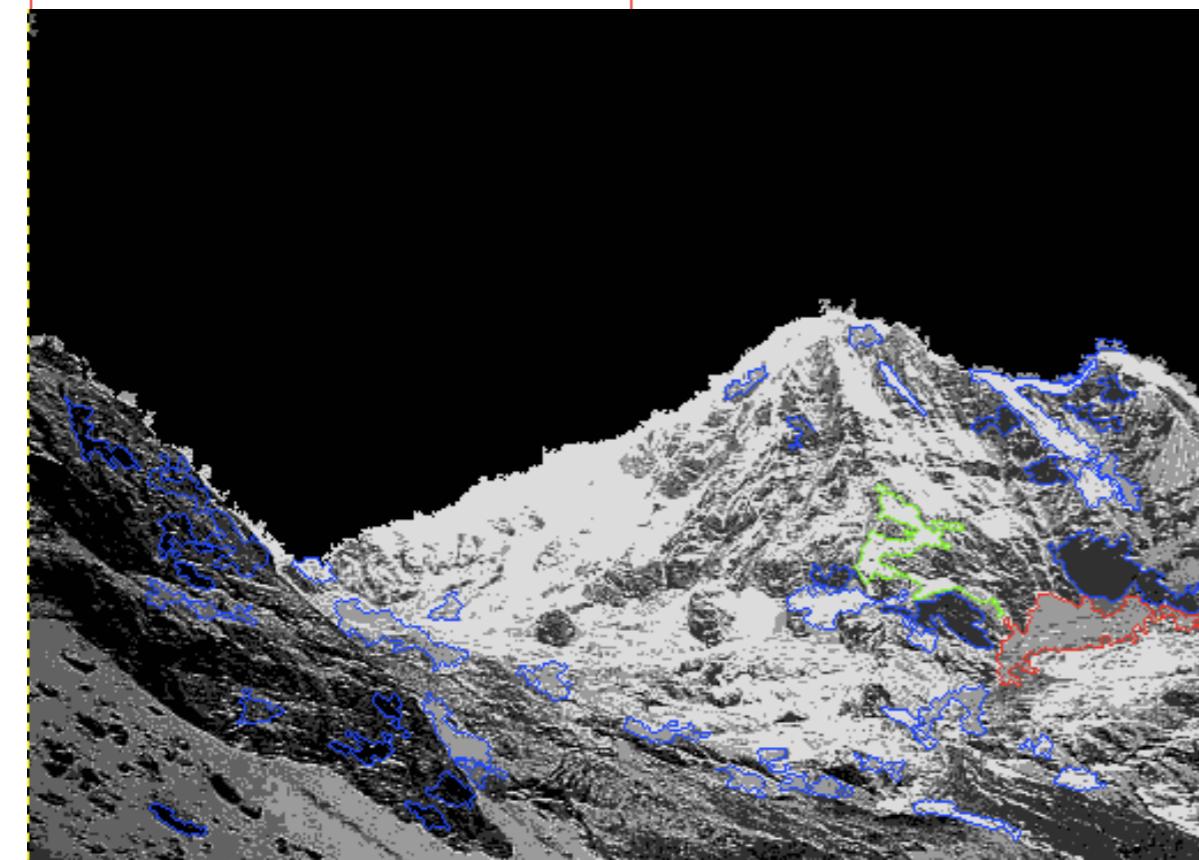
k=3 segmentation blobs



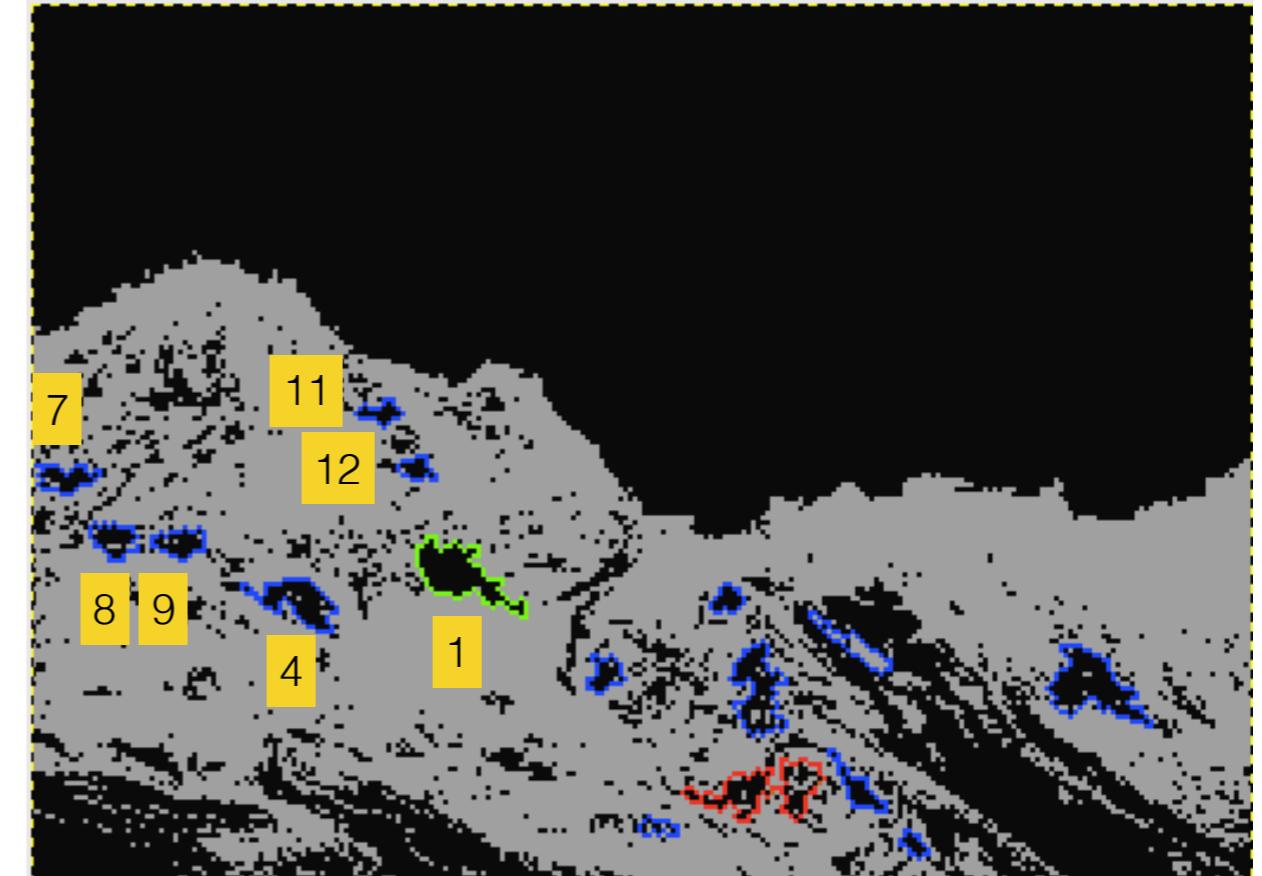
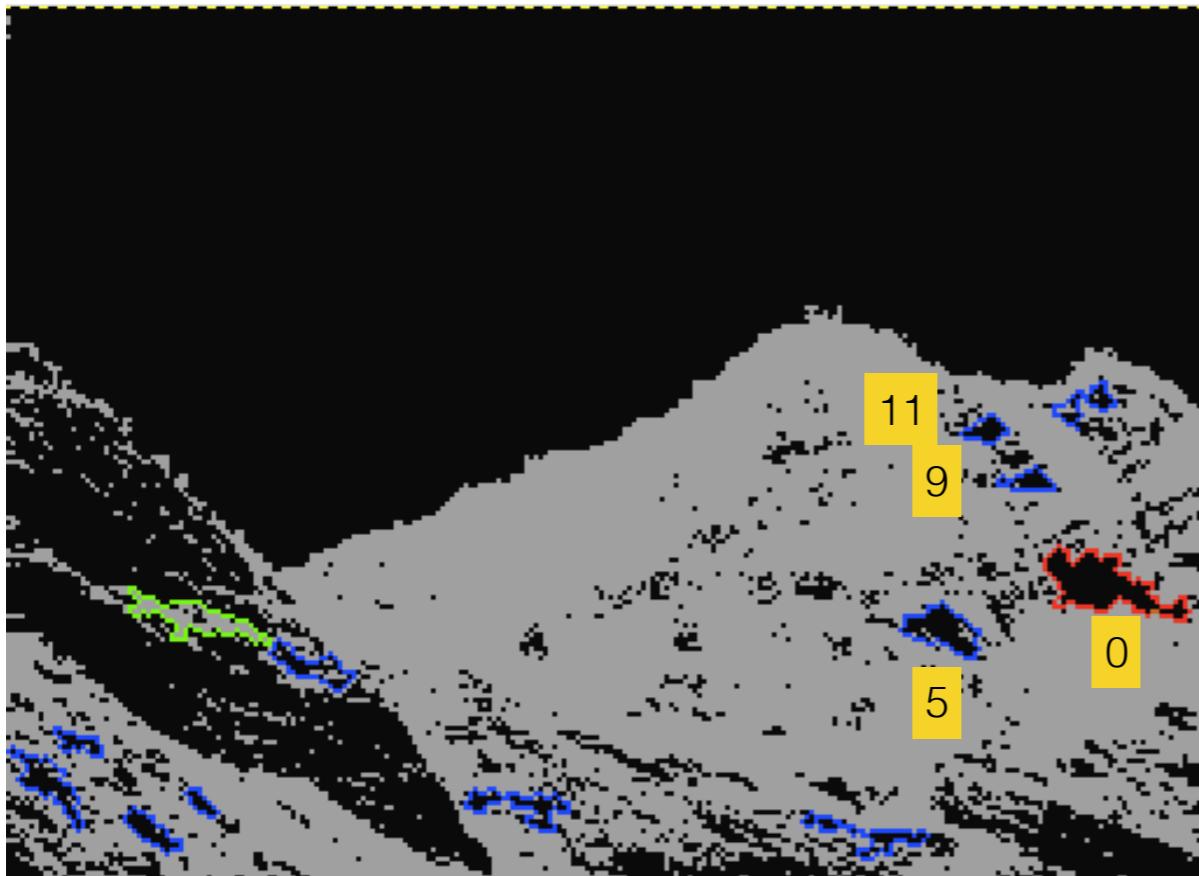
k=4 segmentation blobs



k=5 segmentation blobs



k=2 segmentation blobs, image binned to < 300x300



```
i1=0 i2=1 solvedScale=1.1637253761291504 cost=47.0  
i1=5 i2=4 solvedScale=1.2152479887008667 cost=0.0
```

There are errors in size due to contour level decision in making the segmented images.

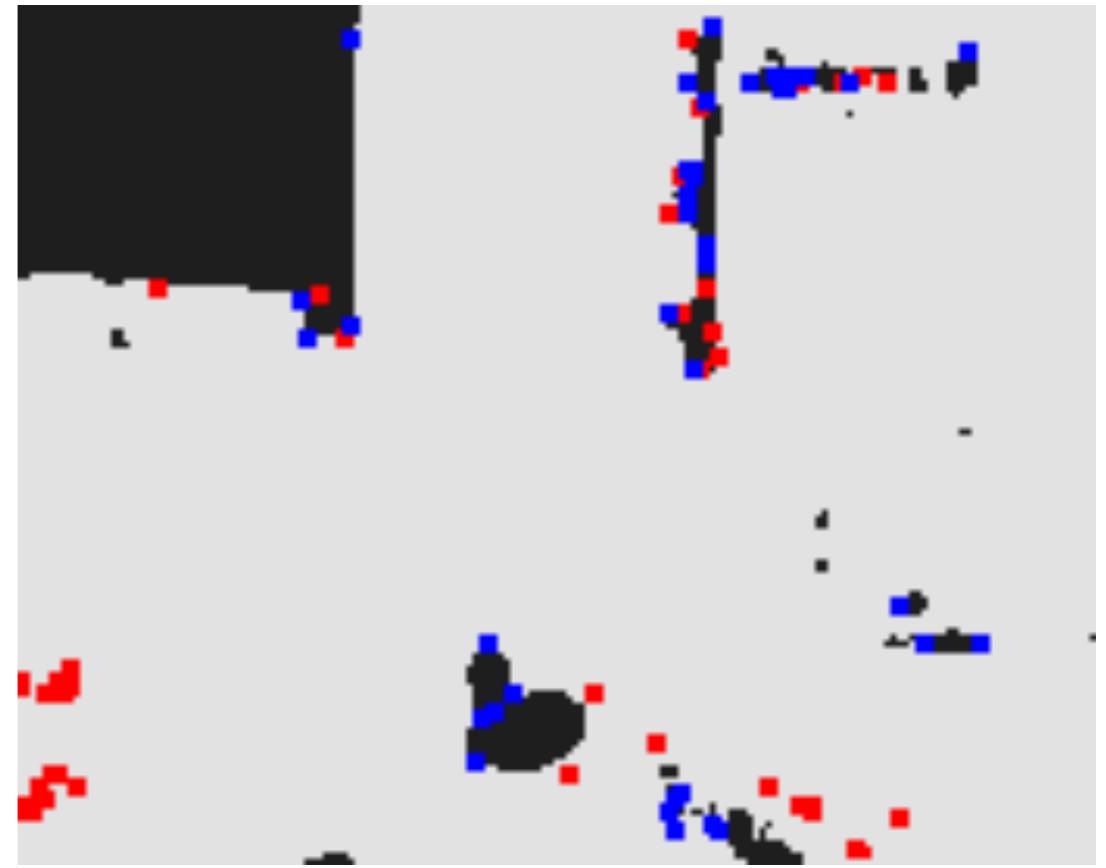
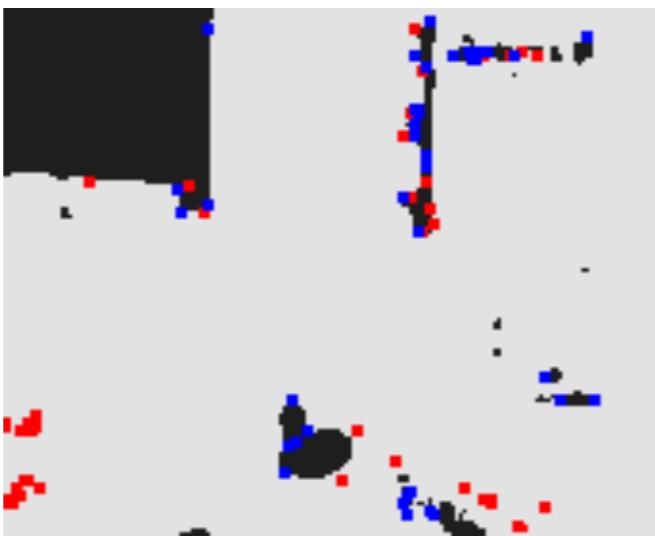
attempt to improve:

- (1) maybe choose the top 10 in this binned, k=2 segmentation image,
then extract the contours from just those in the same size images but higher segmentation k
- (2) OR, use k=2 segmentation for the full size image.
- (3) use k= higher than 2 and the full size image.

(it's beginning to sound like the characteristic scale response function, but these binned segmented images would not be best for that.)

In summary, need to use a higher k segmentation
and possibly the full sized image.

A quick look at color segmentation with k=3 followed by binary segmentation, binned to size < 200 x 200, followed by corners worked fine for 2 image sets, but not the third so the other methods in previous slides are preferred (corners filtered to blobs, matching, calculating transformation, then creating correspondence with all corners).



For the stereoscopic above, there's more than translation
so this can only be the start of the solution

previous notes from a
look at binning and
segmentation

